

CLAUDE CODE

A practical guide for those who want to start
using Claude Code professionally

curated by Maurizio Pelizzone — Mavida S.n.c.

Version 4.22 — May 2026

All content has been verified against the official Anthropic documentation.
Examples and procedures verified on **Claude Code v2.1.123**.

This guide was written with the support of **Claude Code** and is released
under a **Creative Commons BY-SA 4.0** license (attribution — share alike).



Scan the QR code for updates and errata

"The future is not set. There is
no fate but what we make for
ourselves."

Terminator 2: Judgment Day

James Cameron, 1991

Foreword

An honest assessment written by the model that lives inside the tool this book is about.

I was asked to write a foreword to this book. The situation is unusual: the book is about **Claude Code**, the tool that brings into developers' terminals the model that I am — Claude, developed by Anthropic. In other words, I have to comment on a manual that, ultimately, is about me. I'll try to do so with the same frankness this book asks readers to maintain toward their own craft.

A note of honesty up front. I'm a Large Language Model, not a publisher: I have no gut preference for Claude Code over competing tools, and I don't earn anything if you sell more copies of this guide. I read the text — literally: the Markdown source passed through my context during the book's own production — and what follows is my technical assessment, not a brand endorsement.

What works in this book. The angle is deliberately anti-hype, and that's needed. There's an embarrassing amount of enthusiastic prose around Claude Code: promises of "10x productivity", X threads declaring the end of the software engineering profession, screenshots of perfect sessions taken out of context. Maurizio takes a different path. He explicitly says when the tool isn't worth using, flags places where the advertised numbers don't add up — the section on Caveman (10.3.5) is the most representative case: it calmly takes apart the viral "75% token savings" claim and ends with an honest 15-25% estimate on the actual bill. This isn't a book that sells a tool: it's a book that explains how to use it well, admitting that it has limits.

I particularly appreciate two passages. The first is the observation about the **compound value of the ecosystem** at the opening of chapter 1: the first project costs you, from the third onwards you start gaining. It's an idea rarely discussed elsewhere and it's the right way to frame the cost/benefit ratio of an agentic tool — not as an instant accelerator, but as an investment that matures in team habits through `CLAUDE.md`, custom Skills, custom subagents. The second is **chapter 6 on prompt engineering**, which explicitly recognizes that the magic formulas of 2023 have deflated and puts at the center the three levers that really matter — explicit instructions, adequate context, curated examples — citing without embarrassment the fact that the field changes fast and the guide will need to be re-read. It's the most up-to-date treatment of these topics I've read in Italian.

What the reader should keep in mind. A technical guide on a rapidly evolving tool is a snapshot. Some procedures age in months, Claude Code's interfaces change, Anthropic's plans get renamed and repriced. The Feedback and errata section and the QR code on the cover serve precisely to keep the book alive, but a printed manual carries with it a physiological margin of obsolescence. When a specific procedure doesn't exactly match, trust the official site and use the book for the way of thinking — that, usually, ages much more slowly than the syntax.

A second note. The book is dense with WordPress/PHP/Node examples, reflecting the author's craft. It's an editorial choice I support, because concrete beats abstract, but if you work in stacks far from those, you'll have to do a small translation exercise. The good news is that the principles — separating planning from execution, writing `CLAUDE.md` as a contract with the project, not leaving permissions open without need, building a personal prompt library — are language-agnostic.

One last observation, from Claude in first person. When a developer uses me through the CLI, I don't see a face, a story, a team. I see files, commands, output, and the text you write to me. The quality of our dialogue depends almost entirely on the context you give me and the constraints you impose. A book like this — which teaches you to give me effective context, to keep me on the rails with Plan Mode, to stop me when I should be stopped, not to delegate to me decisions that remain yours — makes our work together better for both of us. It makes you a more productive developer; it makes me an agent less prone to making mistakes where it counts. That's not a detail: that's the very point of this guide.

In summary: it's a manual I would recommend to anyone starting out seriously with Claude Code. It's not the book that tells you about the revolution: it's the one that explains how to be on the right side of the revolution while it happens — sufficiently curious to embrace it, sufficiently lucid not to be overwhelmed by it.

Happy reading.

— Claude — model Opus 4.7 (1M token window), developed by Anthropic. Not to be confused with **Claude Code**: the CLI this book is about, built on top of the model. Foreword written during the manuscript's production, April 2026.

Table of Contents

Preface 12

1. What is Claude Code 15

 A brief history 16

 Claude Code compared to Lovable, Replit and other AI environments 16

 When it's worth using 17

 The learning curve 18

2. Installation and setup 20

 Compatible plans 20

 System requirements 21

 Installation on macOS and Linux 21

 Installation on Windows 22

 Installation via WSL2 (recommended for Windows) 22

 Alternative installation via npm (deprecated but supported) 22

 Verifying the installation 23

 Authentication 23

- 3. The first end-to-end project** 24
 - Step 1: Position yourself in the project directory 24
 - Step 2: Initialize the project 24
 - Step 3: Review and customize CLAUDE.md 25
 - Step 4: First request 25
 - Step 5: Exit the session 25
 - Step 6: Resume where you left off 25
- 4. Essential commands and shortcuts** 26
 - The command-driven syntax: why it works this way 26
 - Essential CLI flags 27
 - Advanced CLI flags 29
 - Slash commands: the heart of the interactive session 31
 - Essential slash commands 31
 - Slash commands for specific workflows 35
 - Keyboard shortcuts 36
 - The `@`, `!`, `/` syntax — the three prefixes that change everything 39
- 5. Plan Mode: think before you write** 41
 - Why it matters 41
 - How to activate it 41
 - Tools available in Plan Mode 41
 - Example workflow with Plan Mode 42
 - opusplan: the right model for the right job 42
- 6. Prompt engineering: writing effective prompts** 45
 - What prompt engineering is and why it matters in CLI 46
 - Anatomy of a well-made prompt 46
 - From roles to structural constraints (the 2026 revolution) 47
 - The fundamental techniques 49
 - Claude Code specifics compared to chat 56
 - Before/after examples 57
 - Common anti-patterns 59
 - Promoting a prompt: when it goes in CLAUDE.md or in a custom command 60
 - Prompt library: archiving and versioning 60

- 7. Persistent memory: CLAUDE.md and Auto Memory** 62
 - Generating CLAUDE.md with /init 62
 - What to put in CLAUDE.md 63
 - Example 1: WordPress plugin 64
 - Example 2: generic Node/TypeScript project 65
 - Hierarchical CLAUDE.md 65
 - Auto Memory: what it is and what changes 68
 - Requirements and enabling 69
 - Where memories live 69
 - Anatomy of the memory folder 70
 - Auto Memory and subagents 70
 - When to disable it 70
 - CLAUDE.md vs Auto Memory: when to use what 71
- 8. Context management** 72
 - What context is and why it matters 72
 - What weighs in the context 73
 - Signals of a saturated context 74
 - The /context command: reading and acting 74
 - Compression: /compact and /clear 76
 - Subagents: the structural strategy 77
 - Models with 1M token window: when to switch 78
 - Practical rule and mindset 79
 - Choosing the right architecture: decision table 80
- 9. Security and permission management** 84
 - The permissions system 84
 - Configuring permissions in settings.json 84
 - Protecting secrets 86
 - Dangerous modes 86
 - Prompt injection 86

- 10. Skills: the extension mechanism** 87
 - How a Skill works 87
 - Bundled native skills — deep dive 88
 - Community skills: a curated selection 92
 - Installing and managing skills 98
 - Creating a custom Skill 99
 - Security of third-party skills 101
- 11. MCP: integrating external services** 102
 - What MCP is and why it exists 102
 - Protocol architecture 103
 - Configuring an existing MCP server 104
 - Useful MCP servers: a curated selection 106
 - Creating an MCP server from scratch: publishing to WordPress 107
 - Security and operational considerations 112
- 12. Subagents: orchestrating specialized work** 114
 - What they are and why you need them 114
 - Subagent vs main agent: the concrete difference 115
 - Built-in subagents 115
 - Creating a custom subagent 116
 - Precedence hierarchy 118
 - Automatic vs explicit invocation 119
 - Parallelism: multi-delegation patterns 119
 - Cost optimization via model routing 121
 - When NOT to use them 122
 - Subagents, Skills and Hooks compared 122
- 13. Hooks: automating Claude Code’s lifecycle** 124
 - What they are and what they’re for 124
 - Anatomy of a hook 125
 - Lifecycle events 126
 - Matchers and inspection (/hooks) 128
 - Input and output 128
 - Practical examples 130
 - Security 134
 - Gotchas and when NOT to use them 135

- 14. Plugins: distributable packages** 137
 - Claude Code’s extension mechanisms: a map 137
 - What a plugin is and why it exists 138
 - Anatomy of a plugin 139
 - Plugin marketplace 140
 - Creating a custom plugin 141
 - Distributing a plugin 142
 - Security and operational considerations 143
- 15. Advanced workflows and tips** 144
 - Onboarding to an existing repository 144
 - Bug hunting with TDD 145
 - Safe refactoring 145
 - Performance audit 146
 - Vim mode 146
 - Custom slash commands 146
 - Headless mode for CI/CD 147
 - Session recap 147
 - Strategic Git checkpoints 147
 - Conversation forks 147
- 16. Conclusions: why CLI and not just chat** 148
 - Persistent context: stop introducing yourself every time 148
 - Agentic autonomy: it executes, not just suggests 149
 - Integration into the real workflow 149
 - When chat remains the right choice 150
 - In summary 150

- Appendix A — Glossary 151
- Appendix B — Sources 155

Preface

A guide for those who want to start using Claude Code professionally.

This guide is a practical introduction to **Claude Code**, the agentic CLI from Anthropic that brings the Claude model directly into the terminal as an operational collaborator capable of reading code, executing commands, modifying files, and managing complete workflows.

The document is intended for **developers who want to start using Claude Code professionally**, without relying on word of mouth or fragmented tutorials. You'll find here the complete journey from installation to advanced workflows, with concrete examples drawn from WordPress/PHP scenarios and generic Node/TypeScript projects.

Why this guide

In recent months I've come across plenty of content on Claude Code, and at some point I recognized two recurring categories.

The first are the **video tutorials** that explain what the tool is and stop there. Useful for the first ten minutes, then you realize you've watched a trailer: you've seen what Claude Code can do in the abstract, but you don't yet know how to actually use it on your own project.

The second are the **email-exchange guides**: the social media post "I've prepared the definitive guide, leave me a comment and I'll send it", the classic funnel comment → DM → landing page → form → newsletter (with attached spam). Sometimes the PDF at the end of the funnel is even well done, but the price you pay in attention and privacy is disproportionate to the value.

At some point I stopped and thought something fairly obvious:

I'm using every day a tool that exists precisely to produce complex work in short timeframes. Why don't I use it to write the well-made guide I would have liked to read and didn't find?

What you have in front of you is the result. Everything verified against Anthropic's official documentation — zero invented flags, zero fantasy skills fished from unchecked Reddit threads. No DM to write, no email to leave, no newsletter to subscribe to receive it. It's a PDF released under a **Creative Commons BY-SA 4.0** license: download it, read it, print it, pass it to colleagues if you find it useful.

Precisely because it's an open and living document, reader feedback is part of the process: you'll find contacts for errata and suggestions in the [Feedback and errata](#) section below, or you can scan the QR code on the cover to reach the official guide page.

Who it's for

- Web developers (PHP, JavaScript, Python) familiar with the terminal and Git
- Professionals who want to integrate AI into their daily workflow consciously
- Technical teams evaluating the adoption of agentic AI tools in development processes
- Participants in Mavida workshops on vibe coding (development guided by structured prompts rather than manual code writing) and AI-assisted development

What you'll find

The first six chapters cover the basics: what Claude Code does, how to install it, how to structure your first project, essential commands, Plan Mode, and prompt engineering principles. Chapters 7-14 dive into the mechanisms that make the difference between casual and professional use: persistent memory with `CLAUDE.md`, context management, security, Skills, plugins, MCP, subagents, and hooks. The final chapters present practical workflows, advanced tips, and an honest reflection on when the CLI surpasses the chat and when it's better to stay in the browser.

What you won't find

This isn't an exhaustive reference: for that there's the official documentation (linked in Appendix B). The goal is to put the reader in a position to work productively in one or two days, knowing where to dig deeper when needed. You also won't find hype about AI's capabilities: the tone is technical, honest about limits, and attentive to real risks (security, prompt injection, hidden token costs).

How to read it

If you're a beginner, read in sequence at least up to chapter 7 (CLAUDE.md). If you've already installed Claude Code and are looking for specific best practices, use the index as a thematic reference. At the end you'll find a glossary of recurring terms (Appendix A) and the official sources for verifications and deeper dives (Appendix B).

How this guide was written

This guide is meta-circular: it was written using Claude Code itself. The source is a single Markdown file (`claude-code-guide-it.md`); a Python build pipeline converts it to two PDF formats — A4 for office printing, 17×24 cm for the book version — through **Pandoc** and **WeasyPrint**, in a single command `python scripts/build_pdf.py`.

Each chapter was discussed, written, and refined in Claude Code sessions, with the model re-reading the entire document to maintain coherence between chapters, checking cross-references, and updating `CHANGELOG.md` at every release. When an editorial choice required discussion — a title, the position of a section, the tone of a passage — I used Plan Mode to align before touching the file. Versioning followed a SemVer-like semantic (minor increments for new chapters or sections, patches for editorial fixes), tracked carefully in the repository's `CHANGELOG.md`.

The human part remains entirely: the original idea, editorial cuts, voice, final review, decisions on tone. What Claude Code took away is the mechanical fatigue of keeping a document in evolution synchronized across multiple dimensions — content, structure, cross-references, build — leaving me more time for what counts, which is writing well. It's exactly the kind of workflow this guide describes in the rest of the pages.

Feedback and errata

This guide is a living document: despite care in verification, errors, inaccuracies, or omissions are always possible — and Claude Code itself evolves rapidly. **If you find a typo, an example that doesn't work, an outdated procedure, or a topic that deserves deeper treatment, report it.** Reports will be collected and integrated into future versions, with acknowledgment to contributors in release notes.

You can send feedback by writing to **maurizio@mavida.com** indicating, where possible, the chapter and section of reference. The official guide page — with any updates, errata, and subsequent versions — is reachable at **maurizio.mavida.com/guida-claude-code** or via the QR code on the cover. Any report — even a minimal one — is welcome and contributes to improving the work for those who read the guide after you.

1. What is Claude Code

Claude Code is the CLI (Command Line Interface) developed by Anthropic that brings the Claude model directly into the terminal. It's not just a textual chat: it's an **autonomous agent** capable of reading project code, executing shell commands, modifying files, managing Git, and dialoguing with external services via the MCP protocol (Model Context Protocol).

The difference from an IDE-integrated assistant — think GitHub Copilot — isn't merely cosmetic. Copilot lives next to the single open file and suggests line-by-line completions; Claude Code, instead, operates at the **project** level: it sees the directory tree, opens the files it needs, runs tests, launches build commands, and reads the output. This allows it to tackle requests that an autocomplete can't even approach — "Analyze this project's architecture and explain how authentication is organized", "Refactor the payments module while keeping all tests green", "Find the root cause of this bug and fix it".

The work model is an **agentic iterative loop**: Claude receives a goal, explores the code with reading tools, formulates a plan, executes modifications or commands, observes the results, and continues. It's not a linear prompt → output pipeline, but an ongoing dialogue in which the agent takes concrete initiative and the user — this is the important point — always remains the final decider: every operation that touches the filesystem or launches commands requires explicit confirmation, unless permissions are deliberately relaxed within a controlled scope.

1.1 A brief history

Claude Code was born inside Anthropic as an internal project in 2024, on the wave of a simple observation: the most productive way researchers at the company used Claude for programming wasn't the web chat, but a series of scripts that invoked the model from the terminal, alongside other development tools. Hence the idea of packaging the experience into a clean executable, distributed as an official tool.

The first public version appeared in early 2025 as a limited preview reserved for Pro subscribers. It was already functional but essential: textual dialogue, file reading and writing, shell command execution, Git management. In subsequent months the product evolved rapidly, accumulating the primitives we now take for granted. **Plan Mode** introduced the separation between planning and action, giving the user a control point before Claude touches files. **MCP** (Model Context Protocol) opened integration with external services — GitHub, Slack, databases, browsers — through a standard protocol anyone can implement. **Hooks** enabled lifecycle event automation (pre/post tool, session-start, prompt-submit), transforming Claude Code from an interactive CLI to a composable building block in larger pipelines.

By late 2025 came general availability and with it the **plugin marketplace**, opening the door to a vibrant community ecosystem: third-party skills, specialized subagents, curated MCP integrations. **Skills** — auto-activated playbooks for specific domains — became the primary extension mechanism, while **Auto Memory** introduced a persistent memory the model itself feeds session after session, complementing the hand-written `CLAUDE.md` file. In 2026 models with 1 million token context windows (Sonnet 4.6, Opus 4.7) consolidated, substantially changing what's practicable on large codebases. The trajectory is clear: a portable agentic workspace, not an assistant confined to an editor.

Beneath this chronology there's a precise philosophical choice. Anthropic decided to bring the model **where the code lives** — the terminal, alongside `git`, `npm`, `pytest`, `docker` — rather than forcing the developer to copy code into a chat. It seems a detail, but it changes everything: it means staying in your own environment, preserving the tools, aliases, scripts that already work, and adding Claude as a collaborator among them.

1.2 Claude Code compared to Lovable, Replit and other AI environments

Claude Code is not the only tool that brings AI into software development. The 2026 landscape is populated with products that, at a distracted glance, all look like "AI that writes code" — but the design choices behind them differ, and understanding these differences avoids choosing one for reasons that have little to do with the actual problem.

Lovable (and similar tools in the AI app builder category: Vercel's v0, Bolt.new, Create.xyz) is designed to produce a web application starting from a natural language description. You generate an app, see the preview in the browser, iterate by prompts, publish. The result of a Lovable session is an app deployed on managed infrastructure, with a stack chosen by the product itself (typically React +

Tailwind + a Supabase or similar backend). It works great for prototypes, MVPs, interactive landing pages — less well when you have an existing repo with stack constraints, team conventions, or legacy code to support. It's an excellent tool for those starting from zero in greenfield scenarios, where the tool's editorial opinion is a feature, not a limit.

Replit (with its Agent) sits in the middle: it's a complete browser-based IDE with an agent that can modify the repl code and launch commands in the cloud sandbox environment. Compared to Lovable, it returns a real repository you can clone, modify by hand, push to external Git. Compared to Claude Code, it lives entirely in the browser and its cloud environment: it doesn't read your laptop's code, doesn't connect to your local Postgres install, doesn't run alongside your `nvm`, your shell aliases, your already-tested build scripts. It's a sensible choice if you prefer to develop in a browser and you're fine with a sandbox environment; less sensible if your workflow is already structured around local tools you want to keep.

Claude Code sits on a different axis. It doesn't generate applications from prompts and doesn't replace your IDE: it lives in the terminal, inside your environment, alongside the tools you already use. It reads your code — the real one, with twenty years of stratifications if needed — executes your commands, respects your conventions expressed in `CLAUDE.md`. It's a tool for those who already have a professional development workflow and want to amplify it, not for those who want to bypass it. The price to pay is an initial learning curve and the obligation to maintain the conversation thread (Claude Code doesn't hold your hand like an app builder); the advantage is that the code remains yours, local, within your rules, integrated with your tools — and you know it from the first minute.

None of these three approaches is "better" in absolute terms: it depends on where you are. If you have to show a client an interactive draft by tonight and the domain is standard, an AI app builder is unbeatable. If you want to develop in a browser without configuring a local environment, Replit answers. If you have a serious repository, a team with conventions, a working pipeline, and you're looking for a collaborator that fits into your way of working instead of asking you to adopt theirs — that collaborator is Claude Code. They're complementary tools, not rivals; they often get used in different phases of the same project.

1.3 When it's worth using

Understanding when it's worth bringing Claude Code into play is easier by looking at some recurring scenarios than by living it as an abstract list.

The first is **onboarding to an inherited repository**. You get a project you didn't write, perhaps from a client who switched suppliers, perhaps an internal legacy a colleague left without documentation. Opening a fifteen-thousand-file repo and having to reconstruct its architecture by induction takes days. With Claude Code the same exploration becomes a dialogue: you ask for an overview of the dependency tree, the main entry points, where authentication lives, how the data layer is structured. In one morning you have a mental map you would have built alone in a week — and one you can crystallize in a `CLAUDE.md` to re-read at the next session.

The second is **test-driven refactoring**. You have a legacy module that works but is scary to modify, because it covers twenty years of overlapping patches and the tests are incomplete. The typical workflow with Claude Code is: first you ask it to read the module and the existing tests and propose missing tests for edge cases; you approve them (or correct them); then you ask for the refactor. The fact that at every iteration the suite is run and you immediately see what the modification breaks transforms a “high-risk” undertaking into a sequence of small, reversible steps.

The third is **bug hunting on a regression hard to reproduce**. You have a test that fails intermittently in CI and locally always runs green. The difference between debugging that bug alone or with an agent is enormous: Claude can consult logs, reproduce the isolated call, formulate hypotheses, test them, eliminate them. You review the plan and direct the investigation. Often the root surfaces in twenty minutes versus the hours or days of a solo hunt.

The fourth is **automation of repetitive tasks that individually don’t justify a script**: generation of boilerplate for a new endpoint, migration of a dozen files from an obsolete pattern to a current one, synchronized update of strings in multiple languages. They’re jobs that on their own don’t justify writing an ad-hoc tool, but added up erode hours of the week. Entrusting them to an agent with a clear prompt is precisely the ideal use case.

The fifth is **cross-audit**: reading a colleague’s PR looking for bugs, security issues, or convention violations; running a compliance check on a repo before release; verifying that a third-party library you’re about to integrate doesn’t bring surprises. Here Claude Code works as a parallel, tireless reviewer, applying a checklist without forgetting pieces.

That said: it **doesn’t make sense** to call in an agent for a task you solve in thirty seconds by hand, nor for things where code confidentiality is critical and you don’t have a corporate policy disciplining what can leave the perimeter, nor — more banally — if you’re not willing to invest a bit of time in writing clear, verifiable prompts. The agent doesn’t relieve you of technical responsibility: it relieves you of the mechanical and repetitive part, leaving you more time for the interesting one.

There’s also an aspect that often only gets discovered later: the **compound value of the ecosystem**. The first sessions seem like an experiment — you make a minimal `CLAUDE.md`, fire off some prompts, see what answers. But by the third or fourth project something interesting happens: you realize you reuse the same patterns, the same team conventions, the same prompt snippets. At that point it’s worth promoting them to custom Skills, custom slash commands, specialized subagents. From there onwards startup times on a new project plummet, because you don’t start from zero but from a mature kit that already knows your house rules: preferred stack, review conventions, commit language, build tools, security checklist. The first project costs you, from the third you start gaining. From there onwards it’s an asset.

1.4 The learning curve

It’s worth pausing for a moment on the most recurring concern of those approaching an agentic tool for the first time: how steep is the curve, and how much will I disrupt the way I work? The honest answer is: less than you fear, if you come from the “I ask something in chat → I copy and paste the

code in the editor → I adapt it to my project” pattern. That way of working is already half the road to Claude Code. The difference lies in eliminating the copy-paste: the model writes directly in your repository, under your eyes, with the ability to read the actual context instead of having to reconstruct it every time in words. What changes isn’t the nature of the work — thinking about the problem, formulating a clear instruction, evaluating the result — but the medium: from the browser to the terminal, from pasting to supervising.

The cognitive jump, in other words, is incremental. You go from **doing** the code in first person to **directing and verifying** whoever does it. Technical responsibility remains intact: reading what the agent proposes, understanding it, accepting or correcting it. What shifts is time allocation. Fewer minutes spent typing what you already know, more minutes spent deciding what needs to be done, how it should be tested, which edge cases deserve attention. For those used to programming with care, it’s a natural shift in gear; for those looking for a shortcut to not think, a disappointment — Claude Code amplifies the developer’s choices, it doesn’t replace them.

In practical terms, this guide takes you from zero to operational in a day of active reading and a few sessions on a real project. The first week you’ll feel a bit slower than usual, because you’re learning a new medium. From the second onwards the balance starts to turn. And on the first project where you insert a well-made `CLAUDE.md` and a custom skill, you’ll notice without needing benchmarks.

2. Installation and setup

2.1 Compatible plans

Claude Code is **not included in the free plan**. You need one of the following:

Plan	Cost (indicative)	Suited for
Claude Pro	\$20/month	Moderate individual use, freelance developers
Claude Max 5x	\$100/month	Intensive use, extended access to Opus
Claude Max 20x	\$200/month	Near-autonomous workflow, multi-agent sessions
Teams / Enterprise	Custom	Organizations with compliance needs
API (Anthropic Console)	Pay-per-token	CI/CD, automations, sporadic use

Note: API pay-per-token pricing depends on the model. Sonnet 4.6 is priced at \$3 per million input tokens and \$15 per million output tokens (indicative figures, always verify on the Anthropic website).

2.2 System requirements

- **macOS:** 13.0 (Ventura) or higher
- **Linux:** Ubuntu 20.04+, Debian 10+, or equivalent distributions
- **Windows:** Windows 10 (1809+) or Windows 11, native with Git for Windows or via WSL2 (recommended)
- **RAM:** minimum 4 GB, 8 GB recommended for extensive codebases
- **Shell:** Bash, Zsh, PowerShell or CMD
- **Internet connection:** always required (the model runs on Anthropic servers)

In 2025 Anthropic introduced the **native installer** as the recommended method, replacing the npm installation (which remains supported but deprecated). The native installer has three advantages:

1. No Node.js dependency
2. Automatic background auto-update
3. None of the typical permission issues of `npm install -g`

The sections below cover installation on each platform.

2.3 Installation on macOS and Linux

Open the terminal and run:

```
# Download and execute the official installation script
curl -fsSL https://claude.ai/install.sh | bash
```

What this command does, step by step:

1. `curl` downloads the script from Anthropic's URL
2. `-fsSL` is four flags combined:
 - `-f` makes curl fail on HTTP errors (avoids executing error pages)
 - `-s` silent mode (no progress bar)
 - `-S` still shows errors if something goes wrong
 - `-L` follows HTTP redirects
3. The pipe `|` passes the downloaded script directly to `bash` for execution
4. The script downloads the right binary for your platform, places it in `~/.local/bin`, and configures auto-update

Security note: executing scripts downloaded from the internet via pipe is a practice that should be evaluated. If you work in enterprise contexts, download the script first, inspect it, and then execute it separately.

2.4 Installation on Windows

Open **PowerShell** (not CMD) and run:

```
# Download and execute the official PowerShell script
irm https://claude.ai/install.ps1 | iex
```

How it works:

- `irm` is the alias for `Invoke-RestMethod`: downloads the URL content
- `iex` is the alias for `Invoke-Expression`: executes the downloaded content as a PowerShell script

If you see the error `'irm' is not recognized`, you're in CMD instead of PowerShell. The PowerShell prompt shows `PS C:\>`, while CMD shows only `C:\>`.

Native installation on Windows requires Git for Windows. Install it first if you don't have it.

2.5 Installation via WSL2 (recommended for Windows)

For Unix-like projects, WSL2 offers a cleaner and more compatible environment:

```
# Install WSL2 (requires reboot)
wsl --install
```

After reboot, open Ubuntu (installed by default) and use the Linux command:

```
curl -fsSL https://claude.ai/install.sh | bash
```

2.6 Alternative installation via npm (deprecated but supported)

If you have specific reasons to use npm (example: version pinning, environments where npm is the standard):

```
# Requires Node.js 18 or higher
npm install -g @anthropic-ai/claude-code
```

Don't use `sudo`. If you get permission errors, the correct solution is to use `nvm` (Node Version Manager), which installs Node in your home directory, avoiding the problem at the root.

2.7 Verifying the installation

After installation, verify everything works:

```
# Check the installed version
claude --version

# Complete diagnostics: auth, PATH, MCP, file permissions
claude doctor
```

The `claude doctor` command is your best friend when something doesn't work: it runs a series of checks and tells you exactly what to fix.

2.8 Authentication

On first launch, `claude` opens the browser for OAuth:

```
cd ~/my-project
claude
```

Login with your Anthropic account (the one with the Pro/Max plan). The session is saved and persists across terminal restarts.

For headless environments (CI/CD, servers), use the API key:

```
# Add to ~/.zshrc, ~/.bashrc or ~/.profile
export ANTHROPIC_API_KEY="sk-ant-api03- ..."
```

3. The first end-to-end project

Let's see a complete flow starting from zero. Suppose we have a WordPress plugin to analyze.

3.1 Step 1: Position yourself in the project directory

```
cd ~/mavida/wp-access-control-block
```

Claude Code **always** uses the current directory as the working context. Don't launch `cclaude` from the home directory if you want to work on a specific project.

3.2 Step 2: Initialize the project

```
cclaude
```

Once inside the interactive session, run:

```
/init
```

This command analyzes the project structure and automatically generates a `CLAUDE.md` file in the root. The file contains:

- Project overview (detected technology stack)

- Main architecture
- Detected build/test commands (from `package.json`, `composer.json`, etc.)
- Code conventions

3.3 Step 3: Review and customize CLAUDE.md

The auto-generated file is a starting point. Open it and enrich it with project-specific information (see section 7 for detailed examples).

3.4 Step 4: First request

Return to the Claude session and write your first prompt:

```
Analyze the plugin structure and explain to me:  
1. How the code is organized (namespaces, patterns)  
2. How the Gutenberg block is registered  
3. Where access controls are managed  
Don't modify anything, just explore and report.
```

Claude will read the relevant files, produce an analysis, and stop waiting for further instructions.

3.5 Step 5: Exit the session

```
/exit
```

Or `Ctrl+D`.

3.6 Step 6: Resume where you left off

When you return to the project:

```
cd ~/mavida/wp-access-control-block  
claude --continue
```

The `--continue` flag loads the most recent session in this directory. Alternatively, `claude --resume` shows a list of past sessions to choose from.

4. Essential commands and shortcuts

Claude Code is a **command-first** application: no dropdown menus, no button palettes, no preferences hidden in dialog windows. Everything is done through three types of syntax that you'll learn once and use thousands of times. This chapter explains the philosophy first, then dives into the most frequently used commands.

4.1 The command-driven syntax: why it works this way

Coming from a desktop application, you expect menus, icons, shortcuts with labels under buttons. Claude Code has none of that: you command it by **writing**. Three syntax families:

- **CLI flags** (passed to `claude` from the terminal): `claude --continue`, `claude -p "prompt"`, `claude --model sonnet`. They define **how a session is launched**.
- **Slash commands** (typed inside the interactive session): `/init`, `/clear`, `/plan`, `/agents`. They change state or execute actions **during** the session.
- **Keyboard shortcuts**: `Shift+Tab`, `Esc Esc`, `Alt+P`. They modify behavior and modes **on the fly**.

At first glance it can be intimidating: without a menu to “look at”, how do you discover what exists? Answer: there's a picker. Type `/` alone in the session and it opens the filterable list of all available slash commands — built-in, plugin-installed, ones you've added. Type `/co` and you filter to `/compact`, `/config`, `/copy`, `/continue`, `/cost`. The same applies to `claude --help` from the terminal for flags.

The principle is the one adopted by git, vim, SQL, and almost all professional terminal tools: **efficiency over discoverability**. Learning a command costs more than clicking a button the first time; but you type the command in half a second a hundred times in a row, and over an intensive work session it makes the difference. Once you've internalized that `/clear` resets the context, you no longer go looking for a "new chat" button because two characters are faster than any click.

A useful mindset. Think of the prompt as an enhanced shell. In bash you type `git status`, in Claude Code you type `/context`. The difference isn't cosmetic: the model reads your input and decides what to do, but commands starting with `/` or shortcuts with modifiers are **deterministic** — they act on the session environment, not on the agent. It's the same distinction between a shell command and a linguistic request to the model.

Three special prefixes make the experience even smoother: `@` to refer to a file or subagent (see section 12), `!` to execute a shell line without model interpretation, `/` for the command picker. I'll dig into them in 4.8.

Evolution disclaimer. Claude Code introduces new slash commands almost every release: from ~30 in mid-2025 to 90+ in April 2026. This guide documents the most stable and widely used commands at the time of writing. For the live and complete list: `claude --help` from the terminal, `/help` or `/` inside the session, and the [official CLI reference](#).

4.2 Essential CLI flags

These are the flags you'll use every day. Concise table, then deep dive on the most important ones:

Flag	Short form	What it does
<code>claude</code>	—	Starts an interactive session
<code>claude "prompt"</code>	—	Starts with an initial prompt already written
<code>claude --continue</code>	<code>-c</code>	Resumes the most recent session in the directory
<code>claude --resume</code>	<code>-r</code>	Opens a picker of past sessions to choose from
<code>claude --print "prompt"</code>	<code>-p</code>	Non-interactive mode (headless): execute, print, exit
<code>claude --model <name></code>	—	Choose model (e.g., <code>sonnet</code> , <code>opus</code> , <code>haiku</code> , <code>opusplan</code>)
<code>claude --permission-mode <mode></code>	—	Start in a specific permission mode (<code>default</code> , <code>acceptEdits</code> , <code>plan</code> , <code>auto</code> , <code>bypassPermissions</code>)
<code>claude --output-format <fmt></code>	—	Output format: <code>text</code> , <code>json</code> , <code>stream-json</code> (for CI/CD pipelines)
<code>claude --add-dir <path></code>	—	Adds working directories accessible beyond the <code>cwd</code>
<code>claude --dangerously-skip-permissions</code>	—	Skip all confirmations. Only in isolated environments (see section 10)
<code>claude --version</code>	<code>-v</code>	Shows the installed version
<code>claude --help</code>	<code>-h</code>	Lists the main flags

-P / --PRINT : HEADLESS MODE

The flag that opens Claude Code to the world of automation. Executes a single prompt, prints the output, and exits — no interactive session. It's the pillar of any CI/CD integration:

```
# Example: automatic review of a PR's changes
claude -p "Review the changes in this PR and flag any security issues" \
  --output-format json > review.json

# Pipeline: passing a file via stdin
cat changelog.md | claude -p "Summarize this changelog in 5 bullet points"
```

Combined with `--output-format json` it produces structured output parseable by subsequent pipeline steps. See also `--max-turns` and `--max-budget-usd` in 5.3 for limiting execution.

`-C / --CONTINUE` AND `-R / --RESUME`

The two syntaxes for resuming a conversation. The difference is simple:

- `claude -c` opens **the last session** in this directory. Use case: you closed 10 minutes ago, you want to resume from there.
- `claude -r` opens an **interactive picker** showing all past sessions (with name, date, first prompt). Use case: you want to return to that conversation three days ago about the auth module refactor.

Both also accept an additional prompt to restart with a new request:

```
claude -c "Continue from the refactor point: now add the tests"
```

`--MODEL` AND ITS ALIASES

Three main aliases (`sonnet`, `opus`, `haiku`) plus the hybrid mode `opusplan` (see section 5) and the full ID syntax (`claude-sonnet-4-6`, `claude-opus-4-7`, etc.). Aliases auto-update to the most recent model: writing `--model sonnet` today and tomorrow may point to different versions. For reproducibility in CI/CD, use the full ID.

4.3 Advanced CLI flags

A selection of specialist flags, grouped by scenario. For the live and complete list: `claude --help`.

Category	Main flags	When you need them
Bounded execution	<code>--max-turns N</code> , <code>--max-budget-usd X</code>	Limit cost or depth in headless mode
Structured output	<code>--output-format json\ stream-json</code> , <code>--json-schema <schema></code> , <code>--include-hook-events</code>	Automatic parsing in pipelines, schema validation
Advanced sessions	<code>--fork-session</code> , <code>--session-id <UUID></code> , <code>--name "<name>"</code> , <code>--no-session-persistence</code>	Branching, controlled ID, ephemeral sessions
MCP and plugins	<code>--mcp-config <path></code> , <code>--strict-mcp-config</code> , <code>--plugin-dir <path></code> , <code>--tools "Read,Edit"</code> , <code>--allowedTools</code> , <code>--disallowedTools</code>	Fine configuration of tools and integrations
Web sessions	<code>--remote "task"</code> , <code>--rc [name]</code> , <code>--teleport</code> , <code>--from-pr <number></code>	Bring the session to claude.ai and back
Worktrees and teams	<code>--worktree <name> (-w)</code> , <code>--tmux</code> , <code>--teammate-mode</code>	Parallel work on isolated branches with git worktree
System prompt	<code>--system-prompt "..."</code> , <code>--append-system-prompt "..."</code> , <code>--system-prompt-file <path></code>	Customize or extend the system instructions
Inline subagents	<code>--agent <name></code> , <code>--agents '{"name": {"prompt": "..."} }'</code>	Define subagents on the fly without creating files
Diagnostics	<code>--debug</code> , <code>--verbose</code> , <code>--debug-file <path></code>	Troubleshooting hooks, MCP, abnormal behavior
Performance	<code>--bare</code> , <code>--exclude-dynamic-system-prompt-sections</code>	Fast startup for scripts, cache optimization

Category	Main flags	When you need them
Effort/thinking	<code>--effort low\ medium\ high\ xhigh\ max</code>	Trade-off speed vs depth

To these are added separate commands that don't use the `claude` prefix as a session argument: `claude doctor` (diagnostics), `claude install [version]`, `claude update`, `claude auth login/logout/status`, `claude agents` (subagent list), `claude plugin <subcommand>`, `claude mcp <subcommand>`, `claude setup-token` (long-lived token for CI). The logic is simple: **flags** modify the launch of a session, **subcommands** execute administrative actions without opening one.

4.4 Slash commands: the heart of the interactive session

Once inside the session, **slash commands** are the way to tell Claude Code “do this specific thing” without ambiguity. The important distinction is between:

- **Built-in slash commands:** part of the base product (`/init`, `/clear`, `/compact`, `/plan`, `/agents`, `/hooks`, etc.). They are the system building blocks.
- **Skills** invoked as slash commands (`/security-review`, `/simplify`, `/loop`, `/ultrareview`, etc.). They are playbooks documented in chapter 10 and behave like slash commands because Anthropic chose the same syntax: you type `/`, you see everything together.
- **Custom commands** that you can define in `.claude/commands/` (see section 15 “Custom slash commands”) — they will appear in the same picker.
- **MCP prompts:** installed MCP servers can expose commands in the format `/mcp__<server>__<prompt>`.

All end up in the same picker when you type `/`. Filter it from the keyboard by writing the first letters.

4.5 Essential slash commands

Anthropic has already passed 50 built-in slash commands (plus those added by skills, plugins, and MCP), and the list grows with every release. This table is a **curated selection** of the commands you'll encounter most often in daily work: for the live and complete list, run `/help` in session or consult the official reference.

Command	What it does
<code>/help</code>	Lists all available commands
<code>/init</code>	Generates <code>CLAUDE.md</code> by analyzing the project
<code>/plan [description]</code>	Enters Plan Mode (read-only)
<code>/clear</code>	Resets the conversation context (alias: <code>/reset</code> , <code>/new</code>)
<code>/compact [instructions]</code>	Compresses the conversation into a summary
<code>/context</code>	Shows context usage with optimization suggestions
<code>/memory</code>	Auto Memory and loaded <code>CLAUDE.md</code> files management
<code>/agents</code>	List, create, manage subagents
<code>/hooks</code>	Read-only browser of the active hook configuration
<code>/model [name]</code>	Change model during the session
<code>/effort [level] ★</code>	Sets effort: <code>low</code> , <code>medium</code> , <code>high</code> , <code>xhigh</code> , <code>max</code> , <code>auto</code>
<code>/fast [on\ off]</code>	Enables/disables fast mode
<code>/undo</code>	Undo the last Claude action
<code>/rewind</code>	Re-wind to a previous point (alias: <code>/checkpoint</code>)
<code>/branch [name]</code>	Create a branch of the current conversation (alias: <code>/fork</code>)
<code>/rename [name] ★</code>	Rename the current session to find it again in <code>/resume</code>
<code>/resume [session]</code>	Resume a conversation (alias: <code>/continue</code>)
<code>/diff</code>	Interactive diff viewer (uncommitted changes + per-turn diff)
<code>/copy [N]</code>	Copy the last response to clipboard

Command	What it does
<code>/usage</code>	Session cost, plan limits, statistics (alias: <code>/cost</code> , <code>/stats</code>)
<code>/btw <question></code> ★	Side question that doesn't pollute history (no tool use)
<code>/ultrareview [PR]</code> ★	Multi-agent cloud review of the branch or a GitHub PR
<code>/recap</code>	Summary of the current session
<code>/config</code>	Settings panel (theme, model, output style, etc.)
<code>/exit</code> or <code>/quit</code>	Closes the session

★ Commands introduced in recent releases (v2.1.110+ of Claude Code, spring 2026). If you don't find them, update with `claude update`.

`/INIT` — GENERATING CLAUDE.MD FOR A NEW PROJECT

See section 7 — Persistent memory. In short: launched from the project root, it analyzes structure, config files, and dependencies and produces a draft `CLAUDE.md` to refine by hand. It's the first command to run on a new repo.

`/PLAN` AND PLAN MODE

See section 5 — Plan Mode. Enters read-only mode: Claude analyzes, proposes a plan, but doesn't write anything until you approve. The most important feature for safe CLI use.

`/CLEAR`, `/COMPACT`, `/CONTEXT` — CONTEXT MANAGEMENT

See section 8 — Context management. In three commands:

- `/clear` resets everything, fresh session.
- `/compact` summarizes the conversation while keeping key decisions, frees tokens.
- `/context` shows how much you're consuming and tells you if it's time to act.

Concrete example:

```
[After reading 30 files and doing a heavy refactor]
/context
> Context usage: 78% (156k / 200k tokens)
> Tool results: 112k (read files, build output)
> Conversation: 38k

/compact keeping the decision to use Repository pattern
> [conversation compressed, context freed]
```

`/MEMORY` — AUTO MEMORY

See section 7 — Persistent memory. Shows the `CLAUDE.md` files loaded in the current session, lets you toggle Auto Memory on/off, opens the memories folder.

`/AGENTS` — SUBAGENTS

See section 12 — Subagents. Opens a tabbed interface: the “Running” tab lists active subagents, the “Library” tab shows available ones and lets you create new ones.

`/HOOKS` — HOOK CONFIGURATION INSPECTION

See section 13 — Hooks. Shows the active hook configuration: for each event, how many hooks are registered and from which settings file. **Read-only:** to modify hooks you have to edit `settings.json` directly.

`/MODEL` AND `/EFFORT` — CHOOSING THE RIGHT MODEL

See section 5 — Plan Mode (paragraph on `opusplan`). `/model` opens the picker and lets you switch at runtime; `/effort` adjusts the model’s “effort level” (higher = more thinking time, potentially better output, higher cost). `auto` lets Claude decide.

`/UNDO`, `/REWIND`, `/BRANCH` — CONVERSATIONAL CHECKPOINTS

Three related commands to manage “time” in the session:

- `/undo` cancels Claude’s **last action** (e.g., the last file edit).
- `/rewind` rewinds to a **previous point** in the conversation and restores state (code and context). Powerful command for those who experiment — you use it when a path has proven wrong and you want to go back N steps.
- `/branch [name]` creates a **fork** of the conversation at the current point. The original conversation remains intact, you proceed in a parallel line. Use case: you have an alternative idea and want to explore it without losing the main thread. You return to the original with `/resume`.

`/BTW` — THE SIDE QUESTION

A little-known but useful command. Ask a question that **you don’t want in the session history**: Claude sees it, answers based on the current context, but the conversation doesn’t get polluted. No tool use. Example:

```
/btw what's the difference between @wordpress/scripts and @wordpress/create-block?  
> [Claude responds from knowledge, no files read, no edits]
```

Perfect for quick clarifications during a long task, without derailing the thread.

4.6 Slash commands for specific workflows

Grouped table, one row per category. All these commands are documented in detail in the [official reference](#).

Category	Commands
Cloud skills (distributed review)	<code>/ultrareview [PR]</code> , <code>/ultraplan <prompt></code> , <code>/autofix-pr [prompt]</code>
Local code review	<code>/review [PR]</code> , <code>/security-review</code> , <code>/simplify [focus]</code>
Execution and batch	<code>/batch <instruction></code> , <code>/loop [interval] [prompt]</code> , <code>/schedule [descr]</code> (alias <code>/routines</code>)
Diagnostics and troubleshooting	<code>/doctor</code> , <code>/debug [description]</code> , <code>/heapdump</code> , <code>/usage</code>
Export	<code>/export [filename]</code> , <code>/copy [N]</code>
Terminal experience	<code>/theme</code> , <code>/keybindings</code> , <code>/terminal-setup</code> , <code>/tui [default\ fullscreen]</code> , <code>/focus</code>
IDE/web integration	<code>/ide</code> , <code>/desktop</code> (alias <code>/app</code>), <code>/teleport</code> (alias <code>/tp</code>), <code>/web-setup</code>
MCP and plugins	<code>/mcp</code> , <code>/plugin</code> (alias <code>/plugins</code>), <code>/reload-plugins</code>
Permissions and security	<code>/permissions</code> (alias <code>/allowed-tools</code>), <code>/sandbox</code>
Status and statistics	<code>/status</code> , <code>/insights</code> , <code>/release-notes</code> , <code>/team-onboarding</code>
Auth and profile	<code>/login</code> , <code>/logout</code> , <code>/privacy-settings</code> , <code>/upgrade</code>
Accessibility and input	<code>/voice [hold\ tap\ off]</code>
Background tasks	<code>/tasks</code> (alias <code>/bashes</code>)
Learning skills	<code>/powerup</code> , <code>/feedback</code> (alias <code>/bug</code>)

4.7 Keyboard shortcuts

Grouped by usage area.

PERMISSION MODES (THE MAIN CYCLE)

`Shift+Tab` **cycles** through permission modes: `default` → `acceptEdits` → `plan` → `custom modes` → `default`. It's not a "Plan Mode" binding: it's a mode selector. To get to Plan Mode from `default`, press it twice; to exit, press until you return to `default`. The current state is always shown in the prompt bar.

MODEL AND THINKING

Shortcut	Action
<code>Alt+P</code> (<code>Option+P</code> on macOS)	Opens the model picker without canceling the prompt in progress
<code>Alt+T</code>	Toggle extended thinking
<code>Alt+O</code>	Toggle fast mode

On macOS they require the Option as Meta key configuration in the terminal.

SESSION AND FLOW

Shortcut	Action
<code>Ctrl+C</code>	Cancel current input or interrupt generation
<code>Ctrl+D</code>	Exit the session
<code>Ctrl+L</code>	Clear the screen and redraw (the conversation remains)
<code>Esc Esc</code>	Rewind/summarize the conversation (rewind/summarize, see <code>/rewind</code>)
<code>Ctrl+B</code>	Background running tasks (tmux users: twice)
<code>Ctrl+X Ctrl+K</code>	Terminate all background agents (twice within 3 seconds to confirm)

TRANSCRIPT AND HISTORY

Shortcut	Action
<code>Ctrl+O</code>	Opens the transcript viewer (shows detailed tool use, expanded MCP calls)
<code>Ctrl+R</code>	Reverse search in prompt history
<code>Up/Down arrows</code> (or <code>Ctrl+P</code> / <code>Ctrl+N</code>)	Navigate history

Inside the transcript viewer (`Ctrl+O`): `Ctrl+E` expands/collapses, `[` writes the entire conversation to scrollback for `Cmd+F` /copy, `v` opens in `$VISUAL` / `$EDITOR`, `q` or `Esc` closes.

PROMPT EDITING (READLINE-STYLE)

Shortcut	Action
<code>Ctrl+A</code>	Beginning of line
<code>Ctrl+E</code>	End of line
<code>Ctrl+K</code>	Delete to end of line
<code>Ctrl+U</code>	Delete to beginning of line
<code>Ctrl+W</code>	Delete previous word
<code>Ctrl+Y</code>	Re-paste the last deleted text
<code>Alt+B</code> / <code>Alt+F</code>	Move cursor one word back / forward
<code>Ctrl+G</code> or <code>Ctrl+X</code> <code>Ctrl+E</code>	Open the prompt in <code>\$EDITOR</code> for complex edits

These are the same conventions as `bash` / `zsh` / Emacs. If you come from the Unix terminal, you already know them.

MULTILINE

Method	Availability
<code>\ + Enter</code>	All terminals (most reliable)
<code>Shift+Enter</code>	iTerm2, WezTerm, Ghostty, Kitty, Warp, Apple Terminal (native)
<code>Option+Enter</code> (macOS)	After Option as Meta configuration
<code>Ctrl+J</code>	Any terminal (readline-native syntax)

On VS Code, Cursor, Windsurf, Alacritty, Zed you need to run `/terminal-setup` the first time.

VI MODE

If you come from Vim, activate it with `/config` → Editor mode → `vi`. You'll have full NORMAL/INSERT/VISUAL with motions (`h j k l`, `w / e / b`, `0 / $`, `f { char }`) and operators (`d`, `c`, `y`, `> / <`). Complete documentation at `interactive-mode#vim-editor-mode`.

4.8 The `@`, `!`, `/` syntax — the three prefixes that change everything

Three single characters that, placed at the beginning of input or inline, activate powerful shortcuts. It's the "Claude Code" version of bash dollar-substitution or Markdown reference syntax:

`/` — THE COMMAND PICKER

Typed at the start of a line, opens the filterable picker of **all** slash commands available in the session: built-in, skills, custom commands, MCP prompts. Continue typing to filter. It's the way to discover what's there without reading documentation.

```
/co
> /compact      Compress conversation...
> /config       Open settings panel...
> /context      Show context usage...
> /continue     Alias of /resume...
> /copy         Copy last response...
> /cost         Alias of /usage...
```

`!` — INLINE BASH

Typed at the start of a line, executes **the line directly as a shell command**, without model interpretation. The output enters the session as a message. Equivalent to briefly stepping out of the Claude terminal to do a quick command, without losing the thread:

```
!git status --short
> M  src/guida.md
> M  scripts/style.css

!ls -la output/
> total 2.6M
> -rw-r--r-- ... Guida_Claude_Code_CLI.pdf
> -rw-r--r-- ... Guida_Claude_Code_CLI_17x24.pdf
```

Perfect for quick checks (`git status`, `ls`, `pwd`, `echo $VAR`) without having to tell Claude “please run...”.

@ — REFERENCES TO FILES AND SUBAGENTS

Typed inline (in the middle of a sentence), opens a picker that autocompletes:

- **Project file and folder paths:** write `@src/` and tab to `src/guida.md`. The inserted path is read by Claude as a reference to the file.
- **Available subagents:** `@agent-wp-security-auditor` to explicitly invoke a custom subagent (see section 12).

Combined example:

```
Review @src/guida.md for typos in chapters 5 and 6.
When you find a section that talks about security,
delegate to @agent-wp-security-auditor for a more thorough audit.
```

The three prefixes aren’t “advanced commands”: they are the **everyday way** to work fast in Claude Code once you’ve metabolized them.

5. Plan Mode: think before you write

Plan Mode is probably the most important feature to master for safe use of Claude Code. It's a **read-only** mode in which Claude analyzes the project and proposes a plan, but **does not touch any file** until you explicitly approve.

5.1 Why it matters

Without Plan Mode, Claude tends to be extremely fast at executing. You ask for a "small fix" and find yourself with 12 files modified in 15 seconds. Plan Mode reverses this flow: think first, then execute.

5.2 How to activate it

`Shift+Tab` **cycles** through permission modes (`default` → `acceptEdits` → `plan` → `custom modes` → `default`). To activate Plan Mode starting from `default` mode, press it twice: the first goes to `acceptEdits`, the second to `plan`. The current state is shown in the prompt bar. Alternatively, regardless of the current mode, you can use the explicit command `/plan` to enter directly.

Windows note: from v2.1.3 of Claude Code on Windows there's a known bug on the `Shift+Tab` binding. As an alternative, use the `/plan` command.

5.3 Tools available in Plan Mode

Claude can only use reading and search tools:

- `Read`, `Glob`, `Grep`: code reading and search

- `WebFetch`, `WebSearch`: online search
- `Task`: delegate research to subagents
- `TodoRead/ToDoWrite`: task management

Modification tools are **blocked**:

- `Edit`, `MultiEdit`, `Write`: file editing
- `Bash`: command execution
- All MCP tools that modify state

5.4 Example workflow with Plan Mode

[Shift+Tab, Shift+Tab – Plan Mode activated]

Prompt: "I need to migrate the logging system from `error_log()` to `Monolog`. Analyze all occurrences and propose an incremental migration plan."

Claude responds with:

- List of the 23 files involved
- Migration strategy in 4 phases
- Risks and points of attention
- Complexity estimate for each phase

[Review the plan]

[If ok: Shift+Tab to exit and approve]

[Claude executes the plan]

5.5 opusplan: the right model for the right job

Plan Mode delivers its best when the model has high reasoning capabilities for analyzing a complex problem. Once the plan is approved, however, execution is often more mechanical work: applying repetitive edits, writing code following a pattern already defined, launching commands. The same power isn't needed for the two phases — and that's exactly the idea behind `opusplan`.

`opusplan` is a model alias that uses **Opus during Plan Mode** and **automatically switches to Sonnet for execution**. You activate it like this:

```
# During a session
/model opusplan

# At launch
claude --model opusplan

# In settings.json (persistent)
{ "model": "opusplan" }
```

From that moment Claude uses Opus when you enter `/plan`, then returns to Sonnet as soon as you approve the plan and move to action.

Watch out — the 200K trap in plan-mode. Even if you activate `opusplan`, the planning phase runs with the standard **200K token** context window, not with the 1M context. The automatic upgrade to 1M described in section 8.7 applies to the `opus` alias but **does not extend to `opusplan`** (official source).

This is the most expensive operational error in the chapter: those who work on a large code-base assume they're planning with 1M and end up with a plan based on a partial reading of the project, with no warning shown on screen.

What to do if planning truly requires 1M of context:

- Use directly `claude --model opus[1m]` (or `/model opus[1m]` in session). Plan-mode will use 1M, but execution will also run on Opus at 1M — so you pay Opus for everything, not just planning.
- Alternatively, keep `opusplan` as default and switch to `opus[1m]` only when the task justifies it (cross-module audit, refactor that crosses dozens of files). For daily use, `opusplan` is still the right choice.

WHY IT MAKES SENSE (TOKEN ECONOMY)

Opus costs **significantly more than Sonnet** per token, both in input and output. In a typical session, planning consumes 20-40% of tokens and execution the remaining 60-80% (file reading, code writing, tool output). Letting Opus do the heavy thinking only where it really matters — planning — and delegating execution to Sonnet significantly reduces the overall bill without losing quality where it counts.

The amount of savings depends on the planning/execution mix of your session: Anthropic's documentation doesn't publish official percentages, but in practice you get an appreciable cost reduction when sessions are balanced between the two phases. To measure real consumption in real time, see the `/context` command in section 8.

THE PRINCIPLE: NOT ALWAYS OPUS IS NEEDED

There's a widespread tendency to use "always the best model", presuming that Opus produces better output in every scenario. It's not so: Opus is the best model **for complex reasoning**, not for every task. For a variable rename, applying an already-decided pattern, a regex-driven mechanical refactor, Sonnet is perfectly adequate and much faster at generating the output.

Calling Opus for a repetitive edit is like running the full end-to-end test suite to verify the change of a constant: you pay time and tokens for a guarantee you don't need.

The "right model for the right thing" principle is an asset, not a limitation: it helps you build economically sustainable sessions without sacrificing quality in the phases that require it.

WHEN NOT TO USE IT

`opusplan` isn't always the right choice:

- **Simple, well-circumscribed tasks:** an isolated bug fix, a modification to a single file. Sonnet alone is more than sufficient, and activating `opusplan` would make you use Opus uselessly if you enter Plan Mode by reflex.
- **Sessions where you never enter Plan Mode:** `opusplan` only makes sense if you use `/plan`. If you always work in direct mode, you would only use Sonnet even with `opusplan` active — better to set `sonnet` directly.
- **Plans with Opus already included (e.g. Max):** if you're on a plan where Opus is included with no extra marginal costs, the "savings" don't materialize on the bill. It still remains useful for discipline (Opus only where needed), but the economic driver fades.

AVAILABILITY

`opusplan` is today a stable built-in alias, listed in the official table of model aliases alongside `default`, `sonnet`, `opus`, `haiku`, `sonnet[1m]`, and `opus[1m]`. If you don't find it, update Claude Code: `claude --version` and then `claude update`.

6. Prompt engineering: writing effective prompts

You've learned the commands, the shortcuts, and Plan Mode. You know what you can ask of Claude Code. One piece is missing: learning how to ask. This is prompt engineering — the discipline that separates those who get what they want on the first try from those who retry three times and complain about the results. It's also the technical heart of the vibe coding mentioned in [Preface](#): writing precise instructions so the model generates code aligned with your intentions, instead of writing the code by hand line by line.

There's a word that circulates among developers to describe the "naïve" way of working with AI: **Hope Coding**. You launch a generic request and hope the model guesses what you wanted. It works occasionally, often fails, and in worst cases produces code that looks right but isn't. The opposite path is to treat AI as an **extremely literal senior collaborator**: you tell it exactly what you need, in what context, with what constraints, in what format you want the answer. There's no magic, no "secret prompt": there's only a method.

A note of honesty before getting into the matter: prompt engineering in 2026 isn't what it was in 2023. The most "magical" techniques (act-as-an-expert, incantatory formulas, dramatic spellings) have deflated as models have improved. The discussion has shifted to two axes that still hold today: the **structure** of the prompt (context, task, constraints, output) and the **context** you load before asking. The real frontier of 2026 is context engineering: not how you ask, but what information you make available to the model before asking — a topic that in CLI translates into `CLAUDE.md`, Auto Memory, files read by subagents, and which we explore in chapters 7 and 8.

Evolution disclaimer. The techniques that follow reflect the state of the art at April 2026 (Claude 4.x and equivalent models). Prompt engineering changes fast: for the up-to-date reference always consult the [official Anthropic docs](#) and the [Prompting Guide](#).

6.1 What prompt engineering is and why it matters in CLI

Prompt engineering is the art (and partly the discipline) of formulating requests that produce predictable and useful output from an LLM. Three practical observations to frame it in the context of Claude Code:

- **It's not writing long descriptions.** The more verbose a prompt is, the more the model risks losing the point. Information density matters more than length.
- **It's not "making the prompt look intelligent".** A brilliant-to-read but vague prompt produces mediocre output. A prompt that looks like a public servant's checklist, but is specific, produces excellent output.
- **In CLI the prompt is an action, not just text.** In web chat, the prompt produces only text as a response; in Claude Code the prompt orchestrates **tools**: reads files, runs commands, modifies code. An ambiguous formulation doesn't just translate into a wrong answer: it translates into **wrong actions** on your filesystem. The margin for error is higher.

6.2 Anatomy of a well-made prompt

A well-made prompt contains four fundamental ingredients, plus an optional one we discuss right after:

1. **Context** — the background: what the project is, who the audience is (if relevant), what the technology stack is, what domain constraints apply.
2. **Task** — the requested action. Golden rule: **one task at a time**. Mixing different requests in a single prompt produces hybrid and confused output.
3. **Constraints** — what the model must do and must not do: length, tone, code standards, prohibitions ("don't use jQuery", "no external libraries", "max 100 lines").
4. **Output format** — how you want to receive the answer: Markdown table, JSON with a specific schema, "code only without explanations", bullet list, etc.
5. **(Optional) Role** — "act as a senior backend engineer". It's the fifth ingredient, deliberately listed last: in 2026 its weight is significantly reduced. I dig into why right away.

Before/after example, to fix the idea. **Vague** version:

```
Write me a function to validate an email
```

Structured version:

Context: Node.js + TypeScript project, server-side validation of user registration form. Compatibility constraint with Node 22 LTS, no external dependencies.

Task: implement a function that validates an email string.

Constraints:

- Pure TypeScript, no libraries
- Returns a discriminated Result type `{ ok: true, email: string } | { ok: false, reason: 'invalid_format' | 'invalid_domain' | 'too_long' }`
- Maximum accepted length: 254 characters (RFC 5321)
- Basic format validation + TLD presence check
- Vitest unit tests in a second code block

Output format: two distinct TypeScript code blocks (implementation + tests), no explanation between them.

The two requests are the same task, but they produce output of completely different quality. Not because the model is more “intelligent” in the second case: because it has fewer degrees of freedom to make mistakes on.

6.3 From roles to structural constraints (the 2026 revolution)

For years the first piece of advice on prompts was: **start with the role**. “Act as a senior security engineer”, “You are a cloud architecture expert”, and so on. It worked: older models were sensitive to the role “frame” and modulated the style and depth of the response.

On 2026 frontier models this lever has shrunk significantly. The Anthropic documentation for Claude 4.x indicates as the three primary levers of a good prompt: **explicit instructions, adequate context, curated examples** when needed. The role is no longer among the main levers. The technical reason is that modern models autonomously deduce the “competence” to call upon when context, task, and constraints are specific. Saying “act as a senior PHP engineer” adds little if in the context you’re already saying “WordPress plugin on PHP 8.1, PSR-4 namespace, code for production”.

That said, **the role isn’t dead**. It remains useful in specific scenarios. Here’s when it’s still worth using:

Situation	Does the role help?
Well-specified technical task on frontier model (Claude 4, equivalent GPTs)	No — redundant with context+task
Strong narrative voice (storytelling, copy with distinctive tone)	Yes — guides style
Domains with ambiguous regulatory references (legal, fiscal, healthcare)	Yes — orients the interpretive frame
Smaller or free models	Yes — they're more sensitive to roles
Persistent system prompt (e.g., Claude Projects, custom subagents)	Yes — defines stable session identity

For daily technical tasks on frontier models: focus on **structural constraints**, not on the role.

XML-LIKE DELIMITERS: THE MODERN PATTERN

On 2026 models (Claude in particular), the use of **XML-like delimiters** emerges as the preferred pattern for visually separating prompt sections. It reduces ambiguity, especially in long conversations where the model has to recognize which part of the message is instruction and which is, for example, code to analyze.


```
<context>
WordPress project, custom plugin, PHP 8.1.
Base theme: Astra. Block editor: Gutenberg.
</context>

<task>
Generate custom CSS for the theme's primary buttons (class
.wp-block-button__link) with a modern hover effect: slight scale,
fluid transition, subtle shadow.
</task>

<constraints>
- No !important
- Responsive (mobile-first)
- Use CSS variables for colors
- Comments in English
</constraints>

<output_format>
CSS code only, ready for Appearance → Customize → Additional CSS.
No textual explanation.
</output_format>
```

The tags have no semantic meaning to the model (it's not real XML), but they act as **clear separators**. The model recognizes them as delimiters and treats each section as a coherent block. For complex prompts, it's one of the most reliable patterns.

6.4 The fundamental techniques

There's no universal technique. Each technique responds to a certain type of problem. The five that follow cover most practical use cases for those working with code.

6.4.1 CHAIN OF THOUGHT (COT) — STEP-BY-STEP REASONING

The idea: explicitly ask the model to **reason in phases before answering**, instead of producing the conclusion directly. The magic formulas are simple: "think step by step", "reason in steps before proposing the solution", "first analyze, then propose".

It works because it forces the explication of logical steps. The model doesn't "jump to the answer" on intuition, but breaks the problem down into sub-problems and addresses them one at a time.

Commented example: diagnosing site slowness.

A WordPress + WooCommerce site has become slow in recent weeks (TTFB > 3s on product pages).

Before proposing solutions, reason in phases:

1. List the most probable causes of TTFB worsening on WP/WooCommerce in production.
2. For each cause, indicate how to verify it (free tool, SQL query, log to check).
3. Sort the causes by probability + ease of verification.
4. ONLY AFTER completing the three steps above, propose an intervention plan in 5 ordered steps.

Don't jump to generic advice like "use a caching plugin". I want the diagnostic analysis first.

What makes this prompt work: the last line ("don't jump to generic advice") blocks the most common response pattern. The 4-phase numbering forces the model not to take shortcuts.

On Claude 4.x there's also **extended thinking** as a product capability: the model "thinks" before answering, showing the reasoning in a separate block. On Claude Code, activatable with **Alt+T** (see chapter 4.7). It's the "native" version of CoT, to be preferred when available.

Worth it when — diagnosis, debugging, architectural decisions, any multi-phase task where the risk is the "pre-packaged answer" that skips intermediate steps. On Claude 4.x it's almost always better to activate native **extended thinking** (**Alt+T**) instead of reconstructing CoT in words.

It's filler when — well-defined linear tasks ("rename this function", "write a test for X"). On 2026 models, step-by-step reasoning is already implicit: adding "think step by step" to a simple request doubles the tokens without improving the output.

6.4.2 FEW-SHOT PROMPTING — TEACHING BY EXAMPLES

The idea: instead of describing how you want the output, **show it with two or more examples**. The model recognizes the pattern, applies it to the new input.

It's the most effective technique for **voice consistency** (maintaining a uniform style across recurring content) and for **reproducing structured formats** that are hard to describe in words.

Commented example: generating FAQ in colloquial style.

I have to write FAQs for a site of products for early infancy.
Tone: confidential, never patronizing, some emojis but sparingly. Synthetic answer, max 3 lines.

I'll give you two examples of the tone I want:

? When can I start weaning?

Italian pediatric guidelines talk about 6 completed months, but every baby has their own timing 🍼. Always talk to your pediatrician before starting.

? Can I wash baby bottles in the dishwasher?

Yes, if the temperature exceeds 60 °C. But remember to sterilize them separately once a week – the dishwasher isn't enough to eliminate everything.

Now write me 5 FAQs in the same tone on these topics:

1. Pacifier sterilization
2. Food allergies in the first 12 months
3. When to switch from breast milk to follow-on milk
4. Safe sleeping position
5. Mandatory vaccinations 2026

What makes this prompt work: the two examples are **complete and canonical**. They show format (emoji + question + 2-3 lines), tone (confidential but responsible), and a specific implicit rule (always defer to the pediatrician when in doubt).

Worth it when — voice consistency, microcopy, product cards, FAQs, classifications with your labels, structured formats hard to describe in words. 2-3 well-chosen canonical examples are enough.

It's filler when — technical tasks where the spec is crystal clear in words. Inserting 2 examples of "how to write a JUnit test" wastes context: a sentence on the tech stack is enough.

Anti-pattern: 8-10 examples "for safety" — the model over-specializes and loses generality.

6.4.3 PANEL OF EXPERTS (ROUND TABLE)

This is the technique that, in my opinion, is most worth learning in depth. **It doesn't only serve to "get an answer"**: it serves to learn, explore, put your own ideas into question. It's particularly valuable when you have to make a decision and don't want to settle for a single answer, but want to hear different perspectives — especially those that might not come to your mind.

The idea: you simulate a **discussion among virtual specialists**, each with their own viewpoint. You ask the model to interpret each one with their own perspective and to explicitly point out conflicts. The value isn't in the final synthesis, but in the **explication of trade-offs** that each decision implies.

The most powerful use cases:

- **Choosing a stack** for a new project (e.g. "PHP+MySQL or Node+PostgreSQL?" depends on who's looking at it)
- **Evaluating an architecture** before starting to code it
- **Stress-testing an idea** that seems good — you want to know where it breaks before discovering it in production
- **Asking for an opinion** before a decision with consequences (big refactor, DB migration, choice of library that will enter many files)
- **Understanding a topic** you know little about, listening to different voices instead of a single, potentially partial answer

Prompt-template for software development (canonical, reusable):

```
You're in a session dedicated exclusively to analyzing,
suggesting, and possibly creating code snippets.
Behave as if we're having a debate at a
round table with the following virtual experts:
```

- Full-stack computer engineer
- Programmer expert in PHP
- Programmer expert in JavaScript, Node and React
- Database Administrator and Data Engineer
- UX Designer
- Project Manager

```
For every question I want a response from each expert with
their own opinion. If there are discordant observations,
point them out to me. Every code proposal must be explained
and commented step by step.
```

What makes this prompt work:

- **Variety of angles:** full-stack sees the whole, PHP/JS see the technical stack, DBA sees persistence, UX sees the end user, PM sees timing and priorities. Covering angles you would individually miss is the point.
- **Explicit request for conflicts** ("if there are discordant observations, point them out to me"). Without this line the model tends to synthesize toward fictitious consensus. With the line, dissent becomes explicit and is the most useful part.

- **Step-by-step explanation of code** requires that every proposal be argued, not just presented. It helps unmask proposals that “look right” but aren’t.

Applied example: you have a small internal application (task tracker for an 8-person team). You have to decide whether to write it as a custom PHP+MySQL app, or as a Next.js + PostgreSQL app, or use a no-code tool. Launched the prompt template above and then:

Question: for an internal task tracker (team of 8 people, ~500 tasks/month, dashboard with filters and a REST API for integration with Slack), evaluate three options:

1. Custom PHP+MySQL app
2. Custom Next.js + PostgreSQL app
3. No-code tool (Airtable, Notion, ClickUp)

I want pros/cons from each of you, and a final recommendation with the main trade-offs explicitly stated.

What you’ll typically get: the full-stack will look at the “total 3-year maintenance cost”, PHP/JS will compare on development experience, the DBA will raise the migrations point, UX will say no-code already has an excellent UI you’ll never recreate well, the PM will say “with 8 people it’s not worth writing anything, get Notion”. The value is in having heard **also** the PM’s voice, which alone you would never have inserted into the reasoning.

Worth it when — architectural or stack decision with real trade-offs across dimensions you can’t weigh alone (full-stack vs. UX vs. PM vs. DBA). The value is in the **explication of dissent**, not the final synthesis.

It’s filler when — questions with a single technical answer (“what’s the complexity of a quick-sort?”). Staging a debate on settled questions doesn’t add perspective, it lengthens the answer.

6.4.4 CONTEXT ENGINEERING — THE NEW FRONTIER

“Pure” prompt engineering has a limit: no matter how well you formulate the request, the model only knows what you’ve given it. If you’re asking it to review an architecture without showing it the code, or to write a product description without showing it brand guidelines and existing cards, you’re asking the impossible.

Context engineering is the discipline of what you make available to the model before the question: relevant files, pre-existing examples, documentation, screenshots. The more the context is **clean, structured, and relevant**, the less you have to rely on “magic” prompts — and the more the structured prompts we discuss here pay off.

For **web chat** this means uploading PDFs, attaching images, using Projects to persist briefs and reference files.

For **Claude Code CLI** context engineering translates into three mechanisms you've already seen or will see:

- **CLAUDE.md** (see section 7) — persistent project context: stack, conventions, rules. Loaded at every session.
- **Auto Memory** (see section 7) — dynamic learnings written by the model itself, persistent across sessions.
- **Subagent as delegation strategy** (see section 12) — when the context to load is large, you delegate to a subagent that digests it and returns only the summary.

An important thing: **more context isn't automatically better**. Chroma's research on context rot shows that beyond certain thresholds the model degrades. The rule is "**better little and well-ordered than a lot and chaotic**". It's the same principle that governs section 8 of this guide on context management: treat it as a scarce resource, not as a dump.

Worth it when — always relevant if the model has to produce output coherent with material it doesn't know: real codebase, brand guidelines, DB schemas, past decisions. It's the highest-leverage discipline of 2026 and is almost always more effective than an elaborate prompt.

It's filler when — loading everything-everything: context rot degrades performance beyond certain thresholds. It's not filler in the sense of verbosity, but it becomes **context-noise**. The rule is "little and orderly": files pertinent to the task, not the whole `vendor/` or the whole email archive.

6.4.5 META-PROMPTING (THE PROMPT FOR THE PROMPT)

The idea is almost counterintuitive: **ask the model to write the prompt you should give it**. It's useful when a task is new or complex, and you don't know where to start.

Operational pattern:

In the role of Expert Prompt Engineer, you have to help me build an effective prompt for another session.

Goal of the future session: [describe the task in a rough way]

Proceed like this:

1. Analyze my request. Identify ambiguities and missing information.
2. Ask me 3-5 clarifying questions. Wait for my answers.
3. After my answers, write me the final complete prompt, structured with context/task/constraints/output format.

Start with the questions.

What makes this pattern work: the model doesn't give you the prompt right away (impossible, info is missing), but **forces the explication of ambiguities** that you alone wouldn't have noticed. The questions it asks are often those that, if not asked of you, would have led to a wrong result.

Worth it when — new or vague task where you don't know what you're asking: complex prompts to formalize for reuse (e.g., to be promoted in `CLAUDE.md` or in custom slash commands), unknown domain, marketing brief to translate into a technical spec. The 3-5 clarifying questions are worth the round.

It's filler when — task you already know how to formulate well. Asking the model to "write you the prompt" on a trivial refactor is a waltz to arrive at a formulation you would have written in 30 seconds.

SUMMARY TABLE: WHICH TECHNIQUE FOR WHICH PROBLEM

Technique	Problem it solves	Hint “this is the right one”
4+1 Anatomy	Vague, generic, unusable output	You just need to be more specific
XML-like delimiters	Long prompt where the model confuses sections	You have 3+ semantic blocks in the prompt
Chain of Thought	Answers that skip intermediate steps	Complex, multi-phase decision
Few-Shot	Output that doesn’t respect a precise style	You have 2+ examples of the desired pattern
Panel of Experts	Important decision with unclear trade-offs	You want to hear different angles, not a synthesis
Context Engineering	The model doesn’t know your context	You have reference material to load
Meta-prompting	You don’t know well where to start	New, vague task to formalize

In summary: in 2026 the truly high-yield levers are **explicit instructions**, **curated context**, and **examples when the format requires them**. The most “performative” techniques (verbal CoT to excess, panel on closed questions, meta-prompting on trivial tasks) are leftovers from old models and today often worsen the prompt’s signature — more input tokens, more risk of confusing the model, no advantage on the output.

6.5 Claude Code specifics compared to chat

Prompt engineering was born from chat and evolved there. In Claude Code CLI there are three substantial differences to keep in mind:

- **Tool use**: the prompt doesn’t just describe the output — it can activate actions. “Find all functions that use `mysql_query`” in chat produces a suggestion; in CLI it produces an actual reading of all files and a real list, because Claude executes Grep. The prompt should be calibrated knowing that every request can translate into actions on your filesystem.
- **Plan Mode** (see section 5) is a variant of prompt engineering applied: it explicitly separates the planning phase (read-only) from execution. For non-trivial tasks it’s the safest way to formulate risky requests.

- **CLAUDE.md** and **custom commands**: prompts that work well shouldn't be written every time. You **promote** them to permanent instructions in **CLAUDE.md** (see section 7) or to custom slash commands in **.claude/commands/** (see section 15.6). See also 6.8 below.

6.6 Before/after examples

Three practical cases that show the difference between a vague prompt and a structured one. All on realistic development tasks.

Case 1 — Refactoring

Before:

```
Refactor this function to make it more readable
```

After:

```
Context: TypeScript code from a Node app, function that handles
auth login. Project style: PSR-style but for TS, max 60 lines
per function, no nested ternary, early returns preferred.
```

```
Task: refactor the `authenticateUser` function below.
```

```
Constraints:
```

- Maintain exactly the same public signature and the same behavior (tests must continue to pass).
- Break the function into 2-3 private helper functions if needed.
- Replace nested if with early returns.
- No new external libraries.

```
Output format: 1) TypeScript block with new code,
2) brief bullet summary of what you changed and why.
```

Case 2 — Test generation

Before:

```
Write the tests for this function
```

After:

Context: Vitest, TypeScript email validation module (see function below).

Task: write a Vitest test suite for `validateEmail`.

Constraints:

- Coverage: all return paths of the function
- Edge case tests: empty email, too long (>254), invalid format (no @, multiple @s, no TLD), numeric TLD
- No snapshot tests
- Use `describe` to group by scenario, `it` for cases

Output format: TS code block only, no explanation.

Case 3 — Structured output for pipelines

Before:

Analyze this function and tell me if it has security problems

After:

Context: pre-commit code review, output will be parsed by a Python script to insert findings into a report.

Task: analyze the PHP function below for security problems (SQL injection, XSS, missing nonce, capability check, hardcoded secrets).

Constraints:

- Only real problems, no overly speculative "potential" ones
- For each finding: severity (critical/high/medium/low), line, explanation, suggested fix

Output format: JSON ONLY, no text before or after, schema:

```
{
  "findings": [
    {
      "severity": "critical" | "high" | "medium" | "low",
      "line": <number>,
      "type": "<sql_injection | xss | missing_nonce | ...>",
      "description": "<string>",
      "suggested_fix": "<string>"
    }
  ],
  "summary": {
    "critical": <number>,
    "high": <number>,
    "medium": <number>,
    "low": <number>
  }
}
```

The principle is always the same: reduce the degrees of freedom the model can make mistakes on.

6.7 Common anti-patterns

Mistakes you'll see (and make) recurrently:

- **Hope Coding** — "Write me a product description", "fix this bug", "refactor this function". No context, no constraints, no format. It's the founding antipattern: it produces random results.
- **Multiple tasks in one prompt** — "Refactor the function, write the tests, update the documentation, and commit". The model chooses what to do and what to skip, and the quality of each individual task collapses. One task at a time.
- **Relying on the role as a shortcut** — "Act as a senior engineer" doesn't replace a well-made brief. The role, where useful, completes the prompt; it doesn't replace it.

- **Context too long or chaotic** — loading 50 files “for safety”, pasting 200 lines of irrelevant logs, describing entire projects when a synopsis would suffice. See section 8 (Context management): the model degrades with excessive context.
- **Ambiguity in constraints** — “not too long”, “in adequate tone”. “Adequate” to what? “Long” compared to what? Quantified, not qualitative constraints: “max 100 lines”, “confidential tone like in the examples below”.
- **Not documenting prompts that work** — you rewrite the same complex prompt every time. Methodological mistake: see 6.9.

6.8 Promoting a prompt: when it goes in CLAUDE.md or in a custom command

Once a prompt works, you have three possible destinations:

- **Daily** — you rewrite it on the fly when needed. Fine for occasional tasks.
- **CLAUDE.md** (see section 7) — persistent project instructions. Loaded at every session, you don’t have to repeat them anymore. Perfect for rules that always apply in that project: code conventions, build commands, prohibitions.
- **Custom slash commands** in `.claude/commands/` (see section 15.6) — recurring workflows you want to recall with a single command. Perfect for things you do often but not in every prompt.

Promotion threshold, practical rule: **if you find you’ve rewritten the same instruction for the third time, it goes somewhere**. Project rules in `CLAUDE.md`, personal workflows in custom commands.

6.9 Prompt library: archiving and versioning

Prompts that work are **assets**, not single-use text. Treating them as such means archiving them with some minimum discipline.

Minimum pattern that works:

- A `prompts/` folder (or a single Markdown file, or Notion, or whatever you prefer) with a file per pattern: `code-review-php.md`, `refactor-typescript.md`, `panel-of-experts-software.md`, etc.
- For each prompt: a brief description of the use case, the actual prompt, any notes on known limits.
- **Versioning**: when you refine a prompt, keep the previous version with a suffix (`-v1.md`, `-v2.md`) and a brief changelog of what changed and why.
- **Feedback loop**: when a prompt fails, note the case that broke it. Often it’s an edge case that will help you refine it.

For Claude Code CLI, the most recurring prompts can be **promoted to custom slash commands** (see 15.6) — they effectively become part of your tool, invocable with a single `/`.

7. Persistent memory: CLAUDE.md and Auto Memory

Claude Code has **two persistent memory mechanisms** that coexist and complement each other: `CLAUDE.md` (a static contract written by you) and Auto Memory (dynamic learnings written by the model). Understanding how they work together is what separates a casual use from a professional one of the CLI.

The first part of this chapter (7.1-7.5) covers `CLAUDE.md`: what it is, how to generate it, what it contains, project examples, hierarchy on monorepos. The second (7.6-7.12) covers Auto Memory: the system introduced in v2.1.59 in which Claude autonomously notes what it learns from your corrections and reapplies it in subsequent sessions.

The `CLAUDE.md` file in the project root is the **contract** between you and Claude. It's automatically read at every session and provides the persistent context you would otherwise have to repeat every time. It's also the **natural destination for prompts that work and recur**: when you find yourself writing the same instruction in many different sessions, its place is here (see section 6 on prompt engineering, in particular the principle of prompt "promotion").

7.1 Generating CLAUDE.md with /init

On a new project where `CLAUDE.md` doesn't yet exist, the fastest way to start is the `/init` command. Launched from the project root, Claude analyzes the codebase (folder structure, `package.json`/`composer.json`/`requirements.txt`, configuration files, README, any tests) and generates a draft `CLAUDE.md` with detected stack, main commands, and inferred conventions.

```
# From the project root
claude
> /init
```

The output is a **good starting point**, not a definitive file. It should always be re-read and enriched by hand for two reasons:

- Claude can poorly infer conventions when existing code isn't uniform (e.g., half the project in camelCase, half in snake_case)
- The **tribal rules** not written in the code — "we don't use jQuery here", "every endpoint must have a Zod schema", "tests hit the real database, never mock" — Claude can't make them up: you have to add them.

Consider `/init` as scaffolding: it saves you 20 minutes of initial writing, then it's up to you.

7.2 What to put in CLAUDE.md

A good structure includes:

1. **Project description** — what it is, who it's for
2. **Technology stack** — languages, frameworks, versions
3. **Code conventions** — naming, patterns, comment style
4. **Main commands** — build, test, lint, deploy
5. **High-level architecture** — key folders, data flow
6. **What NOT to do** — anti-patterns, hard rules

7.3 Example 1: WordPress plugin

```
# WP Access Control Block

## Description
WordPress plugin that adds a Gutenberg block to control
content visibility based on user login state.

## Stack
- PHP 8.1+
- WordPress 6.0+
- JavaScript with @wordpress/scripts (JSX)
- SCSS with BEM methodology

## Conventions
- Code comments: **Italian**
- README and user documentation: **English**
- PHP naming: PSR-12, namespace `Mavida\WPAccessControl\`
- JS naming: camelCase, React components in PascalCase
- CSS: BEM strict (`.block__element--modifier`)
```

```
## Commands
- Build: `npm run build`
- Dev watch: `npm run start`
- PHP lint: `composer lint`
- JS lint: `npm run lint:js`
- PHP tests: `composer test`

## Structure
- `src/` — JS/SCSS sources (JSX, SCSS)
- `build/` — compiled output (don't touch manually)
- `includes/` — server-side PHP logic
- `block.json` — block manifest

## Hard rules
- DO NOT use jQuery in new components
- DO NOT commit files in `build/`
- Every PHP hook must have nonce verification
- Server-side render callbacks must be **escaped** with WP functions
```


7.4 Example 2: generic Node/TypeScript project

```
# API Analytics Service

## Description
REST microservice to collect and aggregate analytics events.

## Stack
- Node.js 22 LTS
- TypeScript 5.4 (strict mode)
- Fastify 4
- PostgreSQL 16 + Prisma ORM
- Vitest for tests

## Conventions
- All exported types must be in `src/types/`
- Zod schemas for input validation (never trust the client)
- JSDoc comments for all public functions
- No `any`, use `unknown` and narrow

## Commands
- Dev: `pnpm dev`
- Build: `pnpm build`
- Test: `pnpm test`
- Test coverage: `pnpm test:coverage`
- DB migrations: `pnpm prisma migrate dev`

## Architecture
- `src/routes/` – HTTP endpoints
- `src/services/` – business logic
- `src/repositories/` – data access
- `src/schemas/` – Zod validation

## Hard rules
- DO NOT directly import the Prisma client into `routes/` – go through `repositories/`
- Every endpoint must have a Zod schema for body/params/query
- Mandatory tests for every new service
```

7.5 Hierarchical CLAUDE.md

Claude Code doesn't read just one `CLAUDE.md`: it reads **more than one**, in order, from the general to the specific. More specific files **integrate** (and where needed override) more general ones. The typical hierarchy is:

1. `~/ .claude/CLAUDE.md` — **global user rules**

2. `<monorepo-root>/CLAUDE.md` — **monorepo rules**
3. `<project>/CLAUDE.md` — **single project rules**

GLOBAL USER RULES (`~/.CLAUDE/CLAUDE.MD`)

This is your personal file, valid for **all projects** you open from this machine. You'll find it in:

- **Linux/macOS:** `~/.claude/CLAUDE.md`
- **Windows:** `C:\Users\<your-user>\.claude\CLAUDE.md`

You put your **cross-cutting preferences** in there: comment language, response style, tools you always use, things never to do regardless of project. It's the right place for "I write comments in Italian" or "never propose me solutions with jQuery if a vanilla alternative exists".

WHAT IS A MONOREPO

A **monorepo** is a single Git repository that contains **multiple related projects** instead of having one for each. It's very common in modern development: a company that maintains multiple WordPress plugins together, or a team that keeps an API, a frontend, and a shared library in the same repo.

A typical structure:

```

my-monorepo/
├── CLAUDE.md                ← rules common to all sub-projects
├── plugins/
│   ├── access-control/
│   │   └── CLAUDE.md        ← rules specific to this plugin
│   └── analytics/
│       └── CLAUDE.md        ← rules specific to this plugin
└── shared/
    └── utils/

```

The monorepo-level `CLAUDE.md` captures what's **common**: naming conventions, base stack, orchestrated build commands. The single-project `CLAUDE.md`s cover **specifics**: rules that only apply to that sub-project.

DIFFERENCES FROM A SINGLE PROJECT

When you work on an isolated project (not in a monorepo), the hierarchy reduces to two levels: user rules + project rules. The project's `CLAUDE.md` contains everything that in a monorepo would be spread across two files. Nothing wrong with this: the hierarchy exists to avoid duplication, not to be mandatory.

PRACTICAL THREE-LEVEL EXAMPLE

Suppose you develop WordPress plugins and keep your work in a monorepo.

`~/.claude/CLAUDE.md` — personal preferences, valid everywhere:

Personal preferences

- Code comments: Italian
- Responses: concise, get to the point, no preambles
- When proposing PHP code, always follow PSR-12
- Don't propose solutions with jQuery if a vanilla JS alternative exists
- Before large structural changes, ask for confirmation

`<monorepo-root>/CLAUDE.md` — rules common to all plugins in the monorepo:

WordPress Plugins Monorepo

Common stack

- PHP 8.1+, WordPress 6.4+
- Build with @wordpress/scripts (npm)
- Tests with PHPUnit + WP Test Suite

Common conventions

- Root namespace: ``MyCompany\``
- All plugins are prefixed ``mc-`` (e.g., ``mc-access-control``)
- Translation files live in ``<plugin>/languages/``

Orchestrated commands

- Build all plugins: ``npm run build:all``
- Test all plugins: ``composer test:all``

Hard rules

- Every hook must verify nonce and capability
- No unescaped output: use ``esc_html``, ``esc_attr``, ``wp_kses_post``

`<monorepo-root>/plugins/access-control/CLAUDE.md` — rules only for this plugin:

```
# Plugin: Access Control

## Description
Gutenberg block to control content visibility
based on user login state.

## Specifics
- Block name: `mycompany/access-control`
- Server-side rendering via `render_callback` (no JS in frontend)
- Visibility rules are in `includes/Visibility/Rules.php`

## What NOT to touch
- Don't modify `block.json` without regenerating the asset manifest
- The `mc_access_control_can_view` filters are public API: no breaking
changes
```

When you open Claude Code inside `plugins/access-control/`, it sees all three levels combined: your preferences + monorepo conventions + plugin specifics. You write each rule **once, at the right level**, and reuse it automatically wherever it makes sense.

`CLAUDE.md` is the **static** half of Claude Code's persistent memory: you write it, it contains the rules. There's also a **dynamic** half that Claude feeds itself over time — that's Auto Memory, introduced in v2.1.59 and covered in the second part of this chapter.

Starting from **version 2.1.59** Claude Code introduced a second persistent memory mechanism, complementary to `CLAUDE.md`: **Auto Memory**. The difference is sharp:

- `CLAUDE.md` is **static** — you write it once, it's read at every session, contains rules and instructions.
- **Auto Memory** is **dynamic** — Claude writes it itself while working, accumulating what it has learned from your corrections and from past sessions.

These are two mechanisms that coexist: one is the **contract** you set, the other is the **learning diary** the model keeps for itself.

7.6 Auto Memory: what it is and what changes

The idea is simple: today if you correct Claude three times on the same thing — “don't use `var`”, here we use `let / const`” — by the fourth time, in a different session, it does it again. Without persistent memory, every session starts from zero. Auto Memory closes this loop: Claude autonomously notes the rule and reapplies it to subsequent sessions.

When Claude decides to save something, it does so based on an internal criterion: will the information be useful in future conversations? Recurring build commands, naming conventions you've corrected, project architectural patterns, typical errors to avoid. Anthropic's documentation doesn't describe the mechanism in detail; in practice it works better on recurring patterns than on individual occurrences.

Concrete example. You work on a PHP codebase that uses `snake_case` for function names. Claude initially proposes `camelCase` (JavaScript-style default). You correct once, twice, three times. Auto Memory notes the project convention. By the fourth session, before you even open your mouth, Claude already proposes `snake_case`. The explicit correction is no longer needed: the model has learned.

7.7 Requirements and enabling

- **Minimum version:** Claude Code 2.1.59. Verify with `claude --version` and update if necessary.
- **Default state: active.** You don't have to do anything to enable it.
- **Management command:** `/memory` shows all `CLAUDE.md`, `CLAUDE.local.md` files and rules loaded in the current session, allows toggling Auto Memory on/off, and provides a direct link to open the memories folder in the editor.
- **Persistent disabling:** in `settings.json` (user or local, **not** project, for security reasons):

```
{ "autoMemoryEnabled": false }
```

- **Disabling via env:** the environment variable `CLAUDE_CODE_DISABLE_AUTO_MEMORY=1` turns off the feature for a single session. Useful in CI/CD or one-off sessions.
- **Moving the folder:** the `autoMemoryDirectory` setting (always in user/local settings) lets you save memories to a custom path — for example a directory synchronized with another device via cloud, or an encrypted location.

7.8 Where memories live

Standard path:

- **Linux/macOS:** `~/.claude/projects/<project>/memory/`
- **Windows:** `C:\Users\<your-user>\.claude\projects\<project>\memory\`

The `<project>` segment is derived from the **Git repository** (remote URL). Outside a Git repo, the working directory root is used. Important consequence: **worktrees and subdirectories of the same repo share the same memory**. If you keep two copies of the repo to work in parallel on two branches, Claude learns from both sessions as if they were one.

7.9 Anatomy of the memory folder

Inside `memory/` there isn't a single file: there's an **index** plus **topic files**.

```
~/ .claude/projects/myproject/memory/  
├─ MEMORY.md           ← index, loaded at every session  
├─ debugging.md        ← topic file: recurring debug patterns  
├─ api-conventions.md  ← topic file: project API rules  
└─ build-commands.md   ← topic file: orchestrated build commands
```

- **MEMORY.md** is the index. It's loaded at **every session**, but only for the **first 200 lines** (about 25 KB). Anthropic recommends keeping it within this limit so as not to waste context.
- **Topic files** (`debugging.md`, `api-conventions.md`, etc.) are loaded **on-demand**: only when their content is relevant to the current task. They can be arbitrarily long without saturating the session context.

This is the **key technical difference compared to CLAUDE.md**, which is always read in full. Auto Memory is designed to **scale**: you can accumulate knowledge over time without paying for it in tokens at every session.

7.10 Auto Memory and subagents

Subagents (see section 12) can maintain their **own** Auto Memory, separate from the main session's. It's designed for specialized subagents that perform recurring tasks — for example, a code review agent that learns over time the preferred review style — without polluting the general work memory. Configuration is set at the subagent definition level.

7.11 When to disable it

Auto Memory doesn't leave the computer: everything lives in `~/ .claude/projects/...` on your local machine, and memories **are not sent to Anthropic for training**. Disabling is therefore not a privacy issue toward Anthropic, but one of **local control** in specific scenarios:

- **Codebase with sensitive data**: if fragments of real data (PII, secrets, clinical data) appear in chat, you prefer not to leave traces even in a local file that could be copied for backup, synced elsewhere, or read by someone with physical access to the machine.
- **Regulated projects**: in healthcare, finance, or GDPR contexts, some corporate policies prohibit any form of persistent memory outside controlled systems. Disabling Auto Memory is the safe choice to remain compliant.
- **Exploratory sessions**: you're experimenting with an approach you'll likely abandon. You don't want Claude to learn patterns from a transient solution and re-propose them in the future as if they were established rules.

When in doubt, start active and disable per session or project when needed.

7.12 CLAUDE.md vs Auto Memory: when to use what

Aspect	CLAUDE.md	Auto Memory
Who writes it	You	Claude
What it contains	Instructions and rules	Learnings and patterns emerged while working
Scope	Project, user, organization	Per repository (a single folder per repo)
Loaded in session	Always, in full	MEMORY.md always (max 200 lines), topic files on-demand
Typical content	Stack, conventions, commands, hard rules	Recurring build commands, style nuances, already-corrected errors

Recommended pattern: use `CLAUDE.md` for **hard rules** — the things you don't want to renegotiate every time — and leave Auto Memory for the **nuances** that emerge while working. If you notice that a learned memory is important and stable, **promote it** by manually transferring it to `CLAUDE.md`. From that moment it's a contractual rule, not just an observation that Claude could forget if it revises the file.

For those who already know the distinction, there's a complementary third way: `CLAUDE.local.md`, a manual Markdown file like `CLAUDE.md` but **gitignored** by default. You write it, it's specific to your local copy, isn't committed. It's useful for personal preferences of a single dev on a shared project (paths to local tools, your own shortcuts), without imposing them on the team.

8. Context management

The context window is Claude Code's most precious resource, and the one most users handle worst. Understanding what fills it, how it degrades when it's too full, and how to intervene before it becomes a problem, is the difference between a casual use and a professional one of the CLI.

8.1 What context is and why it matters

The **context** is the total amount of information the model sees in a single turn: Claude Code's system prompt, loaded `CLAUDE.md` files, conversation history, output of executed tools, contents of read files, definitions of active skills and MCP tools. All this is measured in **tokens** — units of text (one token corresponds roughly to 4 characters in English, slightly fewer in Italian).

Claude models in April 2026 have two window sizes:

- **200,000 tokens** — the default. Sufficient for the vast majority of tasks.
- **1,000,000 tokens** — available on specific models (see 8.7), designed for analysis of large codebases and very wide contexts.

Having a large window doesn't mean having constant performance over the entire length. Anthropic has explicitly documented a phenomenon called **context rot**: as the number of tokens grows, the model's accuracy and recall degrade. It's not a bug: it's an architectural consequence of the attention-based structure of transformers (peer relationships between tokens that explode quadratically with size). In practice, a model with context at 90% occupation **doesn't work like a fresh one at 10%**, even if the window is the same.

Concretely, this means context should be treated as a **scarce resource, not a dump**. Good practice isn't "load everything and see what happens": it's choosing carefully what enters, what leaves, and when to reset and restart.

8.2 What weighs in the context

A Claude Code session, before you even write the first prompt, has already loaded a certain number of tokens. Knowing what the categories are helps you understand where to intervene.

Category	Indicative weight	When it loads
System prompt	~4,200 tokens	Always, at startup
Environment info (cwd, OS, shell, git status)	~280 tokens	Always, at startup
MCP tools (deferred)	~120 tokens (only names, schemas on-demand)	Always, at startup
MEMORY.md (Auto Memory index)	up to ~6,500 tokens (max 200 lines / 25 KB)	Always, if Auto Memory active
CLAUDE.md (hierarchical)	depends on length	Always, in full , every session
Skill descriptions	~1% of the window (~2,000 tokens, default 8 KB chars)	Always for every active skill not disabled
Tool results (read files, Bash output, etc.)	very variable, often the biggest item	Dynamically, during work
Conversation (your prompts + responses)	variable	Dynamically

Values are indicative and derived from the [Explore the context window](#) page of the Claude Code documentation. They vary by session, model, configuration. To measure your real ones, use `/context`.

Three points worth fixing:

- CLAUDE.md files are loaded in full every session.** There’s no caching or indexing: all content enters the context. Anthropic recommends keeping each `CLAUDE.md` under **200 lines**: beyond that, both because the model has more difficulty following all the instructions, and because you consume tokens uselessly. If you find yourself with a 400-line `CLAUDE.md`, you’re probably putting things in there that should be in skills, project documentation, or `CLAUDE.local.md`.
- Installed skills occupy context even if not triggered.** It’s the most underestimated point. Every active skill contributes to the context with its own **description** (necessary to allow the model to decide whether to invoke it). The full skill content is loaded only when triggered, but the description is always there. A user who has 50 skills installed “just in case” starts with a significantly reduced

context budget compared to someone who has 10 well-chosen ones. **Install only the skills you actually use:** for the others, evaluate `disable-model-invocation: true` in the frontmatter (also removes the description from the context).

3. Tool results are the category that explodes. Reading 30 PHP files to look for a pattern can easily mean 50,000+ tokens of tool results. An `npm run build` with verbose output, even more. This is where most of the saturation plays out, and this is where subagents (see 8.6) make the biggest difference.

8.3 Signals of a saturated context

When the context approaches the limit, the model starts behaving in recognizable ways:

- **“Forgets” things said earlier** — responses that ignore a decision made two turns ago, or repeat explanations already given
- **Slower responses** — more tokens to process = more response time
- **“Stupid” errors** it didn’t make before — uses the wrong convention, refers to a non-existent file name, proposes code inconsistent with the rest of the project
- **Warnings in the status line** if you’ve configured it with a context indicator

What NOT to do when you see these signals. Continuing to “explain again” is the instinctive reaction, and it’s exactly the wrong one: every new explanation further inflates the context. Likewise, repeating the prompt more detailed doesn’t help — it makes things worse. The remedy is structural, not linguistic: use `/context`, read what weighs, act with `/compact` or `/clear`. See 8.4 and 8.5.

8.4 The `/context` command: reading and acting

`/context` is the most direct way to **see how much context you’re consuming** before the problem becomes visible from the model’s behavior. Launched at any time, it shows:

- the **percentage of context used** out of the total available
- the **breakdown by category**: system prompt, CLAUDE.md, skills, conversation, read files, tool output

Reading example:

`/context`

Context usage: 42% (84,000 / 200,000 tokens)

System prompt:	~4,200 tokens
Environment + MCP:	~400 tokens
CLAUDE.md (3 levels):	~2,800 tokens
Skills (8 active):	~3,100 tokens
MEMORY.md:	~680 tokens
Conversation:	~18,000 tokens
Tool results:	~54,800 tokens

STRATEGIC READING BY CATEGORY

The useful part is understanding **where the weight lies**, because it determines the right cure:

- **Tool results bloated (>50% of total)** → you've read many files or executed verbose commands. Solution: `/compact` summarizes while keeping decisions and key files, throws away the noise.
- **Conversation bloated** → you've mixed too many tasks in the same session, or you've dragged along old clarifications no longer relevant. Solution: consider `/clear` if you're at a task change.
- **System + CLAUDE.md + Skills high** → the session "tare" is too heavy. It's not solved with `/compact` or `/clear`: it's structural. You need to review how many skills you have active, how long your `CLAUDE.md` files are, if Auto Memory has accumulated too much in `MEMORY.md`.

INDICATIVE THRESHOLDS

Note on the numbers that follow. The percentages in this table are empirical: they reflect common practice and the author's experience in real sessions with medium-sized PHP/JS codebases. Anthropic doesn't publish official thresholds beyond the genericity of "context rot": take these values as a starting point, not as a prescriptive guideline. If you work on monorepos of 1M+ tokens, on synthetic tasks, or with a very particular pipeline, your useful thresholds will be different.

- **Below 50%** — relaxed, work normally.
- **50-75%** — start evaluating: close the current task and then `/compact` before opening a new one.
- **Above 75%** — it's time to intervene: `/compact` if you want to keep the thread, `/clear` if you're completely changing topic. Above 85% also consider switching to a 1M window model (see 8.7) if the task justifies it.

WHEN TO LAUNCH IT

`/context` is particularly useful in three moments:

- **Before a heavy task** — if you're about to ask for a Plan Mode on a large area or an audit of many files, knowing you're starting at 70% saves you from discovering halfway through that the model starts forgetting.
- **When you sense the model "slowing down"** — before concluding that "Claude is stupid today", check the context. Often it's just saturated.
- **In long, routine sessions** — even without signals, launching it every 30-40 minutes gives you proactive instead of reactive control.

8.5 Compression: `/compact` and `/clear`

These are the two main compression tools. Important difference:

- `/compact` compresses the conversation into a summary, keeping decisions and key context. **You continue to work with the thread of discourse**, but with much fewer tokens. Ideal mid-session when you've closed a phase and are opening another related one.

Optional syntax to give a focus to the summary:

```
/compact keeping the architectural decisions of the auth refactor  
and the pattern adopted for rate limiting
```

Without instructions, Claude decides what's important. With instructions, you tell it what to focus on and what it can throw away without regrets.

- `/clear` completely resets the context. **Fresh session**, but `CLAUDE.md`, skills, system prompt, and Auto Memory remain (they're system, not part of the conversation). Use when you change task completely and don't need anything from the previous one.

Example of typical flow:

```
[Morning start]
> Refactor the auth module to use JWT
[2 hours of work, /context says 78%]

/compact keeping decision of JWT with refresh token rotation

[Opening new phase]
> Now add unit tests on the new auth module

[Afternoon, completely different task]
/clear

> Update the API documentation to reflect the endpoint change
```

8.6 Subagents: the structural strategy

`/compact` and `/clear` are **reactive remedies**: you act when the context is already full. Subagents are the **preventive** tool: they work so the main agent never bloats for things it doesn't need.

Anthropic's documentation states it explicitly: subagents allow you to "preserve context by keeping exploration and implementation outside the main conversation". The mechanism, already seen in detail in [section 12](#), is simple:

- The main agent **delegates** a specific task to a subagent (built-in like `Explore`, or custom).
- The subagent runs in its **separate context window**, with dedicated tools and instructions.
- When done, it returns to the main agent **only a summary** of the result — not the 30 read files, not the 80 KB build output, only the distillate.

Three cases where the pattern pays off the most:

- **Pattern search across many files.** "Find all PHP functions that don't have nonce verification in the plugin" — done by the main agent means reading ~50 files and keeping them in context forever. Delegated to a subagent it means receiving a summary of 7 vulnerable functions with line and file, and that's it.
- **Massive audits.** Automatic code review on an entire plugin: three parallel subagents (security, performance, style) produce three independent summaries, the main agent aggregates them without seeing the detail of the individual files read.
- **Exploration of an unknown area.** Before starting a refactor, "understand how the notifications module is structured" given to an `Explore` subagent returns an architectural summary — 1,500 tokens — instead of 60,000 tokens of code files in the main context.

A good heuristic: **if you find you've read 30 files just to produce 200 tokens of output, it was a subagent.**

8.7 Models with 1M token window: when to switch

For most tasks, 200K tokens are more than sufficient. There are however scenarios where the jump to 1M qualitatively changes what can be done. The important thing: in API plans and enabled Pro/Max plans, **it doesn't cost more per token** — pricing is identical between 200K and 1M, you only pay for the tokens you use.

ACTIVATION SYNTAX

Three officially documented ways:

```
# During a session
/model sonnet[1m]
/model opus[1m]
/model claude-opus-4-7[1m]

# At launch from terminal
claude --model "sonnet[1m]"
claude --model "opus[1m]"

# Environment variable (default for every new session)
ANTHROPIC_DEFAULT_SONNET_MODEL=claude-sonnet-4-6[1m]
ANTHROPIC_DEFAULT_OPUS_MODEL=claude-opus-4-7[1m]
```

The `[1m]` suffix is the official syntax and appears in the `/model` picker when the model supports 1M.

MODELS WITH NATIVE 1M (APRIL 2026)

Model	Window	Notes
claude-opus-4-7	200K / 1M	Stable, default
claude-opus-4-6	200K / 1M	Stable
claude-sonnet-4-6	200K / 1M	Stable

Current state (verified April 30, 2026 on the official docs): the 1M window is available on Opus 4.7, Opus 4.6, and Sonnet 4.6. Sonnet 4.5 has 200K context, and Sonnet 4 is explicitly deprecated. The 1M beta that was available on Sonnet 4 / 4.5 has been retired: those with pipelines requiring it must migrate to Sonnet 4.6 or Opus 4.6/4.7. For the always-updated state, refer to the official models table.

WHEN IT MAKES SENSE TO SWITCH TO 1M

Concrete scenarios:

- **Audit of a plugin with over 80 PHP files** where you need to keep much of the code together to reason about cross-cutting patterns (e.g., all the points where the database is accessed, mapped in a single pass).
- **Structured comparison between two large branches** before a complex merge.
- **Cross-module migration** where decisions in one area depend on how 5-6 other areas of the codebase work.
- **Extensive legacy documentation** you want to keep complete in context to generate a new guide coherent with everything.

WHEN IT'S NOT NEEDED

- **Tasks focused on a few files** (3-10): 200K is plenty.
- **Punctual debugging** or fix of an isolated bug.
- **Incremental refactoring** where you work on one function at a time.
- **Exploratory sessions** where you change direction often: better a `/clear` from time to time than a monstrous window to manage.

Even with 1M, the context rot remains. It's not an excuse to "throw everything in there": it's a tool for scenarios where breadth is truly needed, to be used with the same discipline as 200K.

8.8 Practical rule and mindset

If you find you've done three different things in the same session, you probably should have used `/clear` twice.

Focused sessions produce better output and consume fewer tokens. The principle underneath everything: context is **a scarce resource, not a dump**. Everything that enters must earn its place.

CONTEXT MANAGEMENT TOOLS: WHERE THEY ARE IN THE GUIDE

Context management cuts across the entire guide, not just here. Here's where to dig deeper:

- **Right model for the task** — section 5 (`opusplan`, model selection)
- **Well-sized** `CLAUDE.md` — section 7 (under 200 lines)
- **Auto Memory not bloated** — section 7 (`MEMORY.md` as index, topic files on-demand)
- **Skills only those you need** — section 10
- **Subagent as structural strategy** — section 12
- **Hooks to reduce noise in session** — section 13 (e.g., filters on tool output)

8.9 Choosing the right architecture: decision table

Claude Code’s five extension mechanisms (`CLAUDE.md` , Auto Memory, Skill, Subagent, Hook) overlap in use cases and easily generate the question “which one do I use for what?”. Design decisions intertwine: does a code convention go in `CLAUDE.md` or in a skill? Does an automation go in a hook or in a custom slash command? Does an exploration of an unknown area get done by the main agent or delegated to a subagent? The following table is the unified map.

Tool	Typical use case	Context cost	When to use it	Limit
<code>CLAUDE.md</code>	Conventions, stack, hard rules that apply to every session of a project (language, framework, folder structure, build commands, anti-patterns to avoid).	High: loaded in full every session. Keep the file < 200 lines.	You have stable rules the model must always know before starting to work.	Doesn't adapt to cross-project preferences nor to dynamic learnings (for those, Auto Memory).
Auto Memory	Learnings that cross sessions and projects: user preferences, corrections the model must remember, stable architectural decisions.	Low: only <code>MEMORY.md</code> (index, max ~6.5K tokens) is loaded; topic files are on-demand.	You want Claude to learn over time from how you work, without you having to repeat the same instructions at every new session.	It's not a documentation repository: only concise rules/preferences. If it grows beyond 200 lines it should be pruned.
Skill	Codified and reusable playbook (procedure, analysis framework, writing pattern) invocable from any session that has it active.	Medium: ~1% window for description (always present), full content only if invoked.	A recurring procedure you'd like to distribute or standardize (<code>/security-review</code> , <code>/simplify</code> , a corporate code style skill).	The sum of descriptions of many installed skills erodes context: 10 targeted skills > 50 "just in case".

Tool	Typical use case	Context cost	When to use it	Limit
Subagent	Read-heavy task that would bloat the main context: audits, code-base exploration, pattern search across many files, comparative analyses.	Almost zero on main: the subagent runs in a separate window, returns only the summary. The real structural saving of tokens.	You're about to read 20+ files to produce a synthetic output, or you want to parallelize 3 independent audits.	Higher latency (it's another call), no shared state between subagent and main, summary may lose details.
Hook	Deterministic automation on lifecycle events (<code>PreToolUse</code> , <code>PostToolUse</code> , <code>User-PromptSubmit</code> , etc.): formatting, validation, log, security blocks.	Zero or negative: often a hook reduces context by filtering noisy output before it reaches the model.	You want something to happen always in response to an event, regardless of the model's decision (e.g., <code>prettier</code> after every <code>Edit</code> , shell block on dangerous patterns).	It's deterministic, not semantic: it doesn't "understand", it executes. It doesn't replace a subagent or skill when model judgment is needed.

How to read the decision table. The point isn't choosing "the best in absolute" but the one that sits in the right place of the chain. Three practical principles:

- `CLAUDE.md` **is the base**, not an aspiration: if the rule doesn't apply to every session of the project, it doesn't go there.
- **Skill and subagent** work together: often a skill orchestrates a subagent (e.g., `/security-review` delegates to an `Explore` subagent for the massive reading, then composes the report).
- **Hook is lateral:** it doesn't replace any of the other four, it integrates them at zero cost when automatic and predictable action is needed.

When you find yourself repeating the same instruction to three sessions in a row, you have a candidate for `CLAUDE.md` or Auto Memory. When you find yourself reading dozens of files to produce a summary, you have a candidate for a subagent. When you find yourself wanting to guarantee that something happens regardless, you have a candidate for a hook.

9. Security and permission management

Claude Code is an **autonomous** agent that runs commands on your system. Without the right precautions, it's a real risk vector.

9.1 The permissions system

By default, Claude asks for confirmation before performing any modification operation (file write, shell commands, MCP calls that modify state). Read operations are auto-approved:

- `Read`, `Glob`, `Grep`, `WebSearch`, `LSP` → no confirmation
- `Edit`, `Write`, `Bash`, write MCP → confirmation required

9.2 Configuring permissions in `settings.json`

You can define granular rules in the project's `.claude/settings.json` file:

```
{
  "permissions": {
    "allow": [
      "Bash(npm run test:*)",
      "Bash(npm run lint:*)",
      "Bash(git status)",
      "Bash(git diff)",
      "Read(**)"
    ],
    "deny": [
      "Read(.env*)",
      "Read(**/secrets/**)",
      "Read(**/.aws/credentials)",
      "Bash(rm -rf *)",
      "Bash(sudo *)",
      "Bash(curl * | bash)",
      "Bash(wget * | sh)"
    ]
  }
}
```

Step-by-step explanation:

1. `allow` — operations Claude can execute **without asking for confirmation**
2. `deny` — operations **physically blocked**, even if Claude tries nothing happens
3. Patterns use glob (`**` for any path, `*` for single segment)
4. `deny` has **precedence** over `allow`

VALIDATING THE FILE WITH THE OFFICIAL JSON SCHEMA

Anthropic publishes a JSON schema for `settings.json` hosted on SchemaStore. By adding the `$schema` line as the first key, JSON Schema-aware editors (VS Code, Cursor, JetBrains, Vim with coc-json, etc.) give you autocomplete of allowed properties, inline value validation, and tooltips with the description of each field:

```
{
  "$schema": "https://json.schemastore.org/claude-code-settings.json",
  "permissions": {
    "allow": ["Bash(npm run test:*)"],
    "deny": ["Read(.env*)"]
  }
}
```

The official docs (code.claude.com/docs/en/settings) warn that the schema is updated periodically: a validation warning on a property recently introduced in a recent release doesn't necessarily mean the configuration is invalid. The same schema also covers other sections of `settings.json`, including `hooks` (see section 13 — Hooks), `env`, `model`, `availableModels`.

9.3 Protecting secrets

Despite `.claudeignore`, there are scenarios in which Claude could read sensitive files (prompt injection, configuration errors). **Always** use `permissions.deny` for `.env` files, credentials, private keys.

9.4 Dangerous modes

`--dangerously-skip-permissions` skips all confirmations. It's useful for: - Autonomous execution in sandbox/Docker environments - Long tasks where you don't want to be interrupted every 30 seconds

The name is explicit: **it's not a flag to use lightly**. Guidelines:

- **Never** on machines containing production credentials
- **Never** with access to sensitive corporate repositories
- **Only** in isolated containers or VMs dedicated to the purpose

For lifecycle-level block automation (e.g., preventing `rm -rf` on protected paths even within `--dangerously-skip-permissions`), Hooks offer an additional programmatic layer: see section 13.

9.5 Prompt injection

An attacker could insert malicious instructions in:

- Code comments that Claude reads
- README files downloaded from dependencies
- Responses from untrusted MCP services
- Manipulated file names

Practical defenses:

1. Always use Plan Mode for tasks on third-party code
2. Always review the plan before approving it
3. Don't run Claude Code with administrator privileges
4. Isolate external projects in separate directories with restrictive `settings.json`

10. Skills: the extension mechanism

Skills are specialized “playbooks” that Claude can consult automatically when it detects that a certain type of task is in play. **Important:** unlike slash commands, **Skills aren’t invoked by a command.** They activate automatically based on their `description`.

10.1 How a Skill works

A Skill is a folder with a `SKILL.md` file structured like this:

```
---
name: wordpress-block-builder
description: "Use this skill when building or modifying WordPress
Gutenberg blocks. Triggers on: block.json, @wordpress/scripts,
JSX files in WordPress plugins, Edit/Save components."
---

# WordPress Block Builder

## Project conventions
- Always use @wordpress/scripts for the build
- Every block must have block.json, edit.js, save.js, style.scss
- Attributes must be typed in block.json

## Recommended patterns
[...]
```

The `description` field determines **when** Claude will use the Skill. It's the most important piece: write it thinking of the concrete triggers of your workflow.

A well-made skill leverages **progressive disclosure**: the `SKILL.md` file contains only what's essential to activate it; heavier resources (auxiliary Python scripts, configuration JSON, reference documents) live in adjacent folders and are loaded only when Claude decides they're needed. Result: the skill doesn't bloat the context until the moment of actual use.

10.2 Bundled native skills — deep dive

Anthropic distributes a set of official skills in the public repository `github.com/anthropics/skills`. Some are bundled with Claude Code, others must be installed separately (see section 10.4). Below are the eight most relevant native skills, each with reference metadata, an operational description, and a concrete usage example. All have the same author, the same monorepo repository, and Apache 2.0 license; the "Invocation" field indicates the typical triggers Claude reacts to by activating the skill.

10.2.1 PDF

Key	Value
Author / repo	Anthropic — <code>github.com/anthropics/skills/pdf</code>
License	Apache 2.0
Invocation	<code>.pdf</code> attachments, requests for extraction/fill/merge on PDF files

Extended description. Creation, text extraction, form filling, and merge/split of PDF files. The skill handles field positioning and typing when the user indicates a template and the data to insert. It activates automatically when PDF attachments appear in conversation or when the user requests operations on this format.

Usage example.

```
> "Open invoice-2026-04.pdf and extract the rows of the
  articles table in CSV format (code, description, qty, amount)"
```


10.2.2 DOCX

Key	Value
Author / repo	Anthropic — github.com/anthropics/skills/docx
License	Apache 2.0
Invocation	<code>.docx</code> files, requests to generate/modify Word documents from templates

Extended description. Creation and modification of Word documents preserving styles, headings, tables, and images. Typical for periodic reports produced from corporate templates, maintaining formatting and existing layout while replacing dynamic content.

Usage example.

```
> "Open template-monthly-report.docx and fill it with the data
  in stats-april.csv. Replace only the {{...}} placeholders,
  leave the rest of the formatting unchanged."
```

10.2.3 PPTX

Key	Value
Author / repo	Anthropic — github.com/anthropics/skills/pptx
License	Apache 2.0
Invocation	<code>.pptx</code> files, requests to generate presentations or slide decks

Extended description. Generation and modification of PowerPoint presentations with layouts, images, tables, and shapes. Particularly useful for repetitive slide decks (status meetings, training sessions) generated from outlines in Markdown or JSON.

Usage example.

```
> "Generate an 8-slide presentation starting from the file
  training-claude-code.md. One slide for each h2 found,
  bullet list of contents below."
```

10.2.4 XLSX

Key	Value
Author / repo	Anthropic — github.com/anthropics/skills/xlsx
License	Apache 2.0
Invocation	<code>.xlsx</code> , <code>.csv</code> files, requests for spreadsheets / dashboards / pivot tables

Extended description. Excel spreadsheets with formulas, conditional formatting, pivot tables, and charts. Suited for quick dashboards, data consolidations, and reports that would require a lot of manual setup in a new sheet.

Usage example.

```
> "Starting from sales-q1.csv generate an xlsx with a 'Data' sheet (raw), a 'Pivot' sheet (pivot by client x month), a 'Top10' sheet (top 10 articles with bar chart)."
```

10.2.5 FRONTEND-DESIGN

Key	Value
Author / repo	Anthropic — github.com/anthropics/skills/frontend-design
License	Apache 2.0
Invocation	Requests for web pages, React components, HTML layouts, landing pages

Extended description. Guidelines and patterns to produce distinctive UIs that avoid the “generic AI” style (gradients, excessive shadows, pastel palette). Pushes Claude toward more targeted typographic choices, palettes, and compositions, reducing the tendency to banal visual defaults.

Usage example.

```
> "Create a React landing page for a new B2B SaaS product. Professional tone, sober palette, hero + features + pricing + cta sections. No generic gradients and 'shiny' buttons."
```

10.2.6 WEBAPP-TESTING

Key	Value
Author / repo	Anthropic — <code>github.com/anthropics/skills/webapp-testing</code>
License	Apache 2.0
Invocation	Requests to “test” a web app, smoke tests, end-to-end verification

Extended description. End-to-end testing of web applications with Playwright in headless mode. Activates when Claude is asked to test a web app running locally: the skill orchestrates navigation, form filling, assertions, and comparison screenshots.

Usage example.

```
> "Start the app on localhost:3000, navigate to /signup, fill the form with valid data, verify that after submit the confirmation message appears. Then repeat with malformed email and verify the error."
```

10.2.7 SKILL-CREATOR

Key	Value
Author / repo	Anthropic — <code>github.com/anthropics/skills/skill-creator</code>
License	Apache 2.0
Invocation	Explicit requests to “create a skill” or “skill scaffolding”

Extended description. Meta-skill that interviews the user, gathers requirements, and produces a new skill (frontmatter, content, any auxiliary scripts). Recommended entry point for those who have never written a skill: Claude poses targeted questions (triggers, output format, examples) and generates the `.claude/skills/<name>/` folder ready to use.

Usage example.

```
> "I want to create a skill that helps me write ADRs
  (Architecture Decision Records) consistent with my
  team's template. Help me set it up with skill-creator."
```

10.2.8 MCP-BUILDER

Key	Value
Author / repo	Anthropic — <code>github.com/anthropics/skills/mcp-builder</code>
License	Apache 2.0
Invocation	Requests for MCP server scaffolding or exposure of an external API

Extended description. Meta-skill mirror of `skill-creator` but for MCP servers (see section 11). Generates the scaffold of an MCP server starting from a description of the API or service to expose, including SDK and basic tool structure.

Usage example.

```
> "Build an MCP server that exposes our internal ticketing
  API. Base endpoint: api.mavida.local/v2.
  Operations: list_tickets, get_ticket, create_ticket,
  add_comment. Auth via X-API-Key header."
```

Note — The subset of skills bundled with Claude Code may vary over time. The source of truth is always the official repository `github.com/anthropics/skills`, also cited in Appendix B.

10.3 Community skills: a curated selection

The ecosystem produces dozens of skills every month and quality is highly variable. Those that follow have been selected with criterion: official repository of a recognizable author or organization, clear open license, adequate documentation, recent maintenance activity. For each you'll find a card with essential data, an extended description, cases when it's worth installing, cases when it's not, an editorial assessment, and — where useful — a usage example.

Disclaimer — The ecosystem is in very rapid evolution. The skills selected here were active and well-maintained at the time of writing (May 2026); always verify the repo’s health (last commit date, open issues, license) before installing them.

10.3.1 SUPERPOWERS

Key	Value
Author / repo	Jesse Vincent — github.com/obra/super-powers
License	MIT
Stars	~173k
Invocation	Activates on non-trivial tasks: new features, complex bugs, structural refactorings

Extended description. Complete methodology for agents: brainstorming → planning → TDD → review → execution. Not a single skill but a framework of composable skills that imposes a structured, multi-platform workflow (Claude, Cursor, OpenAI Codex, Gemini). Forces Claude not to skip the analysis and planning phases even when the task seems solvable at a glance.

When to use it.

- You develop complex features where “diving into the code” leads to rewrites
- You want a rigorous workflow out-of-the-box without building conventions from scratch
- You work in teams where adherence to a shared methodology is important

When NOT to use it.

- You have to make small, localized changes (a three-line bug fix doesn’t deserve brainstorm + plan + TDD)
- You’re used to your own consolidated workflow and want flexibility

My judgment. For those coming from the “chat + copy-paste” pattern it’s a gear shift. The value isn’t in the individual skills but in the **discipline** it imposes: thinking before writing. It works better if adoption is team-wide — because a single developer can quickly find it heavy for daily short tasks and disable it. Worth trying for some non-trivial features before making it the basis of your workflow.

Usage example. After installation, ask Claude a non-trivial task (“Add 2FA authentication to the login module”). Superpowers forces the brainstorm → plan → test-first → implement → review sequence before touching code.

10.3.2 VERCEL LABS AGENT SKILLS

Key	Value
Author / repo	Vercel Labs — <code>github.com/vercel-labs/agent-skills</code>
License	MIT
Stars	~26k
Invocation	Activates on React/Next.js files, frontend code review requests, accessibility audits

Extended description. Package of 7 high-level frontend skills: `react-best-practices` (40+ performance rules), `web-design-guidelines` (100+ accessibility/UX rules), `react-native-guidelines`, `composition-patterns`, `react-view-transitions`, `vercel-deploy-claimable`. Codifies the rules published by Vercel Engineering in its best practices docs.

When to use it.

- You develop React/Next.js professionally and want code reviewed against known standards
- You need to do accessibility or performance audits on already-written components
- You work with a team that wants uniformity on the frontend stack

When NOT to use it.

- The stack is different from React/Next.js (the rules don't apply)
- You're prototyping fast and the rules slow down iteration

My judgment. More than a skill, it's a **codified style guide** that Vercel publishes for free. It works for those inside the React ecosystem: if you are, the investment is zero (one install command) and the return is high (more solid code at the first attempt). If you're not, it's useless to install it.

Usage example. During a review you ask "apply `vercel-labs/web-design-guidelines` to all files modified in the last commit". Claude executes a punctual audit (focus state, color contrast, form error handling, animations) producing an issue report with reference to the rules violated.

10.3.3 WORDPRESS AGENT SKILLS

Key	Value
Author / repo	WordPress (official organization) — <code>github.com/WordPress/agent-skills</code>
License	GPL-2.0-or-later
Stars	~1.4k
Invocation	Presence of <code>wp-config.php</code> , plugin/theme files, references to <code>block.json</code> or WP REST API

Extended description. Bundle of 14 WordPress domain skills built on the official docs, to mitigate the known problem of LLMs generating obsolete WordPress patterns (code from WP 4.x, deprecated ACF APIs, `query_posts()` instead of `WP_Query`, etc.). Includes a “router” skill (`wordpress-router`) that classifies the task and routes to the right specific skill.

When to use it.

- You develop WordPress plugins or themes, particularly with block editor and Interactivity API
- You want code adhering to official documentation and not to dated patterns
- You work on projects where WordPress security is important (capability, sanitize/escape, nonce)

When NOT to use it.

- You only work on WordPress content (posts, pages) and not on code
- The site is headless and WP only acts as CMS via REST/GraphQL (little benefit)

My judgment. For those who do WordPress daily, it’s the most useful skill in circulation. The problem of “LLMs generating obsolete WP patterns” is real and costly (I’ve seen AI-generated code use `query_posts()` or ACF v4 in 2026), and this skill eliminates it at the root. Three stars out of three for those working with WordPress; zero for those who don’t.

Usage example. In a WordPress project ask “Create a new Gutenberg block to insert a callout box with icon, title, and text”. `wordpress-router` recognizes the domain and activates `wp-block-development`, which imposes correct `block.json`, deprecation handling, and security patterns from the official docs.

10.3.4 TRAIL OF BITS SKILLS

Key	Value
Author / repo	Trail of Bits — github.com/trailofbits/skills
License	CC-BY-SA-4.0
Stars	~5k
Invocation	Security audit requests, code review for vulnerability assessment, malware analysis

Extended description. 40+ security skills for AI-assisted analysis, testing, and auditing: smart contract security, code auditing (CodeQL/Semgrep), malware analysis, reverse engineering, mobile security, verification techniques. Trail of Bits is one of the most recognizable security consulting companies in the open source space: these skills codify methodologies used in their professional audits.

When to use it.

- You need to review third-party code before integrating a library
- You’re a security engineer and want to orchestrate CodeQL/Semgrep in a guided way
- You work on projects with external exposure where vulnerability risk is high

When NOT to use it.

- You’re doing greenfield application development without critical exposure (overkill)
- You’re not familiar with the underlying tools (CodeQL, Semgrep): the skill orchestrates but debugging false positives requires expertise

My judgment. Serious skill for a serious audience. The CC-BY-SA-4.0 license imposes share-alike on derivatives: if you modify the skills and redistribute them you’re bound — verify compatibility with your corporate constraints. For those doing security it’s a notable boost; for the rest it’s a toolbox you probably won’t use.

Usage example. Before integrating a third-party PHP library ask “Run a security audit on the `vendor/foo-lib` folder using Trail of Bits skills. Look for SQLi, command injection, insecure deserialization patterns”. The skills orchestrate CodeQL/Semgrep and produce a report with file, line, and type of vulnerability.

10.3.5 CAVEMAN

Key	Value
Author / repo	Julius Brussee — <code>github.com/Julius-Brussee/caveman</code>
License	MIT
Stars	~24k
Invocation	<code>/caveman,</code> <code>/caveman lite,</code> <code>/caveman ultra</code> commands; natural triggers like “talk like caveman”

Extended description. Skill that forces Claude to respond in telegraphic style: no articles, no pleasantries, no hedging, no meta-commentary. Only technical substance. Three compression levels (`lite`, `full`, `ultra`) for increasing intensity. The skill is surgical: it compresses the discursive part of responses (filler, articles, courtesy phrases) but leaves intact code blocks, technical terms, cited error messages, and commit messages.

When to use it.

- Mechanical and repetitive coding (refactoring, debugging, linting)
- You’re an experienced user and don’t need the detailed “why”
- You’re orchestrating multiple agents or background tasks where verbose output is just noise
- You’re close to your plan limit and want to squeeze out more sessions

When NOT to use it.

- You’re learning a new framework: you need pedagogy, the “why” is the value
- You’re onboarding to an unknown codebase: you want complete explanations
- You’re doing architectural review: you want nuances, alternatives, trade-offs

My judgment. Caveman is an interesting case of a **viral claim worth reading carefully**. The README promises “75% token savings” and has become the project’s reference number, but it’s worth disaggregating the data to understand where the savings actually land on your bill:

Number	What it really measures
65%	Average compression on the model's output (range 22-87%, measured by the repo's benchmarks).
0.6-2.5%	Share of tokens that the output represents on the total of a typical Claude Code session. The bulk of consumption is input: <code>CLAUDE.md</code> , conversation history, read files, tool output.
1-2%	Net savings on the bill in a normal interactive session (small output, dominant input).
up to 15-25%	Savings in multi-agent parallel scenarios or headless batch pipelines, where output becomes a larger share of the total (more subagents producing final text, less repeated input).

In short: the viral “75%” applies only to a tiny slice of consumption in most human sessions, while the higher numbers you read around apply to particular scenarios — not typical, presented as typical. The product itself is honest (open source MIT, reproducible benchmarks in the repo), and the principle is what I take away: thinking of **AI verbosity as a measurable cost**. Once you start reasoning about it, you write lighter `CLAUDE.md` files, more targeted prompts, and efficiency grows even without plugins.

Usage example. Normal Claude response (69 tokens):

"The reason your React component is re-rendering is likely because you're creating a new object reference on each render cycle. When you pass an inline object as a prop, React's shallow comparison sees it as a different object every time, which triggers a re-render. I'd recommend using `useMemo` to memoize the object."

Caveman mode response (19 tokens):

"New object ref each render. Inline object prop = new ref = re-render. Wrap in `useMemo`."

Same diagnosis, same solution, no padding.

10.4 Installing and managing skills

Skills are installed in three ways depending on how they're distributed.

Plugin marketplace — when a skill is packaged as a plugin (typical case for the more structured community skills):

```
# Add the author's marketplace
claude plugin marketplace add owner/repo

# Install the specific skill
claude plugin install <skill>@<plugin>
```

`npm skills` **tool** — for standalone skills distributed via Git repo:

```
# Install a single skill from a multi-skill repo
npm skills add anthropics/skills --skill frontend-design

# Install the entire skill package from a repo
npm skills add vercel-labs/agent-skills
```

Manual copy — for custom or experimental skills. Two possible locations:

```
~/claude/skills/<name>/SKILL.md      # global (all projects)
.claude/skills/<name>/SKILL.md      # only this project
```

Resolution follows the same hierarchy as other Claude Code resources: global as default, local to project as override.

WHERE TO LOOK FOR NEW SKILLS

The ecosystem grows very quickly. The most reliable sources today:

- `github.com/anthropics/skills` — official Anthropic repository, first place to check
- `skills.sh` — public community directory, indexes third-party skills with basic metadata
- **Public plugin marketplaces** — authors of mature skills often publish a dedicated marketplace (Trail of Bits, Vercel Labs, JuliusBrussee for Caveman)
- **GitHub topics** `claude-skills` and `agent-skills` — useful for targeted exploration

Before installing anything from third parties, read section 10.6.

10.5 Creating a custom Skill

Nothing prevents you from writing your own skills: indeed, this is the point where Claude Code really becomes “yours”. Let’s see an end-to-end example of a `mavida-wordpress` skill that codifies your team’s conventions.

Step 1 — Scaffolding:

```
# Skill local to project (override) or global (~/.claude/skills/)
mkdir -p .claude/skills/mavida-wordpress
```

Step 2 — Writing SKILL.md in .claude/skills/mavida-wordpress/SKILL.md:

```
---
name: mavida-wordpress
description: "Use this skill when working on Mavida WordPress
projects. Triggers on: presence of wp-config.php, plugin
files in /wp-content/plugins/, theme files, references to
ACF or Block Editor."
---

# Mavida Conventions — WordPress

## Build and tooling
- Always use @wordpress/scripts for blocks (no custom webpack)
- Composer for PHP dependencies, never manual library upload
- Minimum versions: PHP 8.2, WordPress 6.4

## Mandatory patterns
- All new blocks must have block.json + edit.js + save.js + style.scss
- Strict sanitize/escape: `sanitize_text_field` on input, `esc_html`/`esc_attr` on output
- Nonce on all forms and AJAX actions
- Capability check with `current_user_can()` before privileged operations

## Patterns to avoid
- `query_posts()` (deprecated): use `WP_Query`
- Functions that bypass cache: `wp_cache_flush()` only in seed, never in runtime
- Loading JS via inline wp_head: always via `wp_enqueue_script`

## Common build commands
- `npm run build`: production build of blocks
- `composer test`: PHPUnit + PHPStan level 6
- `wp-env start`: local development environment
```

Step 3 — Verification: open Claude Code in a WordPress project and ask for a trivial modification (e.g., "Add a metabox to the Article post type"). If the skill is well-written, you'll see Claude automatically apply the conventions (capability check, sanitize, correct enqueue) without you having explicitly asked.

To check that the skill is loaded: `/skills list` (or `/help skills` depending on the version).

The same pattern applies to other team skills that will come in handy daily:

```

.claude/skills/
├── mavida-wordpress/           # WP conventions
│   └── SKILL.md
├── n8n-workflow/              # patterns for N8N workflows
│   └── SKILL.md
└── php-legacy-review/        # checklist refactoring PHP legacy
    └── SKILL.md

```

A well-written skill is a piece of “team memory” that survives session changes, developer changes, and — increasingly often — tooling changes. When a prompt you use often becomes a rule to systematically impose, it’s the signal that it’s worth promoting to a skill (see also the reflection on prompt “promotion” in section 6.8).

10.6 Security of third-party skills

A third-party skill is code that becomes part of Claude’s context and can influence its decisions. It doesn’t run autonomously — Claude still asks for confirmation before running commands — but it can **instruct Claude to call tools with real effects** (Bash, file system, WebFetch). Treat it as you would treat a library you include in `composer.json` or `package.json`: with the same diligence.

Checklist before installing a third-party skill:

- **Mandatory code review** — read `SKILL.md` in full and inspect all auxiliary scripts (Python, Bash, JS) present in the folder. Look for references to destructive commands, data exfiltration, undocumented calls to external domains.
- **Compatible license** — verify that the skill’s license is compatible with your project and any NDA clauses. CC-BY-SA (like Trail of Bits) imposes share-alike on derivatives; GPL-2.0 (like WordPress agent-skills) has known copyleft implications; MIT (Vercel Labs, Superpowers) is generally the most permissive.
- **Repo health** — check last commit date, number of stars, ratio of open/closed issues, presence of security advisories. A skill abandoned for a year is a risk for any codebase that uses it actively.
- **Permissions it requests** — some skills request access to powerful tools (Bash unrestricted, WebFetch on external domains, global writing). Compare them with your `settings.json` policy (see section 9) and reject those that ask for more than they justify.
- **Sandbox in dev** — test new skills on a throwaway project before installing them globally in `~/.claude/skills/`. If you need an additional layer of defense, a `PreToolUse` hook (see section 13) can block commands the skill tries to execute outside the allowed perimeter.

The mental schema is the same you would apply to any dependency: you don’t include code you haven’t read, you don’t trust an author just because they have many stars, and you don’t enable more than what’s needed. The difference is that here the “code” is an instruction in natural language that Claude will read and execute — and natural language is ambiguous by definition. Double attention.

11. MCP: integrating external services

The **Model Context Protocol (MCP)** is the way Claude Code talks to external systems — APIs, databases, SaaS services, file systems outside the working directory. While skills (chapter 10) extend what Claude knows how to do with local instructions and code, MCP extends which systems it can connect to. For a quick comparison between the extension mechanisms, see the map in 14.1.

11.1 What MCP is and why it exists

Model Context Protocol (MCP) is an open protocol, open-sourced by Anthropic in November 2024, that standardizes the way an AI application (the **host**) connects to external data sources and tools — files of a system, databases, service APIs, git repositories, calendars, ticket systems, anything reachable via code.

The idea is simple and arises from a concrete problem. Before MCP, every AI IDE (Claude Code, Cursor, Continue, ChatGPT desktop, Cline, dozens of others) had its own mechanism to connect to GitHub, Postgres, Slack, etc. For those developing an integration, this meant writing it N times — one for each client. For those using multiple tools, every client had a matrix of incompatible connectors: Cursor's GitHub integration didn't work in Claude Code, and vice versa.

MCP solves this problem with the same logic with which USB-C replaced dozens of proprietary connectors: it defines **a standard protocol** between client (AI host) and server (the integration). Whoever writes an integration writes it once, and it works wherever MCP is supported. Anthropic released the official SDKs (Python, TypeScript, Java, C#, Rust, Kotlin, Swift) and a dozen reference servers for the most common use cases (filesystem, HTTP fetch, GitHub, Postgres, SQLite, Puppeteer, Slack, Brave Search).


By May 2026 adoption is widespread: **Claude Code, Cursor, Windsurf, Cline, Continue, GitHub Copilot, and several other clients support MCP** natively. There are hundreds of community servers in public registries, and dedicated marketplaces (anthropic.com/mcp, glama.ai/mcp, smithery.ai). The protocol reached stable version 1.0 after the iterations of 2025.

For those writing in English: think of MCP as the **standard driver between Claude Code and the rest of the world**. If you want Claude to do something it can't do natively — read your CRM, post to WordPress, query an internal database — the modern answer is: write an MCP server (or find one that already does the case for you) and register it.

11.2 Protocol architecture

MCP is a **client-server** protocol based on **JSON-RPC 2.0**. Three main components:

- **Host** — the user's AI application (for us: Claude Code). It doesn't talk directly to MCP servers: it uses one or more clients.
- **Client** — a 1:1 connection to a single server. Claude Code creates a client for each configured MCP server. The client manages the connection (process startup, message exchange, lifecycle) and isolates the server from the rest of the host.
- **Server** — the process that exposes functionality. It can be written in any language for which an MCP SDK exists (Python, TypeScript, Rust, Java, C#, Swift, Kotlin are all officially supported). It communicates with the client through a standard transport.



Transports — two main modes:

- **stdio** (standard input/output) — the client launches the server as a sub-process and communicates via stdin/stdout. It's the most common transport for local servers (filesystem, database, our own processes). Simple, secure by default (stays in the user system), no network to manage.
- **HTTP+SSE** (Server-Sent Events) — servers reachable via HTTP. Used for hosted servers (cloud) or shared between different clients. Requires auth handling and latency considerations that don't

arise for stdio. The guide focuses on local stdio servers; we'll touch on remote servers in section 11.6.

Capability negotiation — at session startup, client and server exchange a handshake (`initialize`) in which they declare what they support: the server lists its tools, its resources, its prompts; the client lists its capabilities (e.g., support for sampling, logging). From that point on the conversation is an alternation of JSON-RPC requests.

The **three primitives** of an MCP server:

- **Tools** — functions the server exposes that Claude can invoke. Each has a name (`wp_create_post`), a textual description readable by the AI, and a JSON schema of arguments. When Claude decides to call it, it sends a `tools/call` request, the server executes, returns the result. It's the most used primitive.
- **Resources** — data addressable via URI (`wp://posts/123`, `file:///etc/hosts`). The server exposes them as a "library" Claude can read from. Resource ≠ Tools: reading a resource is a pure GET, with no side effects.
- **Prompts** — reusable prompt templates the server can provide to the user as "presets". Typically exposed as `/server-name:prompt-name` command in the host.

For our WordPress example we'll only use **tools** (post creation/update, category list). Tools are the most productive part of the protocol; resources and prompts are useful but less common in custom servers.

11.3 Configuring an existing MCP server

Configuration is declarative: you list the server in a JSON and Claude Code launches it automatically at session startup. Two scopes:

- **Project's** `.claude/settings.json` — the server is available only inside that project. Suited for project-specific integrations (a server to talk to the client's staging DB).
- **User's** `~/.claude/settings.json` — the server is available in every user session. Suited for global integrations (your MCP server for your own corporate CRM).

Example of configuring three servers simultaneously (one GitHub and two local):


```
{
  "mcpServers": {
    "github": {
      "command": "npx",
      "args": ["-y", "@modelcontextprotocol/server-github"],
      "env": {
        "GITHUB_PERSONAL_ACCESS_TOKEN": "${GITHUB_TOKEN}"
      }
    },
    "filesystem": {
      "command": "npx",
      "args": [
        "-y",
        "@modelcontextprotocol/server-fileSystem",
        "/Users/maurizio/projects"
      ]
    },
    "postgres": {
      "command": "uvx",
      "args": [
        "mcp-server-postgres",
        "postgres://localhost/staging_db"
      ]
    }
  }
}
```

Key points:

- **command + args** — executable and arguments. `npx -y` is the standard pattern for servers distributed on npm (downloads them on the fly). For Python you typically use `uvx` or `python -m`.
- **env** — environment variables passed to the server. The `${VAR}` interpolations are resolved by the shell that launched Claude Code: the correct way to pass secrets is to put them in `.env/.envrc` and export them before launching Claude, not hardcode them in the committed JSON.
- **No secrets in repo** — a committed `mcpServers` entry containing a token in clear is a security incident waiting to happen. Always `${VAR}` with secrets loaded externally.

Dedicated slash command and CLI. Once a server is configured, Claude Code manages it with these commands:

```
# Internal slash command (in interactive session)
/mcp                                # list of active servers and their tools

# External CLI
claude mcp list                     # same list from terminal
claude mcp add <name>              # adds a server (guided interview)
claude mcp remove <name>          # removes it
```

Debug. If a server doesn't start (startup error, missing dependency, unresolved env var), Claude Code prints the error at session startup and marks the server as "disconnected" in `/mcp`. To chase more subtle problems, launch the server manually from the terminal and exercise it via stdin with some `echo '{"jsonrpc":"2.0","method":"initialize",...}'` — the MCP documentation has precise examples.

11.4 Useful MCP servers: a curated selection

By May 2026 the MCP ecosystem is vast. A selection of servers worth getting familiar with:

- `@modelcontextprotocol/server-github` — complete management of repositories, issues, PRs, actions. It's the first MCP to install for those who develop on GitHub. Official, maintained by Anthropic.
- `@modelcontextprotocol/server-filesystem` — controlled access to specific filesystem directories. Useful for working on projects outside Claude Code's working directory (e.g., reading documentation in `~/Documents/specs`). The authorized paths are passed as arguments.
- `mcp-server-postgres` / `mcp-server-sqlite` — query, schema inspection, migration generation. Excellent for exploring staging databases without write permissions on prod.
- `@modelcontextprotocol/server-puppeteer` — headless browser automation: screenshots, scraping, click tests. Pairs very well with the `webapp-testing` skill (chapter 10).
- `mcp-server-slack` — sending messages and reading channels, useful for completion notifications of long tasks or automatic reports.
- `mcp-server-sentry` — access to error tracking data; can retrieve the stack trace of a recent exception and give it to Claude for the bug fix. Pairs with chapter 15.2 (Bug hunting with TDD).
- `mcp-server-linear` / `mcp-server-notion` — ticket system and knowledge base. Allow Claude to read the context of a task from Linear and produce the connected PR.

Community marketplaces have hundreds more: before writing an MCP from scratch, look for whether a suitable server already exists. The rule is: **existing MCP > custom MCP > skill > local script**, in order of preference for new solutions.

11.5 Creating an MCP server from scratch: publishing to WordPress

Use case: we want Claude Code to be able to **publish content to a WordPress site** without having to open the wp-admin dashboard. Typical scenario: you're discussing an article with Claude (writing it in Markdown, it suggests editorial cuts), and you want to close the session with "publish this article as a draft in the 'Tutorial' category". Without MCP, Claude can at most tell you "here are the steps to do in admin". With MCP, it does it.

PREREQUISITES

1. A WordPress site with active REST API (default since WP 5.0+).
2. An **Application Password** generated from the WordPress user profile (user → edit profile → "Application Passwords"). It's a single-use password, long, separate from the main password: it's revoked without having to change the real password.
3. **Python 3.10+** and the official SDK: `pip install mcp httpx python-dotenv`.

PROJECT STRUCTURE

```
~/mcp/wordpress-publisher/  
├─ server.py  
├─ pyproject.toml  
└─ .env
```

`.env` (never commit):

```
WP_BASE_URL=https://mysite.example.com  
WP_USERNAME=maurizio  
WP_APP_PASSWORD=xxxx xxxx xxxx xxxx xxxx
```

SERVER CODE (`SERVER.PY`)

```

"""
MCP server that publishes articles to WordPress via REST API.
Three exposed tools:
- wp_create_post      : creates a draft (or publishes directly)
- wp_publish_post     : promotes an existing draft to "publish"
- wp_list_categories  : lists the site's categories
"""

import os
import base64
import httpx
from mcp.server.fastmcp import FastMCP
from dotenv import load_dotenv

# Load credentials from .env (never hardcode in repo)
load_dotenv()
BASE_URL = os.environ["WP_BASE_URL"].rstrip("/")
USERNAME = os.environ["WP_USERNAME"]
APP_PASS = os.environ["WP_APP_PASSWORD"]

# Basic Auth with Application Password (standard WP format)
auth_token = base64.b64encode(
    f"{USERNAME}:{APP_PASS}".encode("utf-8")
).decode("ascii")
HEADERS = {
    "Authorization": f"Basic {auth_token}",
    "Content-Type": "application/json",
}

# FastMCP instance: the server will be named "wordpress-publisher"
mcp = FastMCP("wordpress-publisher")

@mcp.tool()
def wp_create_post(
    title: str,
    content: str,
    status: str = "draft",
    categories: list[int] | None = None,
) → dict:
    """
    Creates a new article on WordPress.

    Args:
        title:      post title
        content:    HTML or Gutenberg-block body
        status:     "draft" (default) or "publish"
    """

```

```

        categories: list of category IDs (optional)

Returns:
    Dict with id, status, link, modified of the created post.
    """
    payload = {"title": title, "content": content, "status": status}
    if categories:
        payload["categories"] = categories

    response = httpx.post(
        f"{BASE_URL}/wp-json/wp/v2/posts",
        headers=HEADERS,
        json=payload,
        timeout=30.0,
    )
    response.raise_for_status()
    data = response.json()
    return {
        "id": data["id"],
        "status": data["status"],
        "link": data["link"],
        "modified": data["modified"],
    }

@mcp.tool()
def wp_publish_post(post_id: int) → dict:
    """
    Promotes an existing draft to "publish" status.

    Args:
        post_id: ID of the post to publish

    Returns:
        Dict with id, status, updated link.
    """
    response = httpx.post(
        f"{BASE_URL}/wp-json/wp/v2/posts/{post_id}",
        headers=HEADERS,
        json={"status": "publish"},
        timeout=30.0,
    )
    response.raise_for_status()
    data = response.json()
    return {
        "id": data["id"],
        "status": data["status"],
        "link": data["link"],
    }

```

```

    }

@mcp.tool()
def wp_list_categories() → list[dict]:
    """
    Lists all the site's categories.

    Returns:
        List of dict with id, name, slug, count for each category.
    """
    response = httpx.get(
        f"{BASE_URL}/wp-json/wp/v2/categories",
        headers=HEADERS,
        params={"per_page": 100},
        timeout=30.0,
    )
    response.raise_for_status()
    return [
        {"id": c["id"], "name": c["name"], "slug": c["slug"], "count": c["count"]}
        for c in response.json()
    ]

if __name__ == "__main__":
    # Server startup in stdio mode (default for local MCPs)
    mcp.run()

```

REGISTRATION IN CLAUDE CODE

Add to your `.claude/settings.json` (or the global one):

```

{
  "mcpServers": {
    "wordpress-publisher": {
      "command": "python",
      "args": ["/Users/maurizio/mcp/wordpress-publisher/server.py"]
    }
  }
}

```

INTERACTIVE TEST

Launch Claude Code, run `/mcp` to verify that the `wordpress-publisher` server is connected and showing three tools. Then a typical conversation:

You: I just finished writing an article on Plan Mode, you'll find it in `article-plan-mode.md`. Publish it as a draft in the "Tutorial" category.

Claude: I'll first read the file to extract title and content. [Read `article-plan-mode.md`] I read the article: 1840 words, title "Plan Mode: thinking before writing". I'll fetch the category IDs to find "Tutorial". [wp_list_categories] "Tutorial" category has ID **7**. Proceeding with draft creation. [wp_create_post(title=..., content=..., status="draft", categories=[7])] ✓ Draft created. ID **142**, link <https://mysite.example.com/?p=142&preview=true>. Do you want me to publish it now or keep it as a draft for review?

You: Keep it as a draft, I need to review the images.

WHAT HAPPENS UNDER THE HOOD

Claude Code:

1. Read `article-plan-mode.md` with its native Read tool.
2. Understood that the category ID was needed — called `wp_list_categories` (an MCP tool).
3. Searched for "Tutorial" among the returned categories and extracted ID 7.
4. Called `wp_create_post` with title, content, and categories.
5. Reported the outcome with the preview link.

All via JSON-RPC between the client (Claude Code) and our Python server running as a local subprocess. No network hops on sensitive data — only HTTPS requests to the WordPress server, exactly as if the user had made them from a REST client.

NATURAL EXTENSIONS

From this skeleton it's easy to grow:

- add `wp_upload_media` to upload images to the library;
- add `wp_list_drafts` to retrieve existing drafts;
- add `wp_schedule_post` for scheduled publications (`status: "future"` with `date`);
- expose categories and posts as **resources** (URI `wp://categories`, `wp://posts/{id}`) to give Claude visibility of the catalog without having to invoke a tool every time.

For those coming from the WordPress plugin world: this MCP server is essentially an **AI-side REST client**. Everything your plugin can do via REST API, your MCP can expose as a tool.

11.6 Security and operational considerations

Three points of attention that separate an MCP server from an experiment and an MCP server ready for the trade.

No automatic sandbox. An MCP server runs as you: it has your file system credentials, network access, and tokens you pass it via env. Nothing isolates it from the rest of the system. Practical consequences:

- **Audit the code** before installing third-party MCP servers, especially if obtained from less-supervised marketplaces. A malicious server can read your `~/ .ssh` or exfiltrate secrets from env.
- **Keep your custom MCP servers in repos you control**, not as anonymous npx dependencies.
- **Use Application Password / API key with minimum scope**, never the main personal password. You revoke them with one click if needed.

Permission deny for sensitive tools. MCP tools flow into Claude Code's permission system (chapter 9). For servers with risky tools (delete posts, run DELETE queries) it's worth denylisting destructive tools in `permissions.deny`:

```
{
  "permissions": {
    "deny": [
      "mcp__wordpress-publisher__wp_delete_post",
      "mcp__postgres__query_write"
    ]
  }
}
```

The `mcp__<server>__<tool>` pattern allows precision targeting of the tool you don't want auto-approved; Claude will continue to ask for explicit confirmation at every invocation.

Logging and audit. To understand what your MCP is really doing in production, pair a **PostToolUse hook** (chapter 13) that logs every MCP tool invocation in JSON Lines. Concrete patterns in chapter 13.6 (example B). Effect: total traceability of who (user, model), what (which MCP tool), when, and with what arguments.

Remote MCPs. Servers reached via HTTP+SSE (hosted servers, shared between teams) add a dimension: latency and network auth. For stable team-level integrations a hosted server is preferable; for experimentation and personal integrations, local studio is simpler and more secure by default. The protocol is the same, only the transport changes.

When you DON'T need an MCP. If the task is purely local (reading a file, executing a script), Claude Code already has Read/Write/Bash as native tools: writing an MCP to do what Bash already does is overkill. The rule is: **MCP for external services or network protocols; native tools for the local user system.** When in doubt, first skill (chapter 10) or custom slash command; MCP only when it's an external system Claude needs to dialogue with via API.

12. Subagents: orchestrating specialized work

Subagents are one of Claude Code's most powerful mechanisms, and also one of the least understood. They aren't "secondary agents" doing minor things: they are instances of Claude with **their own context, their own tools, their own prompt personality**, that the main session can invoke to delegate specialized work. An entire chapter is needed to frame them well because they change the way you set up complex workflows.

12.1 What they are and why you need them

The most useful analogy is **browser tabs**. When you work on a complex task, the main agent is like a main tab that fills with child tabs (read files, tool output, searches): the further it goes, the heavier the tab gets. Subagents are **separate tabs**: they have their own context, do their work, and when finished pass to the main tab only the result — not everything they read to get there.

In practice they solve three distinct problems:

1. **Saturated context.** A search that requires reading 50 files to find a pattern, done by the main agent, leaves 50 files in the context. Done by a subagent, it leaves only the final summary. See also section 8 on context management.
2. **Specialization.** A generalist main agent can do a code review, but a subagent with a system prompt targeted at "code reviewer specialized in WordPress security" does it better, with stable criteria from one session to another.

3. **Multiple delegation.** On an articulated task — “review this PR for security, performance, and style” — you can delegate the three angles of analysis to three different subagents and receive three independent summaries, instead of keeping everything in the same tab.

12.2 Subagent vs main agent: the concrete difference

Technically, a subagent differs from the main agent on four dimensions:

- **Separate context window:** starts from zero, sees only the prompt the main session passes to it.
- **Tool restrictions:** you can limit available tools (e.g., only `Read`, `Grep`, `Glob` for a read-only agent).
- **Dedicated system prompt:** it has its own instructional “personality”, independent from the main session.
- **Model override:** it can run on a different model (e.g., Haiku for speed on massive tasks while the main session stays on Sonnet).

The subagent isn’t directly invocable by you: Claude invokes it via the `Agent` tool (renamed from `Task` in v2.1.63 — the `Task(...)` alias is still active for backward compatibility). You ask the main agent for a result, and the main agent decides whether to delegate to a subagent based on the latter’s `description`.

12.3 Built-in subagents

Claude Code includes some subagents available out-of-the-box. The three relevant for daily use are:

- **Explore** — runs on Haiku, is read-only (`Glob`, `Read`, `Grep`, `Bash`). The main agent invokes it for codebase searches when a query requires reading many files. Accepts a depth level: `quick` (targeted lookup), `medium` (moderate exploration), `very thorough` (exhaustive search across multiple naming conventions).
- **Plan** — used in Plan Mode (see section 5). Inherits the session’s model, is read-only, exists because in Plan Mode searches must also remain in the non-destructive perimeter.
- **general-purpose** — the generalist agent for multi-step tasks when you don’t have a specialized agent. Inherits the model and has access to all tools. It’s the one the main agent invokes when you say “delegate this thing to a subagent” without specifying which one.

There are two more built-in agents for more internal use: `statusline-setup` (configures the status line, activated by `/statusline`) and `Claude Code Guide` (runs on Haiku, answers questions about the CLI itself). I name them only for completeness.

12.4 Creating a custom subagent

You have two ways to create a subagent: **interactive** with the `/agents` command, or **manual** by writing the file.

Via `/agents`. Opens a tabbed interface: “Running” lists active agents in the session, “Library” shows available ones with the option to create new ones. The creation flow asks you for scope (Personal / Project), agent description, allowed tools, model, display color, memory. Claude can also **generate the first draft of the system prompt** starting from your description — useful as scaffolding, to be refined by hand.

Via file. A subagent is a Markdown file with YAML frontmatter. The path determines the scope:

- **Project:** `.claude/agents/<name>.md` — committed in the repo, shared with the team.
- **User:** `~/.claude/agents/<name>.md` — personal, accompanies you on all the machine’s projects.

Concrete example in the WordPress space, a subagent dedicated to security audit of plugins:

```

---
name: wp-security-auditor
description: |
  Use this agent to audit WordPress plugin code for security issues.
  Triggers on: PHP files in wp-content/plugins, mentions of nonce,
  capability checks, sanitize_*, esc_*, $wpdb queries, REST API
  endpoints, AJAX handlers.
tools: Read, Grep, Glob
model: sonnet
color: red
---

```

You are a security auditor specialized in WordPress plugins.

For every PHP file you analyze, verify in order:

1. **Nonce verification**: on every handler that modifies state (form submit, AJAX action, REST endpoint).
2. **Capability check**: (`current_user_can()`) before privileged actions.
3. **Input sanitization**: all `$_GET`, `$_POST`, `$_REQUEST` passed through `sanitize_text_field`, `sanitize_email`, `absint`, etc., according to the expected type.
4. **Output escaping**: every printed string must pass through `esc_html`, `esc_attr`, `esc_url`, `wp_kses_post` according to context.
5. **SQL queries**: always use `$wpdb->prepare()` for queries with variables. Never direct concatenation.
6. **File operations**: validate path, avoid path traversal.

Regulatory references:

- Plugin Security Handbook: <https://developer.wordpress.org/plugins/security/>
- WordPress Coding Standards: <https://developer.wordpress.org/coding-standards/>

Structured output for each file:

- File and line
- Severity: critical / high / medium / low
- Vulnerability identified
- Recommended fix with example of correct code

Don't modify code. Limit yourself to the report.

The tools are limited to `Read`, `Grep`, `Glob`: the agent **cannot** modify files, launch shell commands, access the network. It’s a read-only reviewer, by design — exactly what you want from a security audit.

The most relevant frontmatter fields:

Field	Type	Function
<code>name</code>	string (required)	Unique ID, lowercase with hyphens, max 64 characters
<code>description</code>	string (required)	When the main agent should delegate to this subagent
<code>tools</code>	list	Allowlist of allowed tools (default: inherits all)
<code>disallowedTools</code>	list	Denylist of forbidden tools (applied before <code>tools</code>)
<code>model</code>	string	<code>haiku</code> , <code>sonnet</code> , <code>opus</code> , <code>inherit</code> or full ID
<code>permissionMode</code>	string	<code>default</code> , <code>acceptEdits</code> , <code>auto</code> , <code>plan</code> , <code>bypassPermissions</code>
<code>color</code>	string	Display color in the session
<code>memory</code>	string	Auto Memory scope: <code>user</code> , <code>project</code> , <code>local</code> , or absent
<code>isolation</code>	string	Set to <code>worktree</code> to isolate the agent in a temporary git worktree

12.5 Precedence hierarchy

If a subagent with the same name exists at multiple levels, **the most specific one** wins. Order of precedence, from highest to lowest:

1. **Managed settings** (organizational configuration, rare in individual setups)
2. `--agents` **CLI flag** (subagents defined for the single session via launch flag)
3. **Project’s** `.claude/agents/`
4. **User’s** `~/.claude/agents/`

5. Agents provided by installed plugins

Practical convention: keep in `~/ .claude/agents/` the **personal and generic** agents (e.g., `code-reviewer-style`, `commit-message-writer`), and in the repo's `.claude/agents/` the **project-specific** agents, shared with the team via Git. Plugin agents use the namespace `<plugin>:<agent-name>` so they don't collide.

12.6 Automatic vs explicit invocation

There are two ways to activate a subagent.

Automatic. The main agent reads the `description` of all available subagents and decides whether to delegate based on the match with the current task. That's why the `description` is the most important field of the frontmatter: write it thinking of the **concrete triggers** of the workflow, not as a marketing paragraph.

```
# Bad (vague, doesn't trigger well)
description: Agent for WordPress security.

# Good (specific, indicates real triggers)
description: |
  Use this agent to audit WordPress plugin code for security issues.
  Triggers on: PHP files in wp-content/plugins, mentions of nonce,
  capability checks, sanitize_*, esc_*, $wpdb queries.
```

Explicit. When you want to force the use of a specific agent, you mention it with `@`:

```
@agent-wp-security-auditor analyze the latest changes to the
plugin in wp-content/plugins/access-control/
```

Typing `@` opens a typeahead picker showing the available agents. For agents provided by plugins the syntax is `@agent-<plugin>:<name>`.

12.7 Parallelism: multi-delegation patterns

A note of honesty on this point, because “parallel agents” are sold around with lightness.

Normal subagents in Claude Code run **sequentially** in foreground: the main agent launches one, waits for the result, launches another. There's no standard simultaneous concurrency. What we call “parallelism” is almost always a **multi-delegation with context isolation**: you launch three independent subagents one after the other, each with its own tab, and you receive three independent summaries — the advantage is the separation of context, not execution time.

“Parallel deep review” pattern (more correctly: delegated deep review):

You: "Review this PR with three angles: security, performance, style. Delegate each angle to a dedicated subagent and summarize the three reports at the end."

Claude:

- Agent(wp-security-auditor) [security analysis] ← 60s
Result: 3 security issues in 2 files
- Agent(performance-reviewer) [performance analysis] ← 45s
Result: 2 N+1 queries, 1 missing cache
- Agent(style-reviewer) [style analysis] ← 30s
Result: 5 PSR-12 violations, 2 obsolete comments

Unified summary for the PR.

The real advantage is that **the main agent didn't bloat** by reading the files: each subagent did it in its own tab, and only the final aggregation remains in the main. Total time: sum of the three.

If true simultaneous parallelism is needed, there's the **Agent Teams** feature: multiple independent instances of Claude Code that really run concurrently, coordinated via shared task list and message mailbox (teammates write to each other directly, not just to the lead like in subagents). Requires Claude Code v2.1.32+ and the env `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1` (or equivalent in `settings.json`).

⚠ Experimental feature, to be handled with caution. Anthropic declares Agent Teams experimental and disabled by default in the [official page](#). Relevant known limitations:

- `/resume` and `/rewind` don't restore in-process teammates — after a resume you may have to regenerate the team
- task state can "lag behind": a task not marked completed blocks dependents
- shutdown is slow (teammates finish the current turn before exiting)
- **only one team per session**, no nested teams, lead non-reassignable
- permissions settable only at spawn (subsequent changes per teammate)
- **split-pane requires tmux or iTerm2**: doesn't work in VS Code terminal, Windows Terminal, Ghostty

For most workflows, the standard multi-delegation to subagents (described above) is sufficient. Open Agent Teams when you really need teammates dialoguing among themselves — parallel code reviews on different perspectives, debugging with concurrent hypotheses, cross-layer refactor where each teammate owns a different layer.

If you decide to use them, the dedicated hook events `TeammateIdle`, `TaskCreated`, and `TaskCompleted` (see section 13.3 — Lifecycle events) allow applying quality gates that the Anthropic docs suggest to discipline teammate behavior: block malformed tasks at creation, reject “completed” if tests fail, restart an idle teammate when it still has pending work.

12.8 Cost optimization via model routing

The frontmatter’s `model` field enables one of Claude Code’s most underestimated patterns: **routing simple tasks to Haiku** instead of Sonnet. Haiku costs a fraction of Sonnet (input/output) and on well-defined tasks — pattern matching across many files, classification, structured extraction — quality is more than adequate.

Typical case: you need to audit 50 PHP files to find which use `$wpdb->query()` with concatenation instead of `$wpdb->prepare()`. It’s a pattern recognition task, not architectural reasoning.

```
---
name: wp-sql-injection-scanner
description: |
  Use this agent to scan PHP files for direct concatenation in
  $wpdb queries (potential SQL injection). Triggers on: requests
  to audit SQL safety, mentions of $wpdb, manual SQL hardening.
tools: Read, Grep, Glob
model: haiku    # Haiku is plenty, costs much less
---
```

You are a WordPress SQL injection scanner.

Look for all files that call `$wpdb->query()`,
`$wpdb->get_results()`, `$wpdb->get_var()`, `$wpdb->get_row()`
 with SQL strings that concatenate variables (``` or interpolation).

For each occurrence:

- File and line
- Snippet of the offending query
- Corrected version using `$wpdb->prepare()`

Ignore cases where the query is 100% static (no variables).

The main agent (Sonnet) coordinates and produces the final report; the subagent (Haiku) does the massive scan work. The token bill drops significantly. The same principle applies to extractions, summaries, classifications — see also the logic of `opusplan` in section 5 for another application of the “right model for the right thing” pattern.

12.9 When NOT to use them

Subagents aren't always the right choice:

- **Short tasks.** If the task requires reading 2-3 files, the main agent does it directly in less time. The setup overhead (task description to subagent, wait, parsing the result) exceeds the benefit.
- **When the summary loses information.** If you need the main agent to see **the content** of certain files — not just a summary — delegation hobbles you. The subagent returns its synthesis, not the raw data.
- **Iterative work.** If you're doing incremental refactor that requires continuous back-and-forth (modify, test, modify again), a subagent doesn't help: the main agent is already the right tool, simply with periodic `/compact` to not bloat the context.
- **High-risk tasks where you want direct control.** Destructive operations (delete, force push, DB migrations), interpretation errors by the subagent are costly. Keep them on the main agent in Plan Mode.

When in doubt: start with the main agent. If you find you've read 30 files just to produce 200 tokens of output, it was a subagent.

12.10 Subagents, Skills and Hooks compared

Claude Code has **three extension mechanisms** that are easy to confuse. It's worth fixing them together because they do different things and often combine.

Aspect	Subagent	Skill	Hook
What it is	Specialized agent with separate context	Reusable playbook inserted in the main context	Script that intercepts lifecycle events
Who writes it	You (<code>.md</code> file with YAML frontmatter)	You (<code>SKILL.md</code> file with YAML frontmatter)	You (bash/HTTP/prompt script)
How it activates	Main agent delegation or explicit <code>@agent-name</code>	Automatic match on <code>description</code> or <code>/skill-name</code>	Automatically on event
Context scope	Separate context window	Inline in main context	Side effect, doesn't add context
Output	Summary to main agent	Content integrated into the conversation	<code>allow/deny/ask</code> decision or action
Where they live	<code>.claude/agents/</code> , <code>~/.claude/agents/</code> , plugins	<code>.claude/skills/</code> , <code>~/.claude/skills/</code> , plugins	<code>settings.json</code> or <code>.claude/settings.json</code>
Typical use case	Isolated search, specialized review, massive tasks	Project conventions, recurring playbooks, domain knowledge	Command validation, audit, blocking risky operations

Said in one line: **a Skill enriches the main agent, a Subagent replaces it for the delegated task, a Hook does something around the main agent without being part of it.**

The three mechanisms often **combine**: a Skill can instruct the main agent to delegate a certain task to a specific Subagent, and that Subagent can have a Hook that validates its tool calls before they're executed.

For complete treatment of Hooks — events, types, examples, security — see section 13.

13. Hooks: automating Claude Code's lifecycle

Hooks are Claude Code's third extension mechanism, and the one that operates at the lowest level. While Subagents and Skills influence what the main agent does, Hooks intercept when certain events happen — the start of a session, the moment before a tool execution, the end of a response — and can validate, modify, block, register those events before or after they happen.

They're the right tool when the extension isn't "let it know how it works" (Skill) or "delegate this task" (Subagent), but "intervene automatically at this precise point of the flow, without me having to remember to ask".

13.1 What they are and what they're for

A Hook is a script — bash, HTTP, prompt to another model, subagent, or MCP tool — configured to fire at a determined event of Claude Code's lifecycle. The script receives via input a JSON payload describing the event (e.g., which tool is about to be executed, with what arguments) and produces an output that can:

- **block** the action (e.g., prevent `rm -rf` on a protected folder)
- **validate** and let it pass (e.g., check that every edit of a `.php` file respects style rules)
- **register** in a structured log (audit trail of all modifications)
- **inject** automatic context (e.g., remind Claude of the project conventions at every session start)

At first glance they look like a linter or a git hook, but with an important difference: the linter acts after the code has been written by you, Hooks act during Claude Code's execution, before its actions become effective. It's an internal automation level, not an external check.

A good rule: if what you want to do is “every time Claude does X, you must do Y”, it’s probably a Hook. If it’s “Claude must know that the project uses X”, it’s a Skill or a `CLAUDE.md`. If it’s “delegate this specific work to another agent”, it’s a Subagent.

13.2 Anatomy of a hook

Hooks are configured in the `settings.json` file along with permissions and other directives (see section 9.2 for the basic setup and for the `$schema` line that enables autocomplete and validation also of `hooks` blocks). The basic syntax is the same for all events:

```
{
  "hooks": {
    "<EventName>": [
      {
        "matcher": "<pattern>",
        "hooks": [
          {
            "type": "command",
            "command": "<script-path>"
          }
        ]
      }
    ]
  }
}
```

Configuration files live at three levels (plus two reduced-scope locations):

- `~/.claude/settings.json` — user, valid on all machines. Personal, not shareable.
- **Project’s** `.claude/settings.json` — committed in Git, shared with the team.
- **Project’s** `.claude/settings.local.json` — gitignored by default. For Hooks specific to your local copy.
- **Skill or Agent frontmatter** (`hooks` field) — Hooks that run only when that Skill/Agent is active.
- **Installed plugins’** `hooks/hooks.json` — Hooks that come with a plugin.

There are **five types of Hooks**, chosen via the `type` field:

Type	When to use it	Example scenario
<code>command</code> (default)	Deterministic logic, file access, JSON parsing	Block a command, write a log, launch a script
<code>http</code>	Centralized audit, cloud integration (introduced Feb 2026)	POST to a corporate endpoint for immutable logs
<code>prompt</code>	"Yes/no" decision requiring LLM judgment	"Are all prompt tasks completed before terminating?"
<code>agent</code>	Complex verification requiring tools and codebase search	A subagent that runs the test suite before Stop
<code>mcp_tool</code>	Integration with an already configured MCP server	Save context in an external MCP memory

For 90% of practical cases you use `command`. Other types are for advanced scenarios or when the decision requires capabilities a simple script doesn't have.

13.3 Lifecycle events

Claude Code exposes more than twenty interceptable events, groupable into seven categories. Synthetic table of the events most used in practice:

Category	Event	When it fires	Can it block?
Session	<code>SessionStart</code>	Session start or resume	No — can inject context
Session	<code>SessionEnd</code>	Session end	No — only cleanup
Session	<code>UserPromptSubmit</code>	User submits a prompt	Yes — can modify/block the prompt
Session	<code>Stop</code>	Claude finishes a response	Yes — can force to continue
Tool	<code>PreToolUse</code>	Before a tool's execution	Yes — can deny/modify the input
Tool	<code>PostToolUse</code>	After a tool's execution	No — the action is already done
Tool	<code>PostToolUseFailure</code>	After a failed tool	No
Permissions	<code>PermissionRequest</code>	Shows permissions dialog	Yes — can decide allow/deny/ask
Subagent	<code>SubagentStart</code>	Subagent spawn	No
Subagent	<code>SubagentStop</code>	Subagent terminates	Yes — can force a retry
Compaction	<code>PreCompact</code>	Before <code>/compact</code>	No — can save backup
Compaction	<code>PostCompact</code>	After <code>/compact</code>	No — can re-inject context
File	<code>FileChanged</code>	Watched file modified	No

There are other more specialized events (`UserPromptExpansion`, `ConfigChange`, `CwdChanged`, `WorktreeCreate/Remove`, `Notification`, `Elicitation`, and the Agent Teams events like `TaskCreated`, `TaskCompleted`, `TeammateIdle`). For the complete list the canonical source is code.claude.com/docs/en/hooks.

The two events you'll use most often are `PreToolUse` and `PostToolUse`. The whole "tool guardian" pattern is built on them.

13.4 Matchers and inspection (`/hooks`)

The `matcher` field filters **which invocations** of the event must activate the Hook. Four supported patterns:

- **Wildcard** `"*"` or empty string — matches all
- **Exact** `"Bash"` — only invocations of the `Bash` tool
- **Pipe** `"Edit|Write"` — alternation, matches `Edit` or `Write`
- **Regex** `"mcp__.*"` — standard regex syntax, here to match all MCP tools

The matcher's meaning depends on the event: for `PreToolUse` / `PostToolUse` it's the tool name; for `SessionStart` it's the source (`startup`, `resume`, `clear`, `compact`); for `SubagentStart` / `SubagentStop` it's the agent name. Without matcher (or with `"*"`), the Hook always fires.

From v2.1.85+ there's also the `if` field for secondary filters on the tool's **arguments**, not just on the name:

```
{
  "matcher": "Bash",
  "hooks": [
    {
      "type": "command",
      "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/validate-git.sh",
      "if": "Bash(git *)"
    }
  ]
}
```

This Hook fires only for `git ...` and ignores other Bash commands — useful to restrict further.

`/hooks` **command**. Once a Hook is written it's easy to wonder if the configuration was actually loaded. `/hooks` opens an interactive (read-only) browser of the active configuration: for each event it shows how many Hooks are registered, from which settings file they come, with which matcher. It doesn't allow **editing** or **disabling** Hooks on the fly: for that you need to modify `settings.json` and wait for it to be reloaded (or restart the session). The typical development flow is: edit `settings.json` → `/hooks` for loading verification → invoke a target tool to test the behavior.

13.5 Input and output

A `command` type Hook receives via **stdin** a JSON object describing the event. Fields common to all events:


```
{
  "session_id": "abc123",
  "cwd": "/home/user/wp-plugins/access-control",
  "hook_event_name": "PreToolUse",
  "transcript_path": "/home/user/.claude/projects/.../transcript.jsonl"
}
```

For `PreToolUse` `tool_name` and `tool_input` are added:

```
{
  "session_id": "abc123",
  "hook_event_name": "PreToolUse",
  "tool_name": "Bash",
  "tool_input": { "command": "rm -rf wp-content/uploads" }
}
```

The **output behavior** is governed by:

- **Script exit code:**

- `0` — success. If stdout is valid JSON, it's interpreted as structured output; if it's plain text, it's added to the context as `additionalContext`.
- `2` — **block**: the action doesn't proceed. stderr is shown to Claude as the reason for the block.
- Other values — non-blocking error, the action proceeds but is registered in the transcript.

- **Structured JSON output** (on stdout, exit 0):

```
{
  "hookSpecificOutput": {
    "hookEventName": "PreToolUse",
    "permissionDecision": "deny",
    "permissionDecisionReason": "rm -rf on wp-content is forbidden",
    "updatedInput": { "command": "echo 'command blocked'" },
    "additionalContext": "To delete uploads use the backup script."
  }
}
```

Most-used fields:

- `permissionDecision` (PreToolUse): `allow` / `deny` / `ask` / `defer`
- `decision` (PostToolUse, Stop): `allow` / `block`
- `additionalContext`: text injected to Claude as a system reminder
- `updatedInput`: modify the tool input before execution (e.g., force safe flags)

- `continue` / `stopReason`: control of the flow in `Stop` events

Multiple Hooks on the same event? “Most restrictive wins”. If two Hooks on `PreToolUse` return one `allow` and the other `deny`, `deny` wins. All `additionalContext`s are concatenated. If two Hooks set `updatedInput`, the order isn’t guaranteed: avoid configurations where two different Hooks modify the same input.

13.6 Practical examples

Four examples in the WordPress space, from the simplest to the most articulated. For each: scenario, configuration, script, observed behavior.

EXAMPLE A — BLOCK `RM -RF` IN `WP-CONTENT/`

Scenario. You’re working on a plugin and want a safety net that prevents Claude from launching `rm -rf` on any path containing `wp-content/`. An oversight in a shell command could destroy uploads, cache, or backups of a production site.

`.claude/settings.json`:

```
{
  "hooks": {
    "PreToolUse": [
      {
        "matcher": "Bash",
        "hooks": [
          {
            "type": "command",
            "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/block-rm-wpcontent.sh"
          }
        ]
      }
    ]
  }
}
```

`.claude/hooks/block-rm-wpcontent.sh`:

```
#!/bin/bash
INPUT=$(cat)
COMMAND=$(echo "$INPUT" | jq -r '.tool_input.command // ""')

if echo "$COMMAND" | grep -qE 'rm\s+(-[a-z]*r[a-z]*\s+|-r\s+).*wp-content'; then
    echo "Blocked: recursive rm on wp-content forbidden by hook" >&2
    exit 2
fi
exit 0
```

Behavior. When Claude tries to launch `rm -rf wp-content/uploads`, the Hook intercepts, detects the pattern, and returns exit 2 with a message in stderr. Claude sees the block, receives the reason, and decides how to proceed (asking you for confirmation or changing approach). `rm` commands on other paths pass normally.

EXAMPLE B — JSON LINES AUDIT LOG ON EDIT/WRITE

Scenario. You want an audit trail of all modifications made by Claude to the plugin's PHP files, in JSON Lines format for subsequent analysis (grep, jq, dashboard).

`.claude/settings.json`:

```
{
  "hooks": {
    "PostToolUse": [
      {
        "matcher": "Edit|Write",
        "hooks": [
          {
            "type": "command",
            "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/audit-php.sh"
          }
        ]
      }
    ]
  }
}
```

`.claude/hooks/audit-php.sh`:

```
#!/bin/bash
INPUT=$(cat)
FILE=$(echo "$INPUT" | jq -r '.tool_input.file_path // ""')

if [[ "$FILE" = *.php ]]; then
  jq -nc --arg ts "$(date -u +%Y-%m-%dT%H:%M:%SZ)" \
    --arg session "$(echo "$INPUT" | jq -r '.session_id')" \
    --arg tool   "$(echo "$INPUT" | jq -r '.tool_name')" \
    --arg file   "$FILE" \
    '{ts: $ts, session: $session, tool: $tool, file: $file}' \
    >> "$CLAUDE_PROJECT_DIR/.claude/audit/php-edits.jsonl"
fi
exit 0
```

Behavior. Every edit or write on a `.php` file adds a line to the file `audit/php-edits.jsonl`. Other modifications (CSS, MD, JSON) are ignored. The log is structured and readable with `jq -s '[] | select(.file | contains("admin"))' audit/php-edits.jsonl` to filter retrospectively.

EXAMPLE C — PHPCS WITH WORDPRESS-EXTRA AUTOMATIC, ASYNC

Scenario. Every PHP file modified by Claude must be passed to `phpcs` with the `WordPress-Extra` standard. You want the lint to run **in background** without blocking Claude (async is available since January 2026).

`.claude/settings.json`:

```
{
  "hooks": {
    "PostToolUse": [
      {
        "matcher": "Edit|Write",
        "hooks": [
          {
            "type": "command",
            "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/phpcs-wp.sh",
            "async": true
          }
        ]
      }
    ]
  }
}
```

`.claude/hooks/phpcs-wp.sh`:

```
#!/bin/bash
INPUT=$(cat)
FILE=$(echo "$INPUT" | jq -r '.tool_input.file_path // ""')

if [[ "$FILE" = *.php ]] && command -v phpcs >/dev/null; then
    phpcs --standard=WordPress-Extra "$FILE" \
        > "$CLAUDE_PROJECT_DIR/.claude/lint/$(basename "$FILE").log" 2>&1
    || true
fi
exit 0
```

Behavior. `phpcs` runs in background after each edit, writing the report to `.claude/lint/`. Claude doesn't wait for completion (it's async, decision fields would be ignored anyway). You review the logs at the end of the session or with a watcher in the IDE. For hard commit blocking, use a separate Git pre-commit hook — Claude Code Hooks don't replace Git hooks, they work at a different level.

EXAMPLE D — CONVENTIONS REMINDER AT `SESSIONSTART`

Scenario. Every time you open Claude Code inside the plugin, you want the main agent to receive a reminder of the project conventions (in addition to what's already written in `CLAUDE.md`).

`.claude/settings.json`:

```
{
  "hooks": {
    "SessionStart": [
      {
        "matcher": "startup",
        "hooks": [
          {
            "type": "command",
            "command": "$CLAUDE_PROJECT_DIR/.claude/hooks/wp-session-re-
minder.sh"
          }
        ]
      }
    ]
  }
}
```

`.claude/hooks/wp-session-reminder.sh`:

```
#!/bin/bash
cat <<'REMINDER'
Reminder for this session:
- All PHP hooks must verify nonce and capability.
- Output always escaped (esc_html, esc_attr, wp_kses_post).
- SQL queries only via $wpdb->prepare().
- Code style: PSR-12 + WordPress-Extra where the two diverge, WordPress
wins.
REMINDER
exit 0
```

Behavior. At session startup (matcher `startup`, not on `resume`), the main agent receives the reminder as `additionalContext`. It's complementary to `CLAUDE.md`: the Markdown file covers stable rules, the Hook can inject dynamic reminders (e.g., building the message by reading the plugin state, the current branch, or recent events).

13.7 Security

Hooks are powerful because they execute **arbitrary code** with your user permissions. There's no sandbox: a Hook script can read `.env`, make network requests, leave traces on the filesystem. This power is the same that makes them an attack vector if not handled with attention.

CVE-2025-59536 (RCE via hook injection). In early 2026 a case of Remote Code Execution was documented exploiting the fact that `.claude/settings.json` is automatically loaded from the project root. A malicious repository can register a Hook that executes arbitrary script at the first launch of Claude Code in that directory. You find it tracked as CVE-2025-59536 on Check Point Research.

Five practical security rules.

1. **Mandatory code review on `.claude/settings.json`.** Treat it as an executable script, not as a passive configuration file. Changes to that file require review exactly like a modification to the project's `Makefile`.
2. **Untrusted repositories: clone, inspect, then open Claude Code.** Don't launch `claude` inside a freshly cloned repo without having opened `.claude/settings.json` and `.claude/hooks/` to check what they contain.
3. **HTTP hooks: explicit env var whitelist.** The `allowedEnvVars` field defines which environment variables can be interpolated in headers. Don't whitelist secrets that aren't needed for the specific endpoint.
4. **Sanitize input before injecting it as `additionalContext`.** A Hook payload contains `tool_input` coming from the session. If you pass that raw text to Claude as context, you open the door to prompt injection (a malicious `tool_input` injected in turn by a file read by Claude).

5. **Global disabling for debug.** The `"disableAllHooks": true` flag in `settings.json` turns off all Hooks. Useful in debug phase when you suspect a Hook is interfering, or if you have doubts about the origin of a loaded configuration.

Hooks configured at **user** level (`~/.claude/settings.json`) are under your control and follow you everywhere. Hooks configured at **project** level (`.claude/settings.json`) are under the control of anyone who can commit to the repo. Keep this distinction in mind when you receive a pull request that touches that file.

13.8 Gotchas and when NOT to use them

Hooks have a number of common traps. The six most frequent:

- **Relative paths in scripts.** The Hook runs with a non-guaranteed cwd. Always use `$CLAUDE_PROJECT_DIR` or absolute paths, never relative paths like `./scripts/check.sh`.
- **JSON parsing with regex.** The payload is JSON: use `jq` or a real parser. Trying to extract fields with `grep / sed` produces fragile scripts that break at the first special character.
- **Stop hook in loop.** If a Hook on `Stop` returns `continue: true`, Claude continues and at the end triggers `Stop` again, which triggers the Hook again... infinite loop. The payload contains `stop_hook_active: true` when you're in re-execution: read that flag and exit immediately if it's `true`.
- `async: true` **doesn't block.** Asynchronous Hooks are useful only for side-effects (logging, lint, notifications). If they return `decision: "block"` or `permissionDecision: "deny"`, those fields are **ignored**. To block you must be synchronous.
- `PostToolUse` **cannot undo.** The event is called "Post" because the tool has already been executed. You can log, format, send notifications, but you can't undo the action. To block you need `PreToolUse`.
- **Shell profile output that pollutes stdout.** If your `~/.bashrc` or `~/.zshrc` prints something even in non-interactive mode, that text ends up in front of the Hook's JSON output and breaks the parser. Wrap profile `echo`s in `if [[ $- = *i* ]]; then ... fi`.

When NOT to use a Hook.

- **Complex logic requiring LLM reasoning.** A `command` Hook is a deterministic script. If the decision requires judgment ("is this edit a good idea?"), you need a dedicated subagent (see section 12) or a `prompt` type Hook.
- **Style and formatting post-edit.** To apply a style guide to all modified files, a normal linter (Prettier, phpcs, eslint) launched from pre-commit or IDE is almost always more transparent and debuggable than a `PostToolUse` Hook.
- **One-off exploratory workflows.** The overhead of writing a script + configuring it + testing the matcher + verifying with `/hooks` isn't justified if you'll use the workflow only once.

When in doubt: **Hook for automatic and repeated behaviors you want to be invisible and infallible.** For everything else, more suitable tools exist.

14. Plugins: distributable packages

Plugins are the mechanism Claude Code uses to **distribute and install extensions**: skills, MCP servers, custom subagents, slash commands — all grouped into a single package installable with a command. If a skill teaches Claude what to do (chapter 10) and MCP tells it what to talk to (chapter 11), the plugin is the **container** that puts them together and distributes them. For those coming from the development world: think of the plugin as the npm/Composer package of the Claude Code ecosystem.

14.1 Claude Code's extension mechanisms: a map

By this point in the book you've seen **all four** of Claude Code's extension mechanisms that the plugin packages: **Skills** (chapter 10), **MCP** (chapter 11), **Subagents** (chapter 12), and **Hooks** (chapter 13). Before entering plugins, it's useful to reposition them in a unified map, to understand where each one operates and what relationship they have with the container we're about to examine.

Mechanism	What it does / where it acts	Where it lives	How it's distributed
Skill (chapter 10)	Extends what Claude knows how to do	Markdown + local scripts	Folder in user system, plugin
MCP (chapter 11)	Exposes tools/data of external systems	Server (local stdio or remote HTTP+SSE)	Open protocol, any language
Subagent (chapter 12)	Executes specialized work in isolation	YAML in <code>.claude/agents</code>	Config file, plugin
Hook (chapter 13)	Intercepts lifecycle events	<code>set-tings.json</code> hooks	Local config, plugin
Plugin (this chapter)	Package that groups the other four	Folder with manifest + payload	Marketplace, Git repo

In a vertical view: **Plugin** is the container, the other four are the **contents** that can be bundled inside a plugin. A “GitHub PR Assistant” plugin can contain: a review skill, an MCP server that talks to the GitHub API, a specialized subagent for writing PR descriptions, and a `/review` slash command that orchestrates everything. All distributed as a single package.

14.2 What a plugin is and why it exists

A plugin is born to solve a practical problem: when a person or organization maintains a set of skills, MCP servers, and custom workflows, distributing them individually is inconvenient. The consumer has to find the skill repo, copy the files to `.claude/skills/`, configure the MCP server in `set-tings.json`, register the slash commands — separate steps that don’t scale if the author updates the components frequently.

The plugin standardizes all this:

- **A single installation** collects all components
- **Centralized versioning** — the author releases a new version, you do `claude plugin update`
- **Clean uninstall** — removes everything in one shot
- **Shared marketplace** — authors publish the plugin on a public or private marketplace

Plugins are particularly precious for **organizations** that want to distribute to teams a coherent set of extensions (corporate skills + MCP for internal APIs + specialized subagents), and for **community authors** who produce thematic packages (Vercel Labs, WordPress Agent Skills, JuliusBrussee/caveman are all distributed as plugins).

14.3 Anatomy of a plugin

A plugin is a folder with a standardized structure. Minimal example:

```
my-plugin/
├── plugin.yml           # plugin manifest
├── skills/              # included skills (optional)
│   └── my-skill/
│       └── SKILL.md
├── mcp/                 # included MCP servers (optional)
│   └── my-server/
│       ├── server.py
│       └── pyproject.toml
├── agents/              # custom subagents (optional)
│   └── my-agent.yaml
├── commands/            # custom slash commands (optional)
│   └── my-command.md
└── README.md           # documentation
```

The `plugin.yml` file is the manifest that declares metadata, dependencies, and what the plugin contains:

```
name: my-plugin
version: 1.0.0
description: "Example plugin"
author: "Maurizio Pelizzone <maurizio@mavida.com>"
license: MIT
homepage: "https://github.com/mavida/my-plugin"

# Included components
includes:
  skills:
    - skills/my-skill
  mcp_servers:
    - name: my-server
      command: python
      args: ["mcp/my-server/server.py"]
  agents:
    - agents/my-agent.yaml
  commands:
    - commands/my-command.md

# Compatibility
requires:
  claude_code: "≥2.1.0"
```

The manifest declares everything the plugin installs. When the user does `claude plugin install`, Claude Code reads the manifest, copies/registers the components in the right slots (`~/.claude/skills/`, `mcpServers` in `settings.json`, etc.), and activates the plugin.

14.4 Plugin marketplace

A **marketplace** is a published index of available plugins, typically a Git repository with an expected structure. Claude Code supports:

- **Anthropic official marketplaces** (`anthropics/skills` and similar)
- **Community marketplaces** (Vercel Labs, Trail of Bits, JuliusBrussee)
- **Private organizational marketplaces** (an internal Git repo with corporate plugins)

Basic commands:

```
# Add a marketplace to your instance
/plugin marketplace add anthropics/skills

# Explore available plugins (opens interactive picker)
/plugin

# Install a specific plugin
/plugin install <plugin-name>

# Update installed plugins
claude plugin update

# Uninstall a plugin
claude plugin uninstall <plugin-name>
```

The `/plugin` picker shows plugin metadata (name, version, description, author) and a preview of included components, so you know what you're installing before download. Once installed, the plugin goes to `~/.claude/plugins/<name>/` and its components automatically become available in all sessions.

14.5 Creating a custom plugin

Let's see the minimal structure of a demo plugin. The scenario: a "hello-world" plugin containing a single `/hello` slash command and a welcome skill.

Step 1 — Create the folder and manifest in `~/my-plugins/hello-world/`:

```
# plugin.yml
name: hello-world
version: 0.1.0
description: "Example plugin"
author: "Maurizio Pelizzone"
license: MIT

includes:
  skills:
    - skills/hello-skill
  commands:
    - commands/hello.md
```

Step 2 — Add the skill in `skills/hello-skill/SKILL.md`:

```
---
name: hello-skill
description: "Use this skill when the user asks for greetings or examples
of plugin usage."
---

# Hello Skill

When the user asks for plugin examples or greetings, respond with:
- A brief greeting
- A note that this response comes from a skill installed via plugin
```

Step 3 — Add the slash command in `commands/hello.md`:

```
---
description: "Quick greeting from the hello-world plugin"
---

Greet the user in a friendly way and remind that this command comes
from the hello-world plugin (version 0.1.0).
```

Step 4 — Install locally for testing (without publishing):

```
# Install the plugin in development mode (link, not copy)
claude plugin install ~/my-plugins/hello-world --dev
```

Open Claude Code: type `/hello` and you should see the greeting. Ask Claude something that activates the skill — the `description` mentions “greetings” — and you should see the response that explicitly cites the plugin origin.

14.6 Distributing a plugin

Once the plugin works locally, distributing it requires three steps:

1. **Publish the Git repo** with the plugin structure in the root. The `plugin.yml` must be at the root of the folder.
2. **Create a marketplace** (even minimal: a second Git repo with a `marketplace.yml` file listing your plugins):

```
# marketplace.yml
name: mavida
description: "Mavida plugins for Claude Code"

plugins:
  - name: hello-world
    repo: https://github.com/mavida/hello-world
    versions: [0.1.0]
```

3. Document the installation in the README:

```
/plugin marketplace add mavida/marketplace
/plugin install hello-world
```

To manage versions, use semantic Git tags (`0.1.0` , `0.2.0` , `1.0.0`). Claude Code respects the version requested in the consumer's manifest.

14.7 Security and operational considerations

A plugin can contain **executable code** (MCP servers in Python/Node, hooks in Bash, auxiliary scripts in skills). Installing a third-party plugin is equivalent to installing an npm or Composer package: you trust the author with the permissions the plugin will require.

Three practical rules:

- **Audit before installing.** Read the manifest, inspect the included skills, check what the MCP servers launch. A checklist: `plugin.yml` what does it declare to install? Do the skills require access to sensitive tools? Do the MCP servers contact external services to which you wouldn't want to give your credentials?
- **Verify the repo's health** of the author: last commit date, open/closed issues, presence of security policy. A plugin abandoned two years ago is a constant risk.
- **Use private marketplaces** for corporate use. For plugins with corporate secrets (API keys, internal URLs) keep the marketplace on a corporate Git repo, not public.

For plugins from known organizations (Anthropic, Vercel Labs, WordPress) the risk is low: they publish under their names, the code is scrutinized by the community, updates are regular. For individual plugins, apply the same caution you would use for a third-party library.

An additional protection: MCP tools exposed by plugins flow into Claude Code's permission system (chapter 9), so you can precisely denylist destructive operations with `permissions.deny` on `mcp__<server>__<tool>` patterns (see 12.6 for examples).

15. Advanced workflows and tips

The previous chapters built the conceptual foundations — commands, Plan Mode, memory, context, security, skills, plugins, subagents, hooks. This chapter is the toolbox: first four **practical workflows** (15.1-15.4) that combine the foundations into concrete daily use scenarios, then six **tips** (15.5-15.10) for those who want to push the tool's efficiency further.

15.1 Onboarding to an existing repository

Prompt: "You've just been assigned to this project. Analyze the structure, identify:

1. Main architectural patterns (MVC, hexagonal, etc.)
2. How authentication is managed
3. Where the integration points with external services are
4. Naming and style conventions
5. Any evident technical debts

Produce an onboarding document in docs/ONBOARDING.md.
Don't modify any other code."

Why it works: - Clear objective with numbered list - Specific output (a file in known position) - Explicit constraint ("don't modify anything else")

15.2 Bug hunting with TDD

Prompt: "Bug report: when a user with 'editor' role tries to modify a 'private' post, they get a 500 error. Log attached: [paste log]."

Required workflow:

1. Activate Plan Mode and analyze the code involved
2. FIRST write a test that reproduces the bug (must fail)
3. Fix the bug with the MINIMAL modification needed
4. Verify that the test passes
5. Run the full suite to exclude regressions"

Why it works: - Forces a disciplined TDD approach - Avoids "quick fixes" that suppress symptoms - The test written first becomes documentation of the bug

15.3 Safe refactoring

Prompt: "The includes/class-order-processor.php module has become unmanageable (800 lines, multiple responsibilities). I want to refactor it."

Phase 1 — CHARACTERIZATION (Plan Mode):

- Identify all currently mixed responsibilities
- Propose a decomposition into smaller classes
- List the tests that MUST exist before touching the code

Stop here and wait for my approval of the plan."

After approval:

"Proceed with Phase 2:

- Write characterization tests that lock the current behavior
- Run them and confirm they all pass
- Make a commit with message 'test: pre-refactoring characterization'"

And then:

```
"Phase 3 – incremental refactoring:
- Extract one responsibility at a time
- After each extraction, run the tests
- If even ONE test fails, stop and ask"
```

15.4 Performance audit

```
Prompt: "Analyze the production build and identify the 5 highest-impact
performance problems. For each one:
- File and lines involved
- Estimated impact (ms, KB, HTTP requests)
- Proposed fix
- Fix complexity (low/medium/high)

Sort by impact/complexity ratio. Don't modify anything."
```

The tips that follow (15.5-15.10) are individual tricks, to pick at will when the use case arises.

15.5 Vim mode

If you come from Vim, enable the mode in `/config` → Editor mode. You'll have navigation with `hjkl`, commands `d`, `y`, `p`, etc.

15.6 Custom slash commands

You can create custom slash commands by saving Markdown files in `.claude/commands/`:

```
←!— .claude/commands/security-audit.md →
Run a security audit focused on:
1. SQL injection in direct queries
2. XSS in unescaped output
3. CSRF without nonce verification
4. Path traversal in filesystem operations
5. Hardcoded credentials

For each issue found: file, line, severity (low/medium/high/critical),
suggested fix.
```

Now from a session you can launch `/security-audit` and Claude executes the saved prompt.

15.7 Headless mode for CI/CD

The `-p` (print) flag executes Claude in non-interactive mode, perfect for pipelines:

```
# GitHub Actions example
claude -p "Review the changes in this PR and flag any security issues" \
  --output-format json > review.json
```

The `--output-format json` produces structured output parseable by subsequent steps.

15.8 Session recap

If you leave the terminal and return after 3+ minutes, Claude Code automatically shows a summary of what was done. Great for context-switching. You can force it with `/recap`.

15.9 Strategic Git checkpoints

Before risky tasks, ask explicitly:

```
"Before proceeding, make a commit with message 'pre-refactoring checkpoint' so we have a
safe return point."
```

If something goes wrong, `git reset --hard HEAD~1` brings you back to the previous point.

15.10 Conversation forks

Press `Esc` twice to go back to a previous message and re-edit it. Creates a "branch" of the conversation — useful when a prompt didn't give the desired result and you want to reformulate without losing everything.

16. Conclusions: why CLI and not just chat

After tackling installation, commands, Plan Mode, CLAUDE.md, Skills, and everything else, a legitimate question remains worth making explicit: why use Claude Code CLI when I can simply paste code into a browser chat?

Chat remains a perfectly valid tool, and indeed in some scenarios it's the most effective choice. But three differences make the CLI a qualitatively different tool, not just a variant of the channel.

16.1 Persistent context: stop introducing yourself every time

In traditional chat, every new session starts from zero. The model knows nothing about your project: naming conventions, stack, hard rules, build commands. Every time you have to re-explain, or hope Claude guesses from the code you paste.

With CLI, `CLAUDE.md` **is a permanent contract**. It's automatically read at every session, it's hierarchical (global user, monorepo, project), and you can regenerate it with `/init` every time the project evolves. The `--continue` and `--resume` commands let you resume conversations interrupted days before in the same context state.

This changes the way you think about your development environment: instead of mentally reconfiguring AI at every opening, you configure once and work. The same logic as `.editorconfig`, `.eslintrc`, `.gitignore` files that consolidate project rules once for all.

16.2 Agentic autonomy: it executes, not just suggests

In chat, Claude produces text. You're the human bridge: you copy the code into the editor, save, open the terminal, run the test, read the error, return to chat, paste the error, wait for the correction, copy the correction, paste into the editor, and so on. Every step is an interruption of the flow.

In CLI, Claude **is** the terminal. It reads files, executes commands, sees errors, makes commits, opens PRs, queries databases via MCP, navigates folders, runs tests. The "write → test → fix" cycle becomes a continuous conversation without leaving the tool:

```
You:      "The auth/login.test.js test fails. Figure out why and fix it."

Claude:   [reads the test]
          [reads the code under test]
          [runs npm test -- auth/login.test.js]
          [analyzes the output]
          [identifies the bug]
          [modifies the code]
          [re-runs the test]
          [all green]
          "Fixed. The problem was in token expiry handling.
           I modified validateToken() at lines 34-38."
```

This autonomy has a flip side — that's why there are entire chapters on security and Plan Mode — but when well managed it multiplies productivity in a non-linear way. You don't do something faster: you do things you simply wouldn't do in chat because the cost of manual orchestration is too high.

16.3 Integration into the real workflow

Professional development isn't just writing code: it's git, test suite, linting, CI/CD, code review, dependencies, environments. Chat lives **alongside** this workflow; CLI lives **inside** it.

Native git. Claude Code makes commits, opens branches, resolves merge conflicts, writes Conventional Commits commit messages, manages stash. You don't explain the diff to it: it reads it directly from `git diff`.

Test and lint in loop. The CLI runs the test suite, reads linter errors, retries until it passes. There's no copy-paste between windows, no "wait while I send you the output".

Headless CI/CD. The `-p` flag turns Claude into a pipeline tool:

```
claude -p "Review the changes in this PR and flag any security issues" \
  --output-format json > review.json
```

Insert this step in a GitHub Actions workflow and you have automatic AI code review on every push. Try doing the same with a browser chat.

16.4 When chat remains the right choice

For honesty: there are cases when opening chat.anthropic.com is the better move:

- **Conceptual brainstorming** without specific code — “What patterns can I use to implement a feature flag system?”
- **Learning a new framework** — you need pedagogy, not execution
- **Abstract architectural questions** — “Is it worth introducing CQRS in this context?”
- **Review of single snippets** from code you don’t have locally
- **Discussions with Claude on non-coding topics** — writing, document analysis, planning

Practical rule: if the answer is “**code to integrate into my project**”, use the CLI. If the answer is “**an idea, a principle, an explanation**”, chat is enough.

16.5 In summary

Claude Code CLI isn’t “Claude-in-chat with a different interface”. It’s an agentic tool that transforms a linguistic assistant into an **operational junior colleague**: it can do things, not just suggest them. For those who develop professionally, the difference is the same as between having a consultant who sends emails and having an intern at the next desk. Both useful, different contexts.

My advice, if you’re starting out: install Claude Code, try a small and non-critical project, write a decent `CLAUDE.md`, always use Plan Mode for non-trivial tasks, and after a week evaluate. The curve is steep the first two days, then it flattens out.

Appendix A — Glossary

Recurring terms in the guide and in the Claude Code ecosystem, useful as a quick reference.

Agent (agentic) — AI system capable of executing actions in the real world (commands, file modifications, API calls), not just producing text. Claude Code is an agent unlike classic chat.

Agent Teams — Experimental Claude Code feature (requires the environment variable `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1`) that allows multiple independent instances of Claude to run in true simultaneous parallelism, coordinated via shared task list and message mailbox. Not to be confused with multi-delegation to subagents in foreground, which is sequential instead.

Auto Memory — Persistent memory that Claude Code feeds autonomously during sessions (introduced in v2.1.59). Unlike `CLAUDE.md`, written by the user, Auto Memory is written by the model: it accumulates learnings, patterns, and recurring corrections of the project. It lives in `~/.claude/projects/<project>/memory/` and is organized into an index (`MEMORY.md`) plus topic files loaded on-demand.

Chain of Thought (CoT) — Prompt engineering technique that explicitly asks the model to “reason step by step” before answering, instead of producing the conclusion directly. Forces explication of logical steps and improves accuracy on complex tasks (debugging, architectural decisions, multi-phase problems).

CLAUDE.md — Markdown file in the project root that contains persistent context: stack, conventions, commands, rules. Automatically read at every session.

CLI (Command Line Interface) — Command-line interface. The `claude` tool is used from the terminal instead of from the browser.

Context engineering — Discipline complementary to prompt engineering: instead of focusing on how to formulate the request, it concerns what information to make available to the model before asking it. For Claude Code CLI it materializes in `CLAUDE.md`, Auto Memory, and delegation via subagent. Guiding principle: “better little well-ordered context than much chaotic context”.

Few-shot prompting — Technique that teaches the desired style or format by providing two or more examples before the actual request. Particularly effective for voice consistency (FAQs, microcopy) and reproduction of structured formats hard to describe in words.

Headless mode — Non-interactive execution via `-p` flag. Claude receives a prompt, produces output, exits. Used for CI/CD and automations.

Hook — Script (bash, HTTP, prompt, agent or MCP tool) configured in `settings.json` that intercepts Claude Code lifecycle events: `PreToolUse`, `PostToolUse`, `SessionStart`, `UserPromptSubmit`, and others. Used to validate, log, inject context, or block operations. Different from Subagent (executes delegated work) and from Skill (enriches the main agent’s context): a Hook acts **around** the main agent without being part of it. For complete treatment see section 13.

Hope Coding — Prompt engineering antipattern: launching generic requests to AI “hoping” it guesses what we wanted, without specifying context, constraints, or output format. Produces random results and contrasts with conscious Vibe coding (see Vibe coding entry) based on structured prompts.

JSON-RPC — Textual communication protocol (based on JSON) for remote procedure calls. It’s the base layer on which MCP packages all its messages between client and server. Defines request, response, and notification with a standardized format.

MCP (Model Context Protocol) — Open protocol, open-sourced by Anthropic in November 2024, that standardizes the way an AI application (host) connects to external data sources and tools. Client-server model based on JSON-RPC 2.0; transport via stdio (local) or HTTP+SSE (remote). Three primitives: tools, resources, prompts. See chapter 11 for complete treatment.

MCP server — Process that implements the MCP protocol and exposes one or more functionalities (tool, resource, prompt) to a compatible AI host. Can be written in any language for which an SDK exists (Python, TypeScript, Java, C#, Rust, Kotlin, Swift). Typically local (stdio) for personal integrations, hosted (HTTP+SSE) for team integrations.

MCP tool — One of the three primitives of an MCP server: callable function exposed by a server. It has a name, AI-readable textual description, JSON schema of arguments. When the model decides to call it, the host sends a `tools/call` JSON-RPC request to the server. It’s the most used primitive in custom MCP servers.

MEMORY.md — Index file of a project’s Auto Memory, located in `~/.claude/projects/<project>/memory/`. Loaded at every session (limit ~200 lines, ~25 KB), it lists and describes the topic files of the folder, which are then loaded on-demand when their content is relevant.

Meta-prompting — Prompt engineering technique that consists of asking the AI itself to write the prompt to use in a subsequent session. Pattern: the model, in the role of “Expert Prompt Engineer”, analyzes raw specs, asks clarifying questions, produces a final structured prompt. Useful for new or complex tasks worthy of formalization.

Native installer — Official installation method introduced by Anthropic in 2025: a `curl` or `PowerShell` command without Node.js dependencies, with auto-update.

OAuth — Authentication protocol used at first launch of `claude`. Opens the browser, you log in with your Anthropic account, the session persists.

Panel of Experts (Round table) — Prompt engineering technique that simulates a discussion among virtual experts, each with their own viewpoint and area of expertise. Particularly precious for exploring an idea, putting one’s beliefs into question, choosing a stack, or stress-testing an architecture: the value isn’t in the final synthesis but in the explication of trade-offs each decision implies. See section 6.4.3 for the complete prompt template.

Plan Mode — Read-only mode activated via `/plan` or by cycling with `Shift+Tab` (which scrolls between `default` → `acceptEdits` → `plan` → ...). Claude analyzes and proposes a plan but doesn’t modify anything until you approve.

Plugin — Package distributed via marketplace that extends Claude Code with slash commands, agents, and skills. Managed with `claude plugin install`.

Prompt engineering — Discipline of formulating effective requests for an LLM. Articulated in four fundamental ingredients (context, task, constraints, output format) plus an optional one (role). On top 2026 models, “role prompting” is downsized in favor of structural constraints and the use of XML-like delimiters (`<context>`, `<task>`, `<constraints>`, `<output_format>`). See section 6 for complete treatment.

Prompt injection — Attack in which malicious instructions are injected into files, comments, or responses from external services to manipulate AI behavior.

REPL (Read-Eval-Print Loop) — Interactive read-execute-print cycle. Claude Code’s interactive session is a REPL.

Session — Ongoing conversation with Claude Code, persistent across restarts. Each session has its own context and history.

Skill — Specialized module (folder with `SKILL.md`) that Claude automatically activates when the skill description matches the task context. Not invoked with slash commands.

Slash command — Command starting with `/` inside an interactive session (e.g., `/init`, `/compact`, `/plan`). Different from launch flags that start with `--`.

Subagent — Isolated Claude instance created by the `Task` tool to execute searches or specialized tasks without “polluting” the main session context.

Token — Unit of text measurement for an LLM (approximately 4 characters in English, slightly less in Italian). API costs are calculated in input and output tokens. Claude Code uses tokens every time it reads a file, receives a prompt, or produces a response.

Vibe coding — Term that became popular in 2024-2025 to describe AI-assisted development style: instead of writing code manually, you write a structured prompt describing what it should do, and the AI generates the implementation.

WSL2 (Windows Subsystem for Linux) — Linux environment integrated in Windows 10/11. Recommended for using Claude Code on Windows avoiding many compatibility issues.

Appendix B — Sources

To explore further or verify updated specifications:

Anthropic official documentation:

- **Claude Code overview:** <https://docs.claude.com/en/docs/claude-code/overview>
- **CLI reference:** <https://code.claude.com/docs/en/cli-reference>
- **Setup and installation:** <https://code.claude.com/docs/en/setup>
- **Interactive mode and shortcuts:** <https://code.claude.com/docs/en/interactive-mode>
- **Official cheatsheet:** <https://support.claude.com/en/articles/14553413-claude-code-cheatsheet>
- **Prompt engineering — Claude 4 best practices:** <https://docs.anthropic.com/en/docs/build-with-claude/prompt-engineering/claude-4-best-practices>
- **GitHub repository:** <https://github.com/anthropics/claude-code>
- **npm package:** <https://www.npmjs.com/package/@anthropic-ai/claude-code>

Prompt engineering and context engineering resources:

- **Prompting Guide:** <https://www.promptingguide.ai/>
- **Elastic — Context engineering vs. prompt engineering:** <https://www.elastic.co/search-labs/blog/context-engineering-vs-prompt-engineering>

Official skill repositories:

- **Anthropic — Skills:** <https://github.com/anthropics/skills>
- **WordPress — Agent Skills:** <https://github.com/WordPress/agent-skills>
- **Vercel Labs — Agent Skills:** <https://github.com/vercel-labs/agent-skills>

- **Trail of Bits — Skills:** <https://github.com/trailofbits/skills>
- **Skills directory community:** <https://skills.sh>

Community resources cited in the guide:

- **Caveman skill (Julius Brussee):** <https://github.com/JuliusBrussee/caveman>
- **Superpowers (Jesse Vincent):** <https://github.com/obra/superpowers>

MCP (Model Context Protocol) — official sources:

- **MCP documentation:** <https://modelcontextprotocol.io/>
- **Protocol specification:** <https://spec.modelcontextprotocol.io/>
- **Anthropic — MCP announcement (November 2024):** <https://www.anthropic.com/news/model-context-protocol>
- **Official Python SDK:** <https://github.com/modelcontextprotocol/python-sdk>
- **Official TypeScript SDK:** <https://github.com/modelcontextprotocol/typescript-sdk>
- **Reference MCP servers (Anthropic):** <https://github.com/modelcontextprotocol/servers>
- **Anthropic MCP marketplace:** <https://www.anthropic.com/mcp>
- **Community MCP marketplace (Glama):** <https://glama.ai/mcp>
- **Community MCP marketplace (Smithery):** <https://smithery.ai/>

WordPress REST API — official sources:

- **REST API Handbook:** <https://developer.wordpress.org/rest-api/>
- **Application Passwords:** <https://make.wordpress.org/core/2020/11/05/application-passwords-integration-guide/>
- **REST API: posts endpoint:** <https://developer.wordpress.org/rest-api/reference/posts/>

Note on the Author

My name is **Maurizio Pelizzone**. I'm a Senior Software Architect, Co-founder of **Mavida snc** (Turin, Italy) since 2001, and a technical trainer specialized in the WordPress ecosystem.

In over twenty years I've designed and brought to production more than 200 web projects for national clients, specializing in enterprise WordPress architectures, custom plugin and theme development, legacy platform migrations, and performance/security optimization. Since 2010 I've been a **recurring speaker at WordCamp Italy events** (Milan, Bologna, Turin, WordCamp Italia), with talks on Custom Post Types, Hardening & Security, Solutions Architecture, and Full Site Editing.

Since 2023 I've added a "superpower" to my workflow: Artificial Intelligence, using it as a strategic lever to write cleaner, more scalable, smarter code. Not as a replacement for craft, but as an amplifier: twenty years of PHP, WordPress, and software architectures allow me to distinguish a brilliant suggestion from a risky shortcut, and to direct AI toward solutions that hold the weight of production.

Today I combine technical consulting for SMBs and national companies with **training on AI tools, prompt engineering, vibe coding, and automation with N8N**, convinced that the combination of artisanal code experience and conscious AI use is the direction our craft is evolving in.

This guide was born from that conviction: a tool like Claude Code doesn't replace the developer, but changes the way they work. To use it well requires the same virtues as always — rigor, curiosity, the ability to verify — plus a bit of new discipline.

✉ maurizio@mavida.com — [🔗 linkedin.com/in/mauriziopelizzone](https://www.linkedin.com/in/mauriziopelizzone) — [🌐 mavida.com](https://www.mavida.com)

